

Learning From Networks

—*Algorithms, Theory, & Applications*

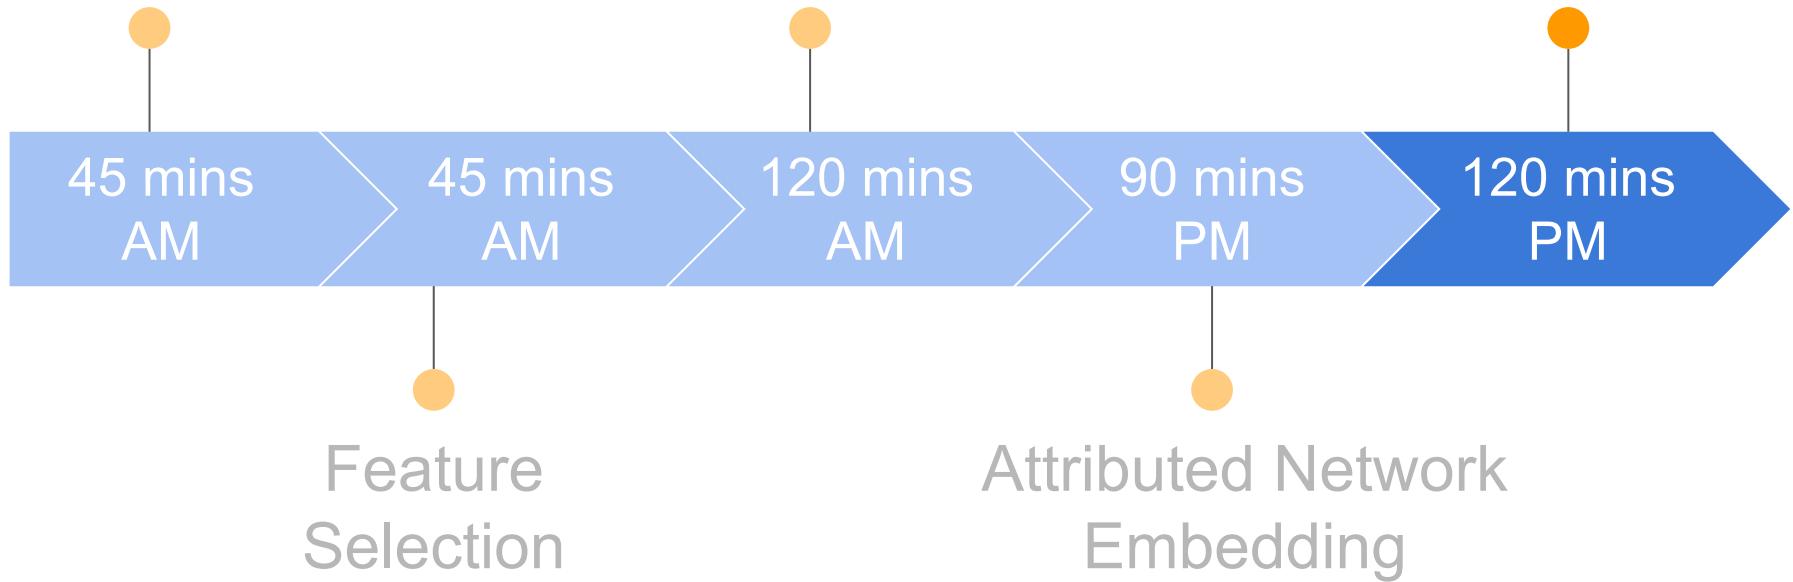
Xiao Huang, Peng Cui, Yuxiao Dong, Jundong Li, Huan Liu, Jian Pei, Le Song,
Jie Tang, Fei Wang, Hongxia Yang, Wenwu Zhu

xhuang@tamu.edu; cuip@tsinghua.edu.cn; yuxxdong@microsoft.com; jundongl@asu.edu;
huan.liu@asu.edu; jpei@cs.sfu.ca; le.song@antfin.com; jietang@tsinghua.edu.cn;
few2001@med.cornell.edu; yang.yhx@alibaba-inc.com; wwzhu@tsinghua.edu.cn;

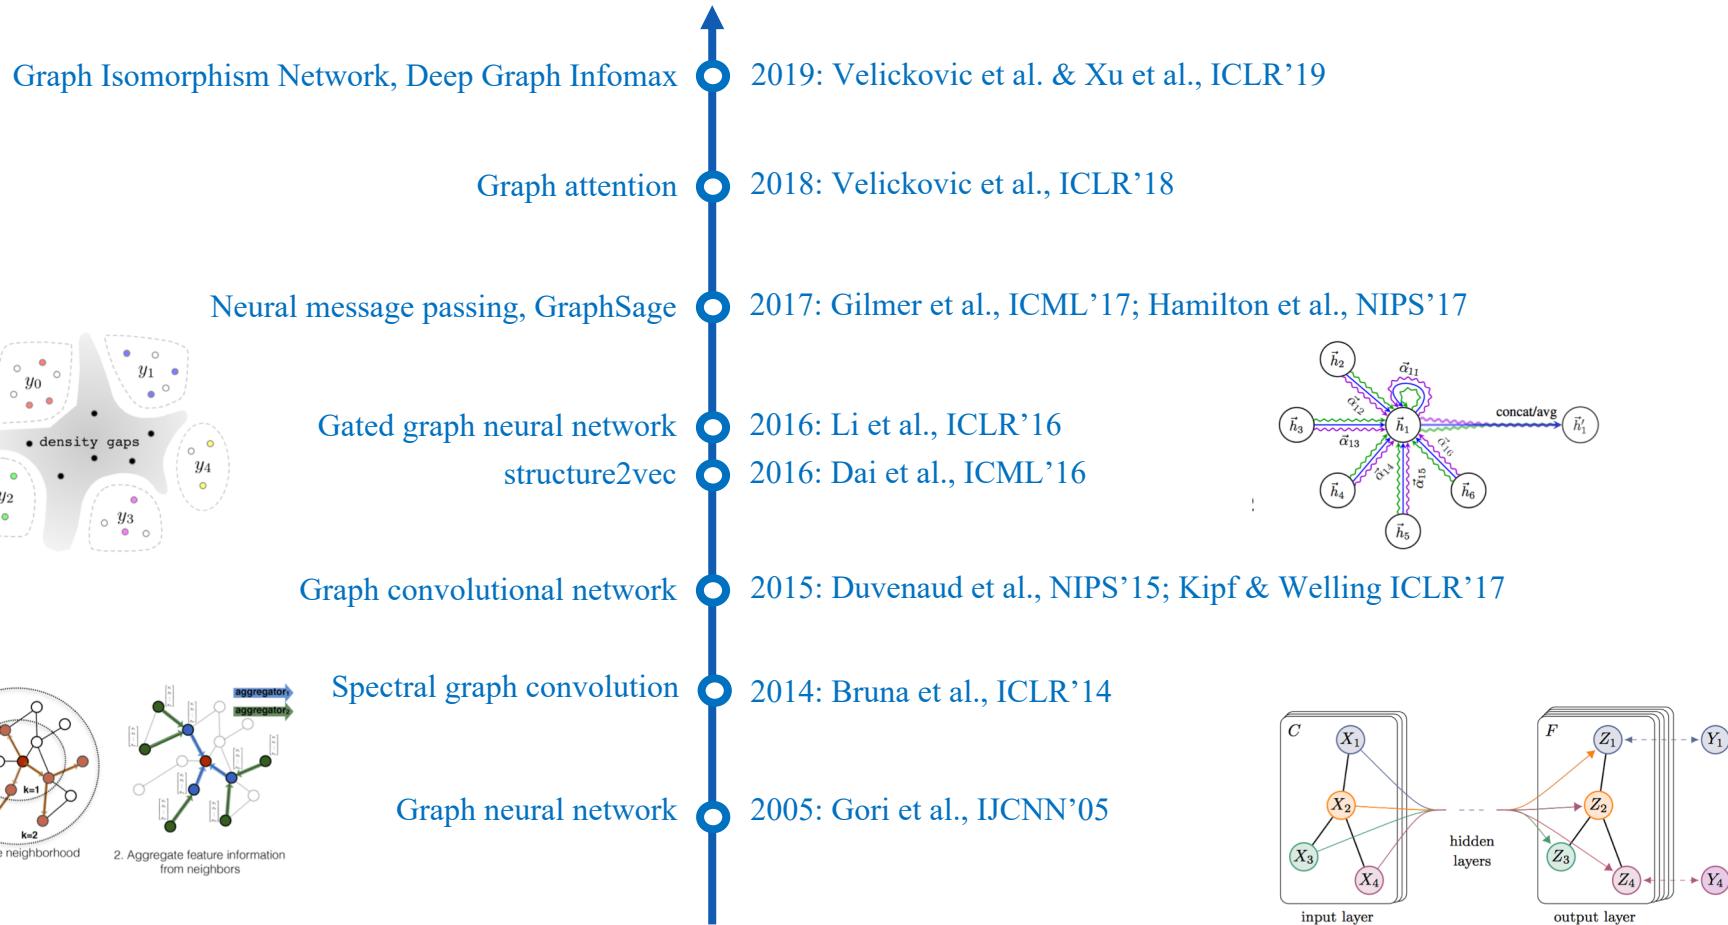
Motivations

Network Embedding

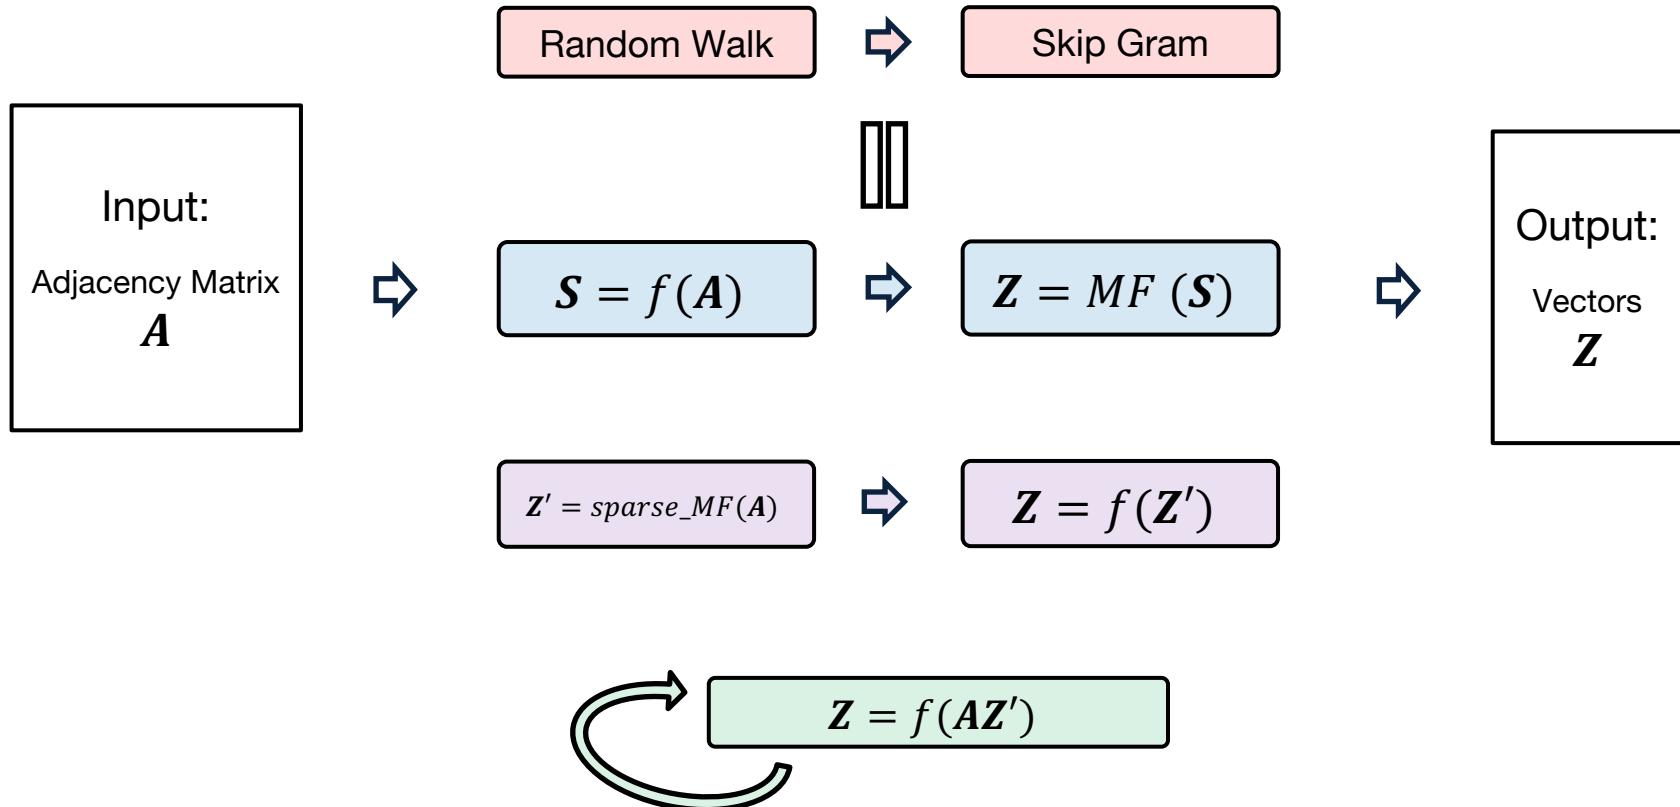
Graph Neural Networks



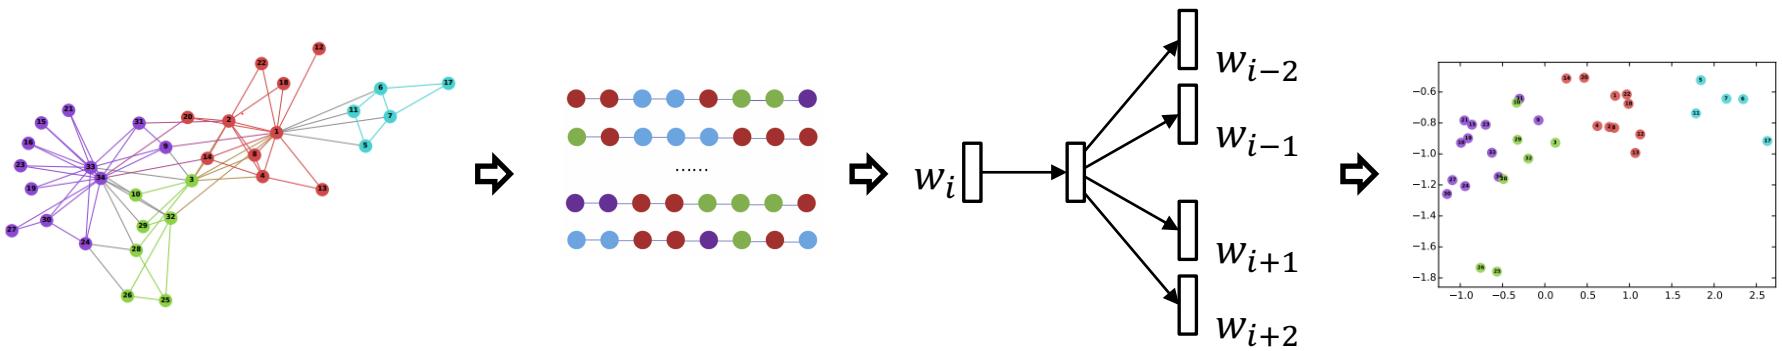
Graph neural networks



Connecting NE with graph neural networks



Network embedding: DeepWalk



Random walk strategies

- Random Walk

- DeepWalk (walk length > 1)
- LINE (walk length = 1)

- Biased Random Walk

- node2vec (2-order random walk)
- metapath2vec (heterogeneous random walk)

1. Perozzi et al. **DeepWalk**: Online learning of social representations. In *KDD'14*. **Most Cited Paper in KDD'14**.
2. Tang et al. **LINE**: Large scale information network embedding. In *WWW'15*. **Most Cited Paper in WWW'15**.
3. Grover and Leskovec. **node2vec**: Scalable feature learning for networks. In *KDD'16*. **2nd Most Cited Paper in KDD'16**.
4. Dong et al. **metapath2vec**: scalable representation learning for heterogeneous networks. In *KDD 2017*. **Most Cited Paper in KDD'17**.

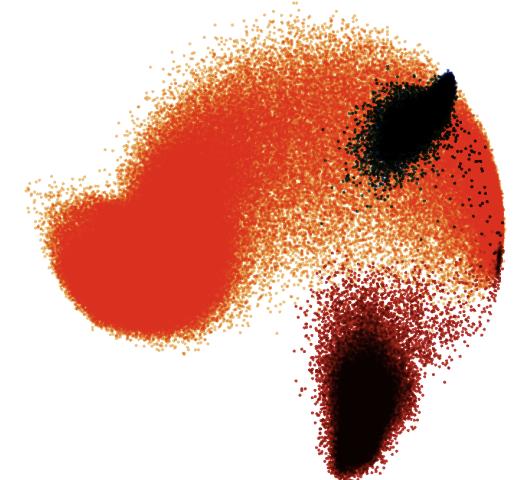
Application: Embedding Heterogeneous Academic Graph

	225,572,477
Papers	
	244,499,947
Authors	
	664,891
Topics	
	4,397
Conferences	
	48,758
Journals	
	25,554
Institutions	

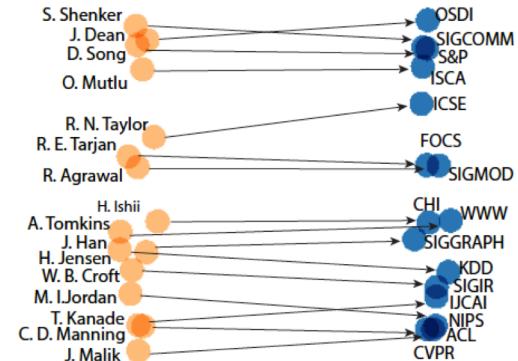
Microsoft Academic Graph



metapath2vec



- <https://academic.microsoft.com/>
- <https://www.openacademic.ai/oagi/>
- metapath2vec: scalable representation learning for heterogeneous networks. In *KDD* 2017.



Application 1: Related Venues

Science

Science, also widely referred to as Science Magazine, is the peer-reviewed academic journal of the American Association for the Advancement of Science (AAAS) and one of the world's top academic journals. It was first published in 1880, is currently circulated weekly and has a subscriber base of around 130,000. Because institutional subscriptions and online access serve a larger audience, its estimated readership is 570,400 people.

Website links: scienmag.org, en.wikipedia.org

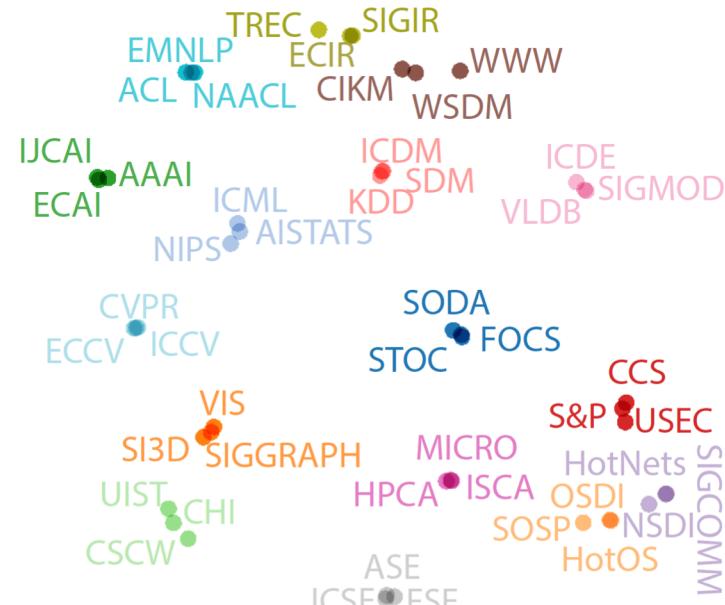
RELATED JOURNALS

[Nature](#)

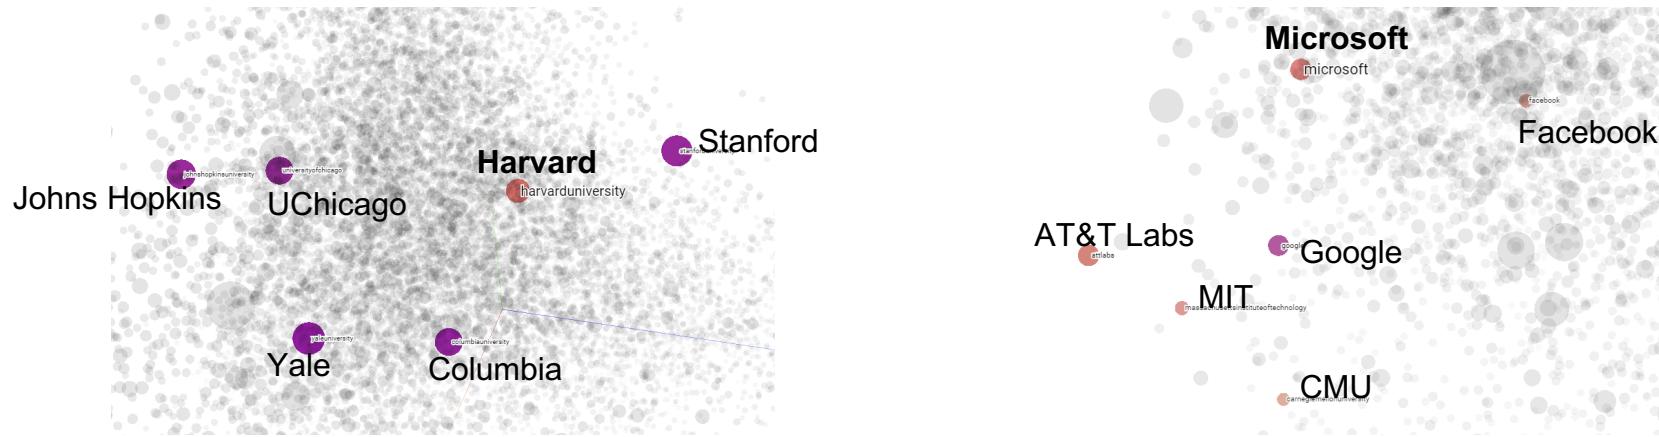
[Proceedings of the National Academy of Sciences of the United States of America](#)

[Nature Communications](#)

- <https://academic.microsoft.com/>
- <https://www.openacademic.ai/oagi/>
- metapath2vec: scalable representation learning for heterogeneous networks. In *KDD* 2017.



Application 2: Similarity Search (Institution)



- <https://academic.microsoft.com/>
- <https://www.openacademic.ai/oag/>
- metapath2vec: scalable representation learning for heterogeneous networks. In *KDD* 2017.

What are the **fundamentals**
underlying random-walk + skip-gram based
network embedding models?

Unifying DeepWalk, LINE, PTE, & node2vec as Matrix Factorization

- DeepWalk $\log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$
- LINE $\log \left(\frac{\text{vol}(G)}{b} \mathbf{D}^{-1} \mathbf{A} \mathbf{D}^{-1} \right)$
- PTE $\log \left(\begin{bmatrix} \alpha \text{vol}(G_{\text{ww}}) (\mathbf{D}_{\text{row}}^{\text{ww}})^{-1} \mathbf{A}_{\text{ww}} (\mathbf{D}_{\text{col}}^{\text{ww}})^{-1} \\ \beta \text{vol}(G_{\text{dw}}) (\mathbf{D}_{\text{row}}^{\text{dw}})^{-1} \mathbf{A}_{\text{dw}} (\mathbf{D}_{\text{col}}^{\text{dw}})^{-1} \\ \gamma \text{vol}(G_{\text{lw}}) (\mathbf{D}_{\text{row}}^{\text{lw}})^{-1} \mathbf{A}_{\text{lw}} (\mathbf{D}_{\text{col}}^{\text{lw}})^{-1} \end{bmatrix} \right) - \log b$
- node2vec $\log \left(\frac{\frac{1}{2T} \sum_{r=1}^T (\sum_u \mathbf{X}_{w,u} \underline{\mathbf{P}}_{c,w,u}^r + \sum_u \mathbf{X}_{c,u} \underline{\mathbf{P}}_{w,c,u}^r)}{b (\sum_u \mathbf{X}_{w,u}) (\sum_u \mathbf{X}_{c,u})} \right)$

\mathbf{A} Adjacency matrix

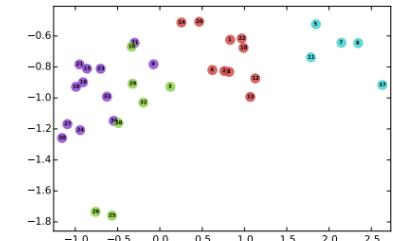
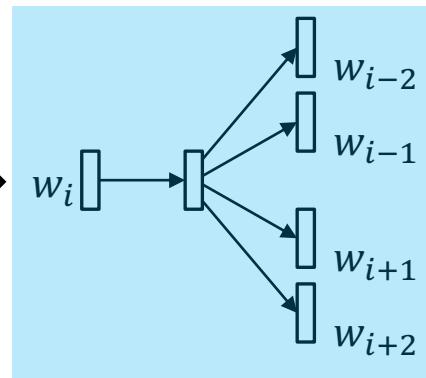
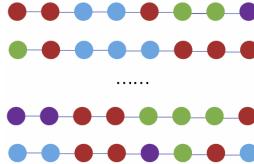
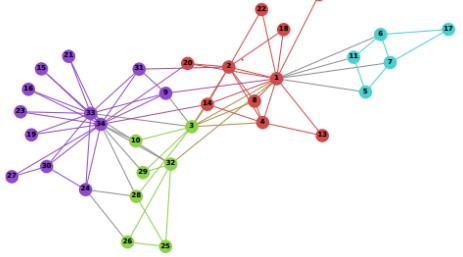
\mathbf{D} Degree matrix

$$\text{vol}(G) = \sum_i \sum_j A_{ij}$$

b : #negative samples

T : context window size

Understanding random walk + skip gram



$$G = (V, E)$$

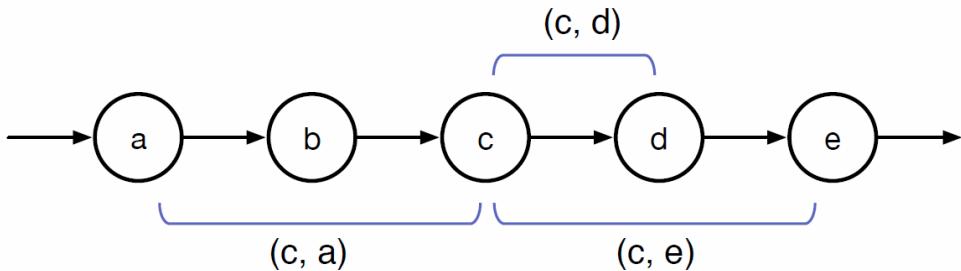
- Adjacency matrix A
- Degree matrix D
- Volume of G : $\text{vol}(G)$

?

$$\log\left(\frac{\#(w, c)|\mathcal{D}|}{b\#(w)\#(c)}\right)$$

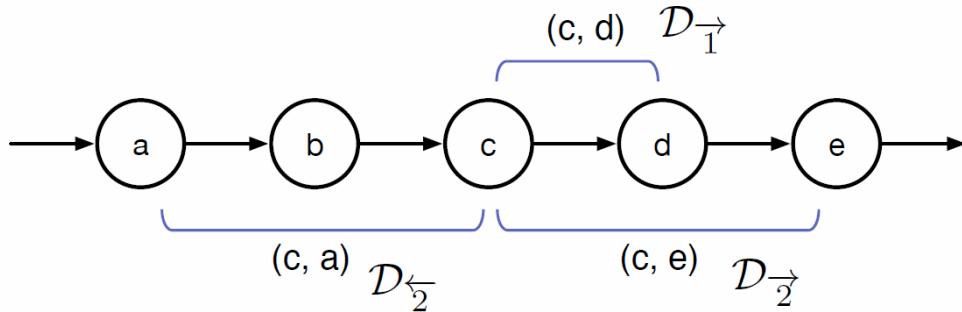
- $\#(w, c)$: co-occurrence of w & c
- $\#(w)$: occurrence of word w
- $\#(c)$: occurrence of context c
- $|\mathcal{D}|$: number of word-context pairs

Understanding random walk + skip gram



Suppose the multiset \mathcal{D} is constructed based on random walk on graphs, can we interpret $\log \frac{\#(w,c)|\mathcal{D}|}{b\#(w)\#(c)}$ with graph structures?

Understanding random walk + skip gram



- Partition the multiset \mathcal{D} into several sub-multisets according to the way in which each node and its context appear in a random walk node sequence.
- More formally, for $r = 1, 2, \dots, T$, we define

$$\mathcal{D}_r^{\rightarrow} = \{(w, c) : (w, c) \in \mathcal{D}, w = w_j^n, c = w_{j+r}^n\}$$

$$\mathcal{D}_r^{\leftarrow} = \{(w, c) : (w, c) \in \mathcal{D}, w = w_{j+r}^n, c = w_j^n\}$$

Distinguish direction
and distance

Understanding random walk + skip gram

$$\log \left(\frac{\#(w,c) |\mathcal{D}|}{b\#(w) \cdot \#(c)} \right) = \log \left(\frac{\frac{\#(w,c)}{|\mathcal{D}|}}{b \frac{\#(w)}{|\mathcal{D}|} \frac{\#(c)}{|\mathcal{D}|}} \right)$$

the length of random walk $L \rightarrow \infty$

$$\frac{\#(w,c)}{|\mathcal{D}|} = \frac{1}{2T} \sum_{r=1}^T \left(\frac{\#(w,c)_{\vec{r}}}{|\mathcal{D}_{\vec{r}}|} + \frac{\#(w,c)_{\leftarrow r}}{|\mathcal{D}_{\leftarrow r}|} \right)$$
$$\frac{\#(w,c)_{\vec{r}}}{|\mathcal{D}_{\vec{r}}|} \xrightarrow{p} \frac{d_w}{\text{vol}(G)} (\mathbf{P}^r)_{w,c}$$
$$\frac{\#(w,c)_{\leftarrow r}}{|\mathcal{D}_{\leftarrow r}|} \xrightarrow{p} \frac{d_c}{\text{vol}(G)} (\mathbf{P}^r)_{c,w}$$

$$\frac{\#(w,c)}{|\mathcal{D}|} \xrightarrow{p} \frac{1}{2T} \sum_{r=1}^T \left(\frac{d_w}{\text{vol}(G)} (\mathbf{P}^r)_{w,c} + \frac{d_c}{\text{vol}(G)} (\mathbf{P}^r)_{c,w} \right) \quad \mathbf{P} = \mathbf{D}^{-1} \mathbf{A}$$

$$\frac{\#(w)}{|\mathcal{D}|} \xrightarrow{p} \frac{d_w}{\text{vol}(G)}$$

$$\frac{\#(c)}{|\mathcal{D}|} \xrightarrow{p} \frac{d_c}{\text{vol}(G)}$$

Understanding random walk + skip gram

$$\log \left(\frac{\#(w, c) |\mathcal{D}|}{b \#(w) \cdot \#(c)} \right) = \log \left(\frac{\frac{\#(w, c)}{|\mathcal{D}|}}{b \frac{\#(w)}{|\mathcal{D}|} \frac{\#(c)}{|\mathcal{D}|}} \right)$$

the length of random walk $L \rightarrow \infty$

$$\frac{\#(w, c)}{|\mathcal{D}|} \xrightarrow{p} \frac{1}{2T} \sum_{r=1}^T \left(\frac{d_w}{\text{vol}(G)} (\mathbf{P}^r)_{w,c} + \frac{d_c}{\text{vol}(G)} (\mathbf{P}^r)_{c,w} \right) \quad \mathbf{P} = \mathbf{D}^{-1} \mathbf{A}$$

$$\frac{\#(w)}{|\mathcal{D}|} \xrightarrow{p} \frac{d_w}{\text{vol}(G)} \quad \frac{\#(c)}{|\mathcal{D}|} \xrightarrow{p} \frac{d_c}{\text{vol}(G)}$$

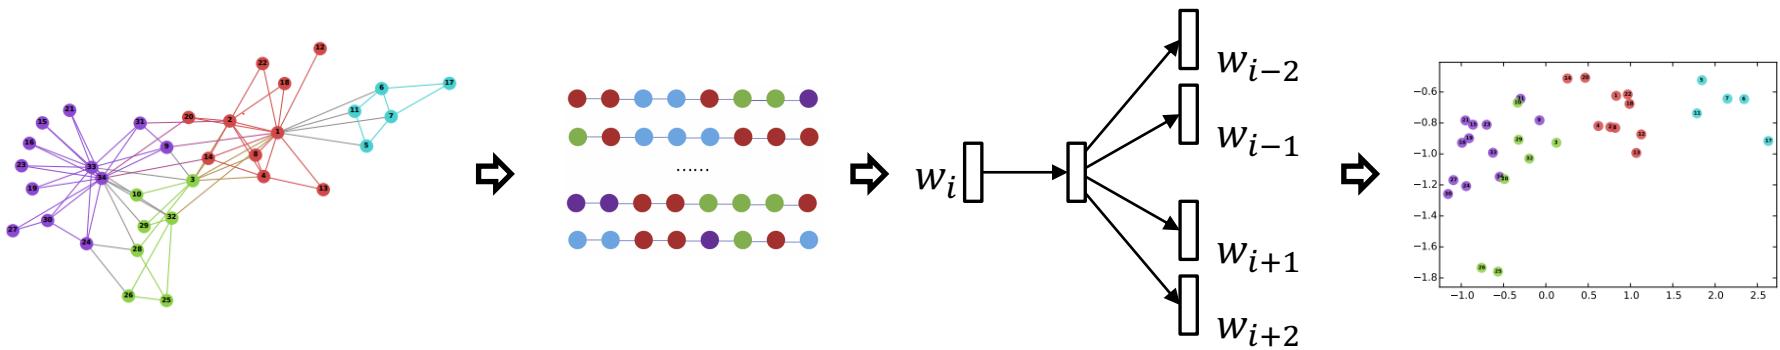
$$\begin{aligned} \frac{\#(w, c) |\mathcal{D}|}{\#(w) \cdot \#(c)} &= \frac{\frac{\#(w, c)}{|\mathcal{D}|}}{\frac{\#(w)}{|\mathcal{D}|} \cdot \frac{\#(c)}{|\mathcal{D}|}} \xrightarrow{p} \frac{\frac{1}{2T} \sum_{r=1}^T \left(\frac{d_w}{\text{vol}(G)} (\mathbf{P}^r)_{w,c} + \frac{d_c}{\text{vol}(G)} (\mathbf{P}^r)_{c,w} \right)}{\frac{d_w}{\text{vol}(G)} \cdot \frac{d_c}{\text{vol}(G)}} \\ &= \frac{\text{vol}(G)}{2T} \left(\frac{1}{d_c} \sum_{r=1}^T (\mathbf{P}^r)_{w,c} + \frac{1}{d_w} \sum_{r=1}^T (\mathbf{P}^r)_{c,w} \right) \end{aligned}$$

Understanding random walk + skip gram

$$\frac{\#(w, c) |\mathcal{D}|}{\#(w) \cdot \#(c)} \xrightarrow{p} \frac{\text{vol}(G)}{2T} \left(\frac{1}{d_c} \sum_{r=1}^T (\mathbf{P}^r)_{w,c} + \frac{1}{d_w} \sum_{r=1}^T (\mathbf{P}^r)_{c,w} \right)$$

$$\begin{aligned} & \frac{\text{vol}(G)}{2T} \left(\sum_{r=1}^T \mathbf{P}^r \mathbf{D}^{-1} + \sum_{r=1}^T \mathbf{D}^{-1} (\mathbf{P}^r)^\top \right) \\ &= \frac{\text{vol}(G)}{2T} \left(\sum_{r=1}^T \underbrace{\mathbf{D}^{-1} \mathbf{A} \times \cdots \times \mathbf{D}^{-1} \mathbf{A}}_{r \text{ terms}} \mathbf{D}^{-1} + \sum_{r=1}^T \mathbf{D}^{-1} \underbrace{\mathbf{A} \mathbf{D}^{-1} \times \cdots \times \mathbf{A} \mathbf{D}^{-1}}_{r \text{ terms}} \right) \\ &= \frac{\text{vol}(G)}{T} \sum_{r=1}^T \underbrace{\mathbf{D}^{-1} \mathbf{A} \times \cdots \times \mathbf{D}^{-1} \mathbf{A}}_{r \text{ terms}} \mathbf{D}^{-1} = \text{vol}(G) \left(\frac{1}{T} \sum_{r=1}^T \mathbf{P}^r \right) \mathbf{D}^{-1}. \end{aligned}$$

Understanding random walk + skip gram



DeepWalk is asymptotically and implicitly factorizing

$$\log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

\mathbf{A} Adjacency matrix

\mathbf{D} Degree matrix

$$\text{vol}(G) = \sum_i \sum_j A_{ij}$$

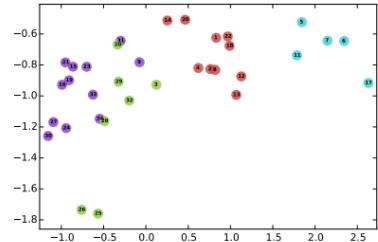
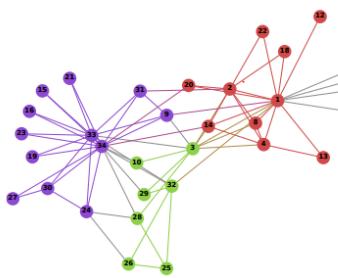
b : #negative samples

T : context window size

Unifying DeepWalk, LINE, PTE, & node2vec as Matrix Factorization

- DeepWalk $\log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$
- LINE $\log \left(\frac{\text{vol}(G)}{b} \mathbf{D}^{-1} \mathbf{A} \mathbf{D}^{-1} \right)$
- PTE $\log \left(\begin{bmatrix} \alpha \text{vol}(G_{\text{ww}}) (\mathbf{D}_{\text{row}}^{\text{ww}})^{-1} \mathbf{A}_{\text{ww}} (\mathbf{D}_{\text{col}}^{\text{ww}})^{-1} \\ \beta \text{vol}(G_{\text{dw}}) (\mathbf{D}_{\text{row}}^{\text{dw}})^{-1} \mathbf{A}_{\text{dw}} (\mathbf{D}_{\text{col}}^{\text{dw}})^{-1} \\ \gamma \text{vol}(G_{\text{lw}}) (\mathbf{D}_{\text{row}}^{\text{lw}})^{-1} \mathbf{A}_{\text{lw}} (\mathbf{D}_{\text{col}}^{\text{lw}})^{-1} \end{bmatrix} \right) - \log b$
- node2vec $\log \left(\frac{\frac{1}{2T} \sum_{r=1}^T (\sum_u \mathbf{X}_{w,u} \underline{\mathbf{P}}_{c,w,u}^r + \sum_u \mathbf{X}_{c,u} \underline{\mathbf{P}}_{w,c,u}^r)}{b (\sum_u \mathbf{X}_{w,u}) (\sum_u \mathbf{X}_{c,u})} \right)$

NetMF: explicitly factorizing the DeepWalk matrix



DeepWalk is asymptotically and implicitly factorizing

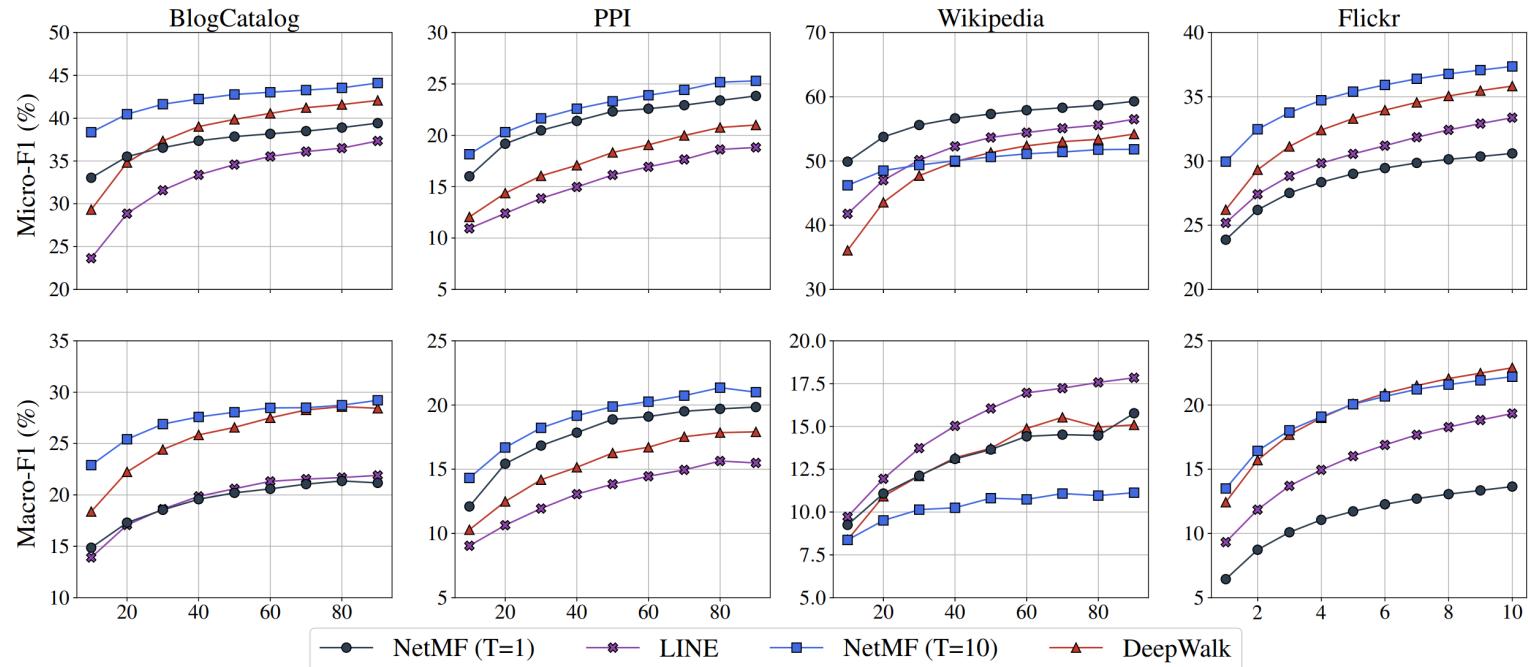
$$\log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

NetMF

1. Construction
2. Factorization

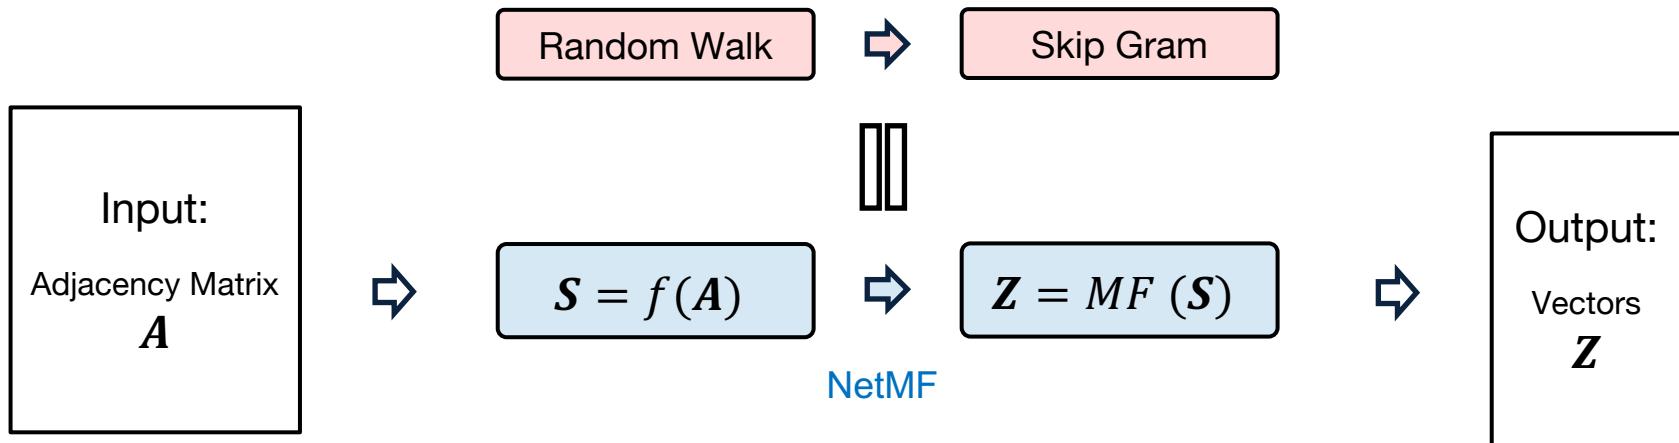
$$\mathbf{S} = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

Results



Predictive performance on varying the ratio of training data;
The x-axis represents the ratio of labeled data (%)

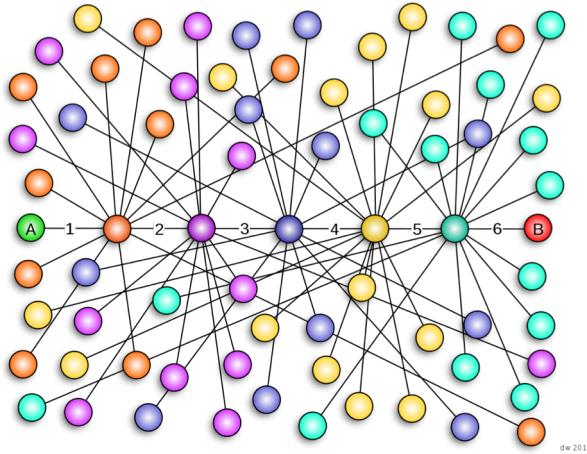
Connecting NE with graph neural networks



Incorporate network structures A into the similarity matrix S , and then factorize S

$$f(A) = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

Challenges



$$\Rightarrow \quad S = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (D^{-1} A)^r \right) D^{-1} \right)$$

dense

NetMF is not practical for very large networks

NetMF

How can we solve this issue?

1. Construction
2. Factorization

$$\mathbf{S} = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

NetSMF--Sparse

How can we solve this issue?

1. **Sparse** Construction
2. **Sparse** Factorization

$$\mathbf{S} = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

Sparsify S

For random-walk matrix polynomial $L = D - \sum_{r=1}^T \alpha_r D (D^{-1} A)^r$

where $\sum_{r=1}^T \alpha_r = 1$ and α_r non-negative

One can construct a **($1 + \epsilon$)-spectral sparsifier \tilde{L}** with $O(n \log n \epsilon^{-2})$ non-zeros

in time $O(T^2 m \epsilon^{-2} \log^2 n)$

$O(T^2 m \epsilon^{-2} \log n)$ for undirected graphs

Suppose $G = (V, E, A)$ and $\tilde{G} = (V, \tilde{E}, \tilde{A})$ are two weighted undirected networks. Let $L = D_G - A$ and $\tilde{L} = D_{\tilde{G}} - \tilde{A}$ be their Laplacian matrices, respectively. We define G and \tilde{G} are $(1 + \epsilon)$ -spectrally similar if

$$\forall x \in \mathbb{R}^n, (1 - \epsilon) \cdot x^\top \tilde{L} x \leq x^\top L x \leq (1 + \epsilon) \cdot x^\top \tilde{L} x.$$

- Dehua Cheng, Yu Cheng, Yan Liu, Richard Peng, and Shang-Hua Teng, Efficient Sampling for Gaussian Graphical Models via Spectral Sparsification, COLT 2015.
- Dehua Cheng, Yu Cheng, Yan Liu, Richard Peng, and Shang-Hua Teng. Spectral sparsification of random-walk matrix polynomials. arXiv:1502.03496.

Sparsify S

For random-walk matrix polynomial $L = D - \sum_{r=1}^T \alpha_r D (D^{-1} A)^r$

where $\sum_{r=1}^T \alpha_r = 1$ and α_r non-negative

One can construct a $(1 + \epsilon)$ -spectral sparsifier \tilde{L} with $O(n \log n \epsilon^{-2})$ non-zeros

in time $O(T^2 m \epsilon^{-2} \log^2 n)$

$$\begin{aligned} S &= \log^\circ \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (D^{-1} A)^r \right) D^{-1} \right) \\ \alpha_1 = \dots = \alpha_T = \frac{1}{T} \quad \Rightarrow \quad &= \log^\circ \left(\frac{\text{vol}(G)}{b} D^{-1} (D - L) D^{-1} \right) \\ &\approx \log^\circ \left(\frac{\text{vol}(G)}{b} D^{-1} (D - \tilde{L}) D^{-1} \right) \end{aligned}$$

NetSMF --- Sparse

- ▶ Construct a random walk matrix polynomial sparsifier, $\tilde{\mathbf{L}}$
- ▶ Construct a NetMF matrix sparsifier.

$$\text{trunc_log}^\circ \left(\frac{\text{vol}(G)}{b} \mathbf{D}^{-1} (\mathbf{D} - \tilde{\mathbf{L}}) \mathbf{D}^{-1} \right)$$

- ▶ Factorize the constructed matrix

	Time	Space
Step 1	$O(MT \log n)$ for weighted networks $O(MT)$ for unweighted networks	$O(M + n + m)$
Step 2	$O(M)$	$O(M + n)$
Step 3	$O(Md + nd^2 + d^3)$	$O(M + nd)$

NetSMF---bounded approximation error

$$\begin{aligned} & \log^\circ \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right) \\ &= \log^\circ \left(\frac{\text{vol}(G)}{b} \mathbf{D}^{-1} (\mathbf{D} - \mathbf{L}) \mathbf{D}^{-1} \right) \rightarrow \mathbf{M} \\ &\approx \log^\circ \left(\frac{\text{vol}(G)}{b} \mathbf{D}^{-1} (\mathbf{D} - \tilde{\mathbf{L}}) \mathbf{D}^{-1} \right) \rightarrow \tilde{\mathbf{M}} \end{aligned}$$

Theorem

The singular value of $\tilde{\mathbf{M}} - \mathbf{M}$ satisfies

$$\sigma_i(\tilde{\mathbf{M}} - \mathbf{M}) \leq \frac{4\epsilon}{\sqrt{d_i d_{\min}}}, \forall i \in [n].$$

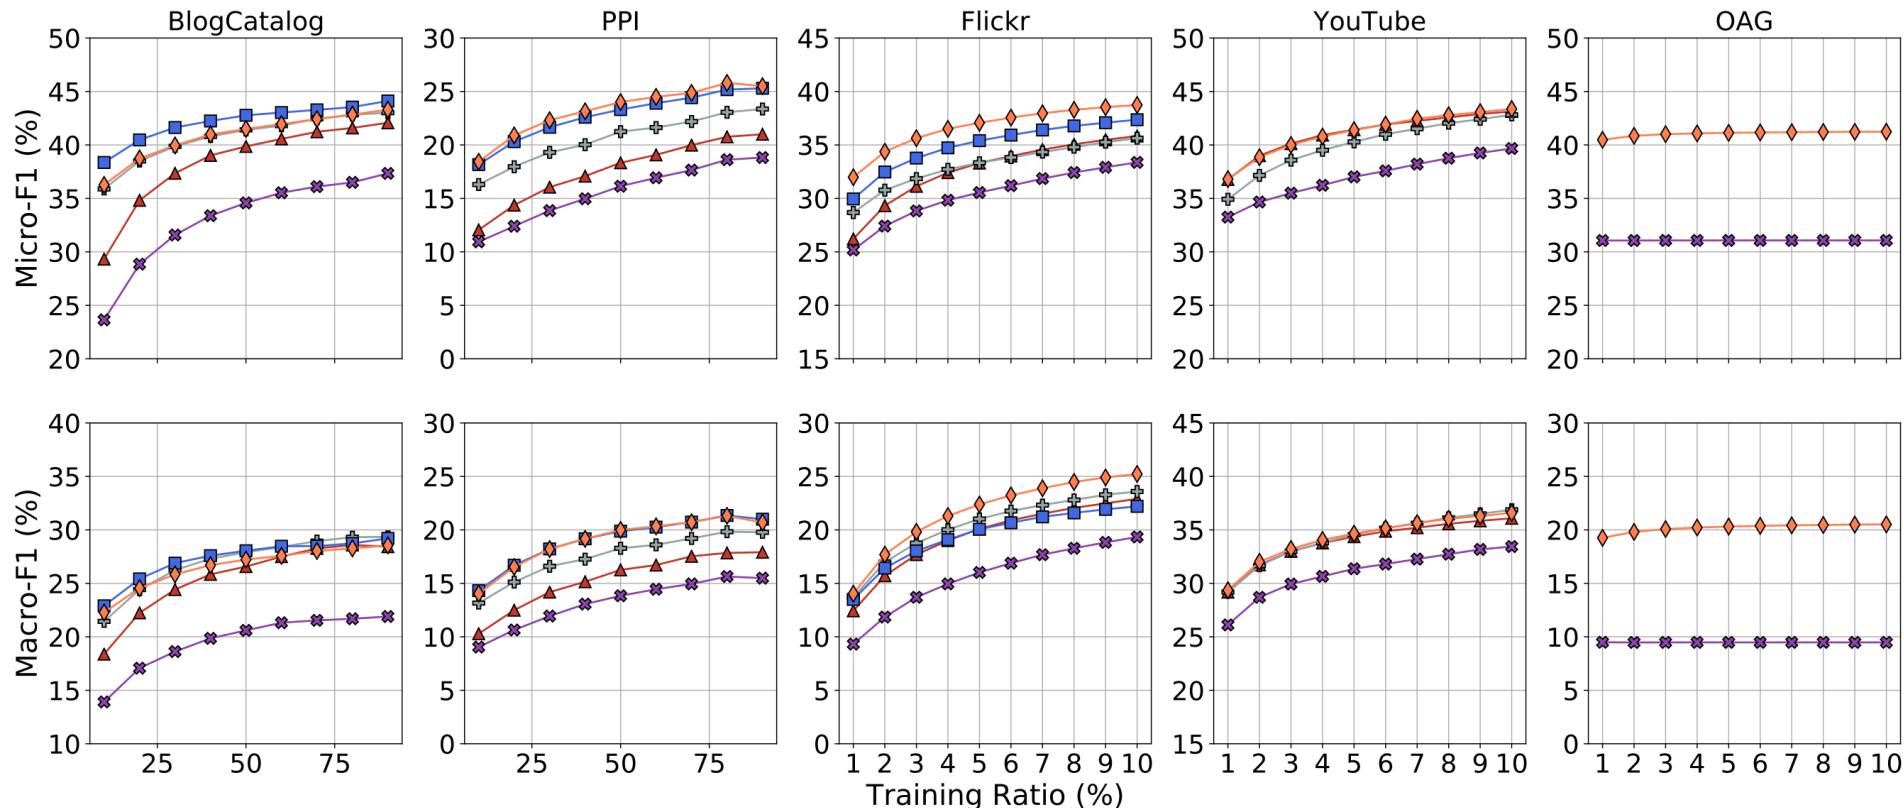
Theorem

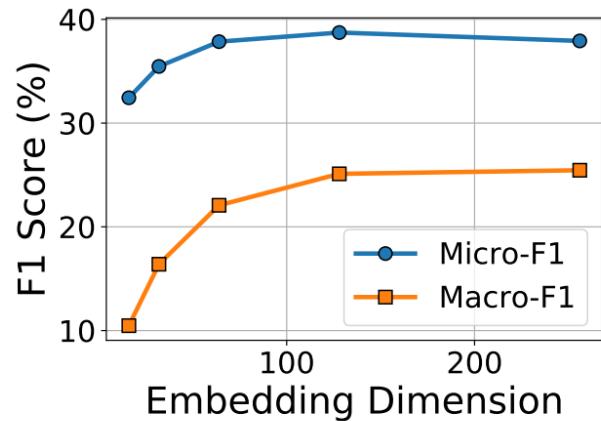
Let $\|\cdot\|_F$ be the matrix Frobenius norm. Then

$$\left\| \text{trunc_log}^\circ \left(\frac{\text{vol}(G)}{b} \tilde{\mathbf{M}} \right) - \text{trunc_log}^\circ \left(\frac{\text{vol}(G)}{b} \mathbf{M} \right) \right\|_F \leq \frac{4\epsilon \text{vol}(G)}{b \sqrt{d_{\min}}} \sqrt{\sum_{i=1}^n \frac{1}{d_i}}.$$

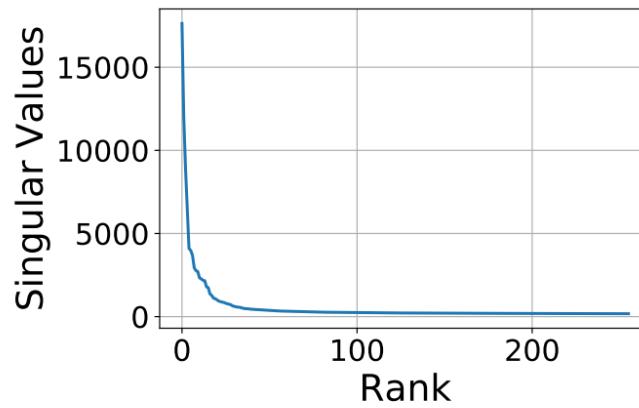
Dataset	BlogCatalog	PPI	Flickr	YouTube	OAG
$ V $	10,312	3,890	80,513	1,138,499	67,768,244
$ E $	333,983	76,584	5,899,882	2,990,443	895,368,962
#labels	39	50	195	47	19

	<i>LINE</i>	<i>DeepWalk</i>	<i>node2vec</i>	<i>NetMF</i>	<i>NetSMF</i>
BlogCatalog	40 mins	12 mins	56 mins	19 mins	13 mins
PPI	41 mins	4 mins	4 mins	1 min	10 secs
Flickr	42 mins	2.2 hours	21 hours	5 days	48 mins
YouTube	46 mins	4.3 hours	4 days	×	4.1 hours
OAG	2.6 hours	–	–	×	24 hours



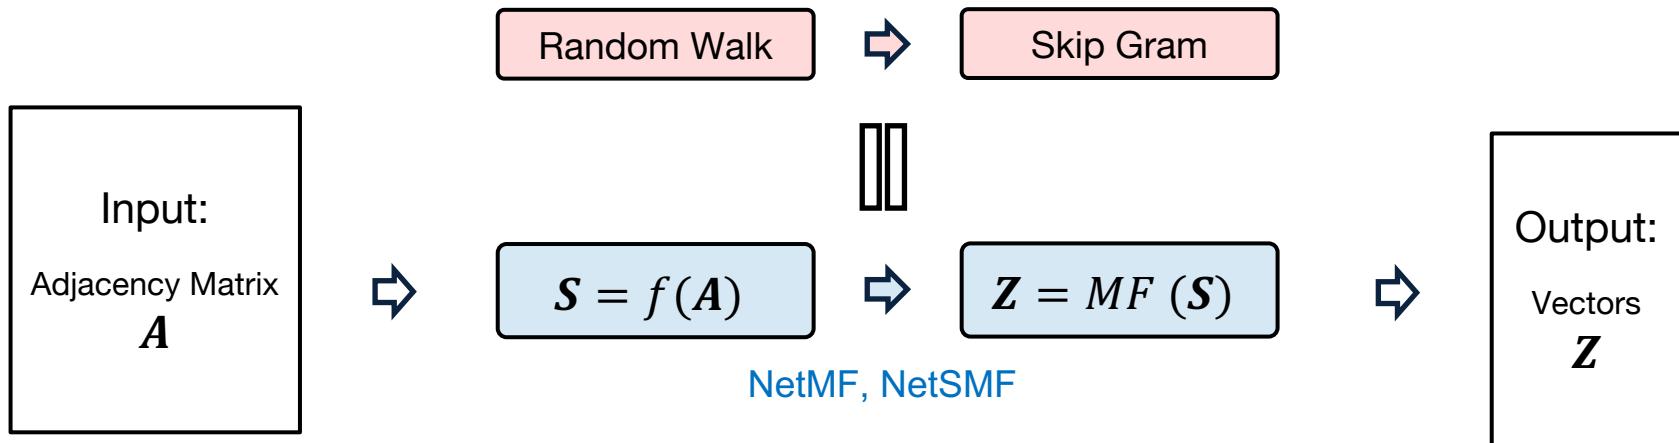


(a)



(b)

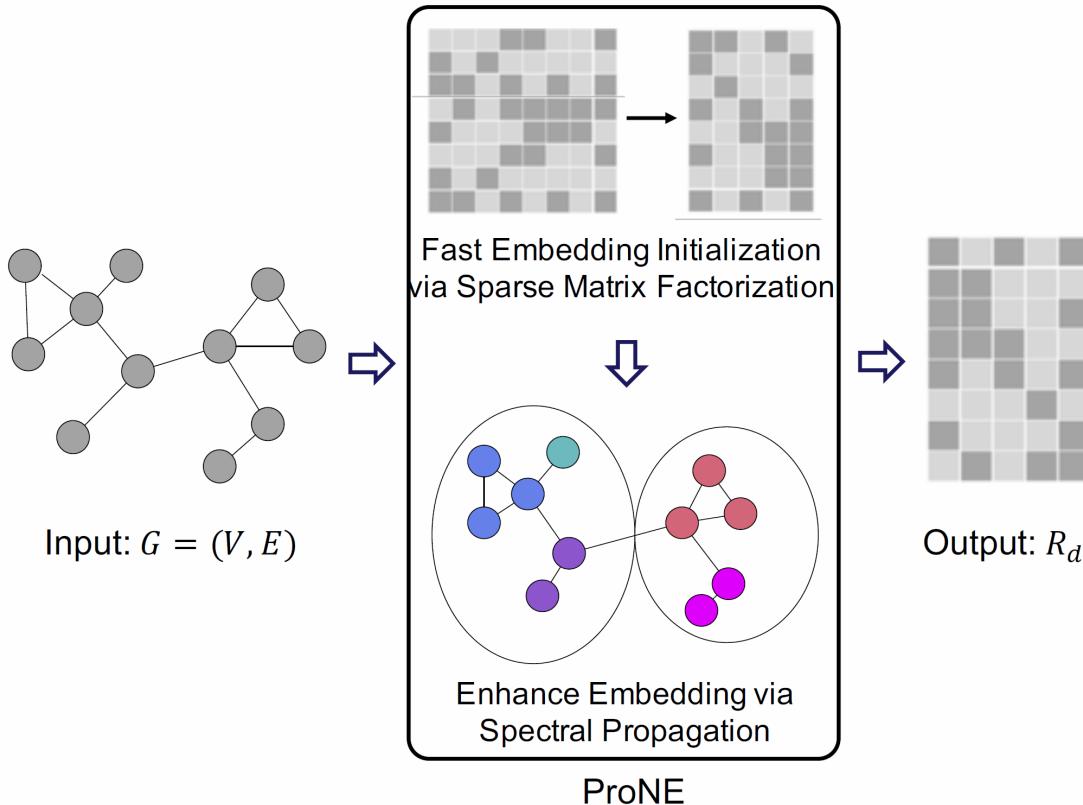
Connecting NE with graph neural networks



Incorporate network structures A into the similarity matrix S , and then factorize S

$$f(A) = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

ProNE: More fast & scalable network embedding



Embedding enhancement via spectral propagation

$$R_d \leftarrow D^{-1}A(I_n - \tilde{L}) R_d$$

$\tilde{L} = U g(\Lambda) U^T$ is the spectral filter of $L = I_n - D^{-1}A$

$D^{-1}A(I_n - \tilde{L})$ is $D^{-1}A$ modulated by the filter in the spectrum

Chebyshev expansion for efficiency

- To avoid explicit eigendecomposition and Fourier transform
 - Chebyshev expansion $T_{i+1}(x) = 2xT_i(x) - T_{i-1}(x)$ with $T_0(x) = 1, T_1(x) = x$

$$\begin{aligned} \tilde{L} &= U \text{diag}([g(\lambda_1), \dots, g(\lambda_n)]) U^T & \Rightarrow & \tilde{L} \approx B_0(\theta)T_0(\bar{L}) + 2 \sum_{i=1}^{k-1} (-)^i B_i(\theta)T_i(\bar{L}) \\ &\approx U \sum_{i=0}^{k-1} c_i(\theta)T_i(\bar{\Lambda}) U^T \\ &= \sum_{i=0}^{k-1} c_i(\theta)T_i(\bar{L}) \end{aligned}$$

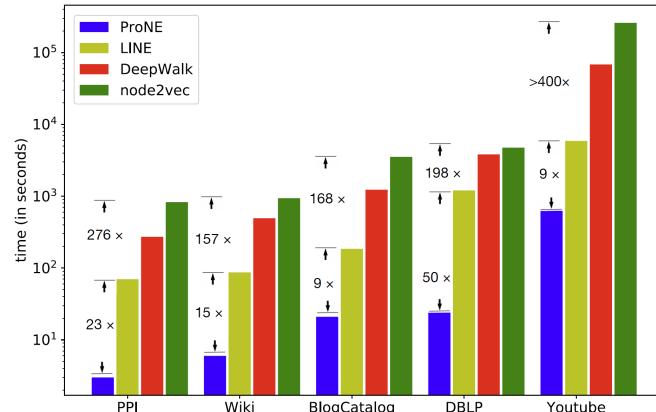
Efficiency

Dataset	DeepWalk	LINE	node2vec	ProNE
PPI	272	70	828	3
Wiki	494	87	939	6
BlogCatalog	1,231	185	3,533	21
DBLP	3,825	1,204	4,749	24
Youtube	68,272	5,890	>5days	627

20 Threads 1 Thread

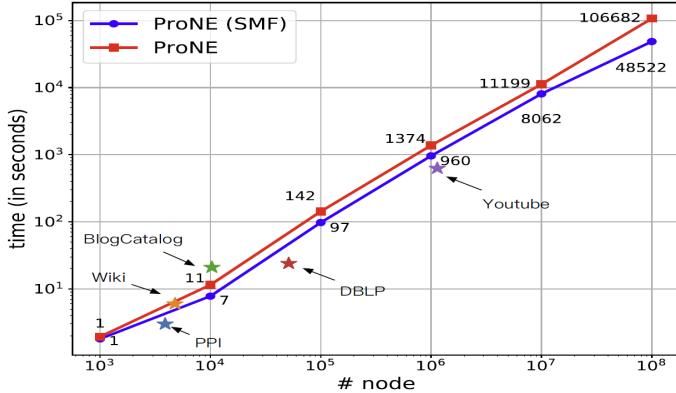
19hours 98mins 10mins

1.1M nodes



**ProNE offers 10-400X speedups
(1 thread vs 20 threads)**

Scalability & Effectiveness

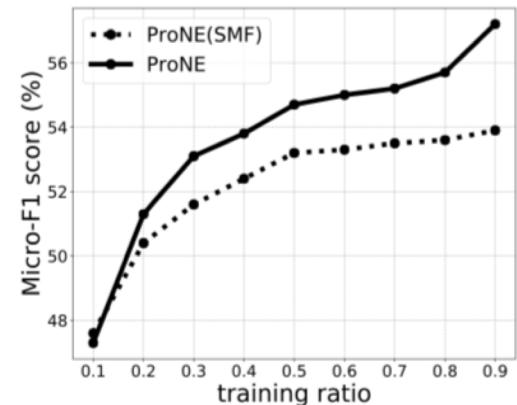
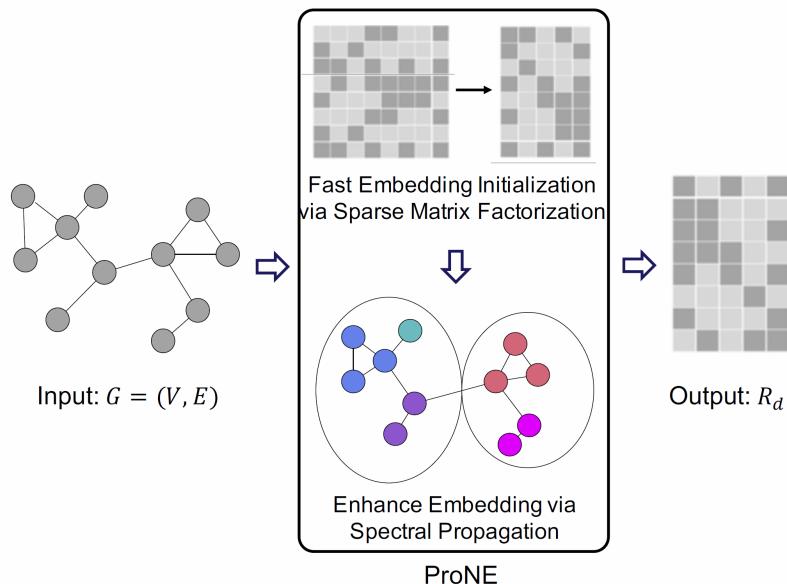


(a) The node degree is fixed to 10 and #nodes grows

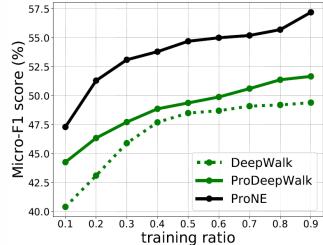
Dataset	training ratio	0.1	0.3	0.5	0.7	0.9
PPI	DeepWalk	16.4	19.4	21.1	22.3	22.7
	LINE	16.3	20.1	21.5	22.7	23.1
	node2vec	16.2	19.7	21.6	23.1	24.1
	GraRep	15.4	18.9	20.2	20.4	20.9
	HOPE	16.4	19.8	21.0	21.7	22.5
	ProNE (SMF)	15.8	20.6	22.7	23.7	24.2
Wiki	ProNE	18.2	22.7	24.6	25.4	25.9
	($\pm\sigma$)	(± 0.5)	(± 0.3)	(± 0.7)	(± 1.0)	(± 1.1)
	DeepWalk	40.4	45.9	48.5	49.1	49.4
	LINE	47.8	50.4	51.2	51.6	52.4
	node2vec	45.6	47.0	48.2	49.6	50.0
	GraRep	47.2	49.7	50.6	50.9	51.8
BlogCatalog	HOPE	38.5	39.8	40.1	40.1	40.1
	ProNE (SMF)	47.6	51.6	53.2	53.5	53.9
	ProNE	47.3	53.1	54.7	55.2	57.2
	($\pm\sigma$)	(± 0.7)	(± 0.4)	(± 0.8)	(± 0.8)	(± 1.3)
	DeepWalk	36.2	39.6	40.9	41.4	42.2
	LINE	28.2	30.6	33.2	35.5	36.8
Youtube	node2vec	36.3	39.7	41.1	42.0	42.1
	GraRep	34.0	32.5	33.3	33.7	34.1
	HOPE	30.7	33.4	34.3	35.0	35.3
	ProNE (SMF)	34.6	37.6	38.6	39.3	39.0

**Embed 100,000,000 nodes by one thread:
29 hours with performance superiority**

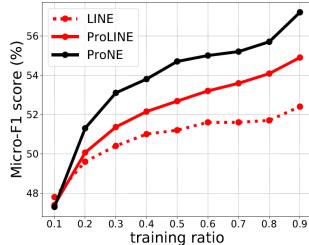
Embedding enhancement



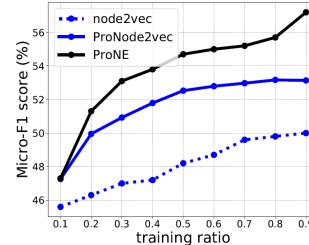
A general embedding enhancement framework



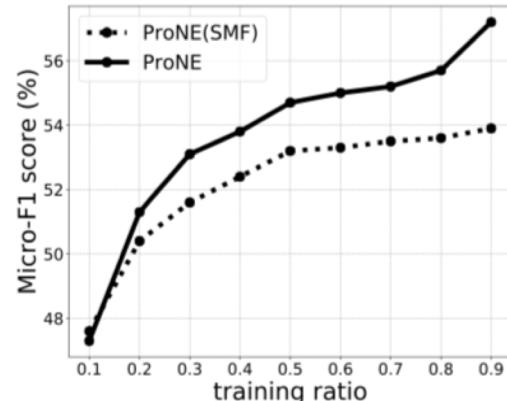
(a) ProDeepWalk



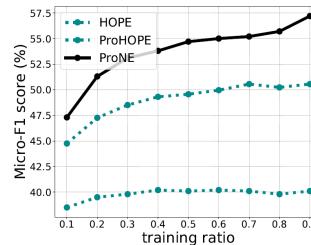
(b) ProLINE



(c) ProNode2vec

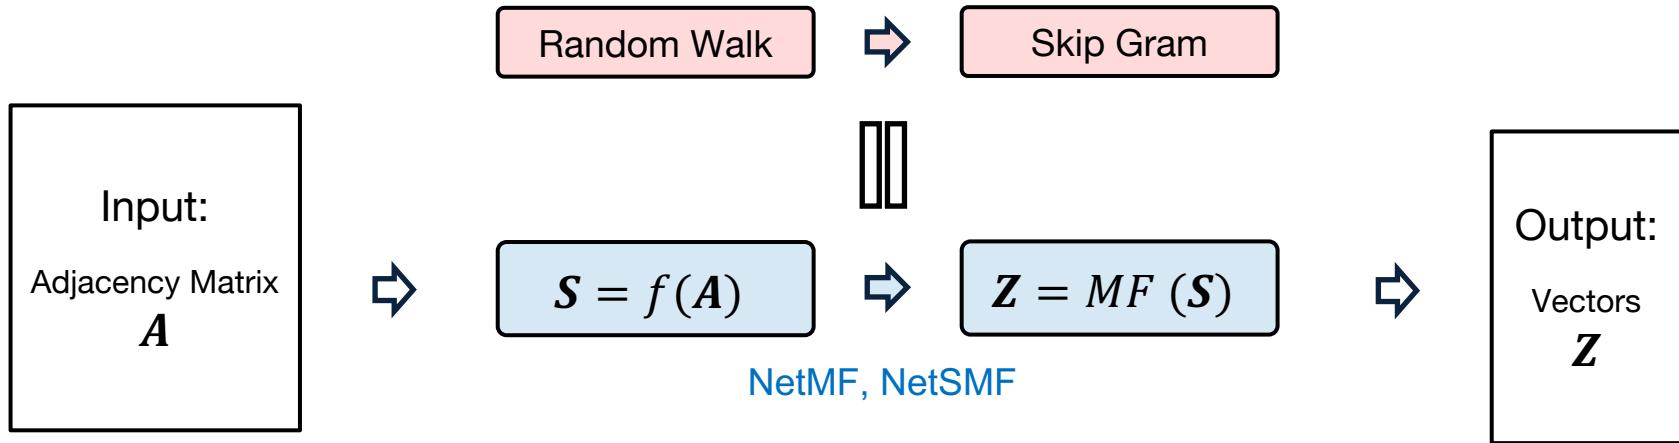


(d) ProGraRep



(e) ProHOPE

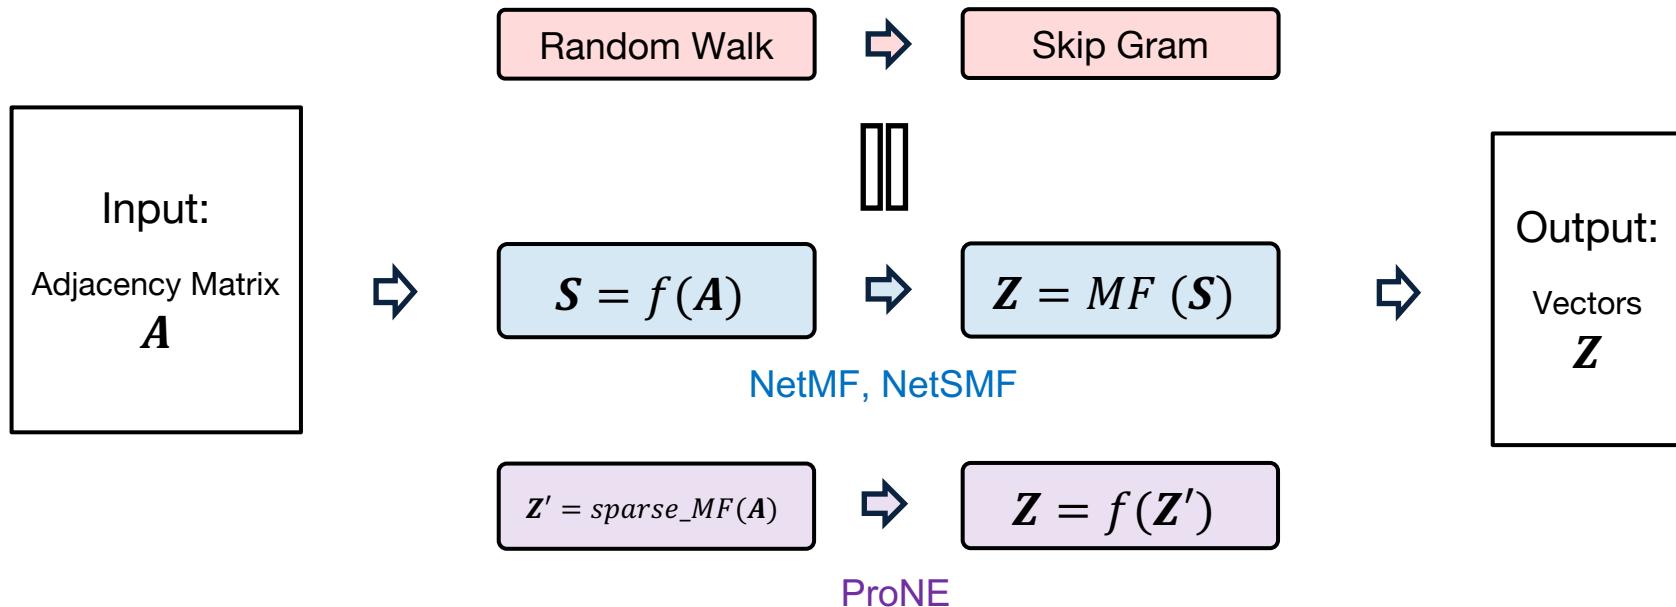
Connecting NE with graph neural networks



Incorporate network structures A into the similarity matrix S , and then factorize

$$f(A) = \log \left(\frac{\text{vol}(G)}{b} \left(\frac{1}{T} \sum_{r=1}^T (\mathbf{D}^{-1} \mathbf{A})^r \right) \mathbf{D}^{-1} \right)$$

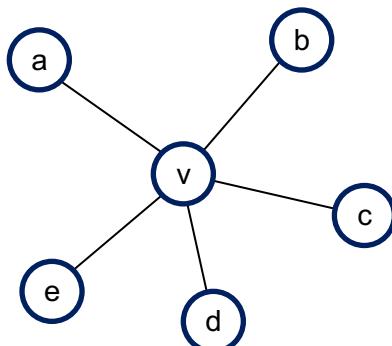
Connecting NE with graph neural networks



Factorize A , and then incorporate network structures via spectral propagation

Connecting NE with graph neural networks

$$\text{ProNE: } R_d \leftarrow D^{-1} A(I_n - \tilde{L}) R_d$$



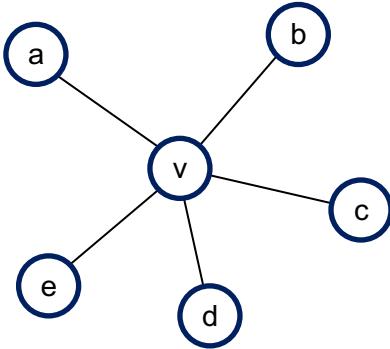
$$\mathbf{h}_v = f(\mathbf{h}_v, \mathbf{h}_a, \mathbf{h}_b, \mathbf{h}_c, \mathbf{h}_d, \mathbf{h}_e)$$

1. Defferrard et al. Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. In **NIPS 2016**
2. Zhang et al. ProNE: Fast and Scalable Network Representation Learning. In **IJCAI 2019**

Graph Neural Networks

- Input: an undirected weighted network $G = (V, E)$ with $|V| = n$ & $|E| = m$
 - Adjacency matrix $\mathbf{A} \in \mathbb{R}_+^{n \times n}$
 - $A_{i,j} = \begin{cases} a_{i,j} > 0 & (i,j) \in E \\ 0 & (i,j) \notin E \end{cases}$
 - Degree matrix $\mathbf{D} = \text{diag}(d_1, d_2, \dots, d_n)$
 - Node feature matrix $\mathbf{X} \in \mathbb{R}^{n \times q}$
- Output: for each node, its k -dimension latent feature representation vector $\mathbf{Z}^{n \times k}$
 - Latent feature embedding matrix $\mathbf{Z} \in \mathbb{R}^{n \times k}$

The Core of Graph Neural Networks

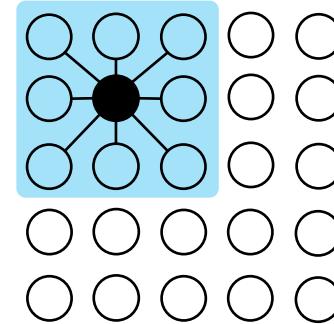
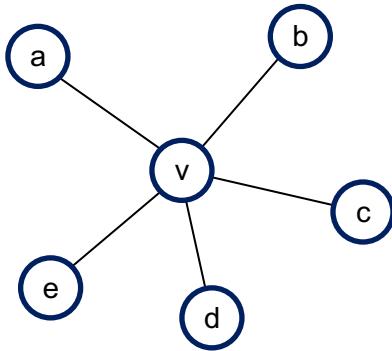


$$\mathbf{h}_v = f(\mathbf{h}_a, \mathbf{h}_b, \mathbf{h}_c, \mathbf{h}_d, \mathbf{h}_e)$$

Neighborhood Aggregation:

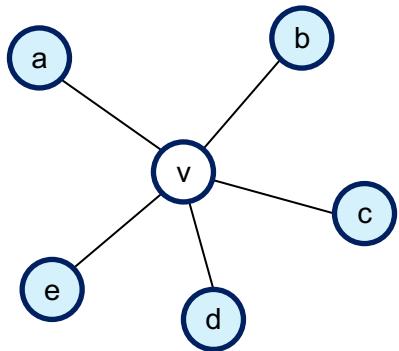
Aggregate neighbor information and pass into a neural network

Graph neural networks



- **Neighborhood Aggregation:**
 - Aggregate neighbor information and pass into a neural network
 - It can be viewed as a center-surround filter in CNN--graph convolutions!

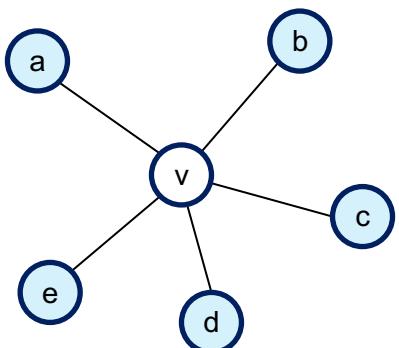
GNN: Graph convolutional networks



$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v) \cup v} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

1. Kipf et al. Semisupervised Classification with Graph Convolutional Networks. ICLR 2017
2. Defferrard et al. Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering. In NIPS 2016

GNN: Graph convolutional networks



$$h_v^k = \sigma(W^k \sum_{u \in N(v) \cup v} \frac{h_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

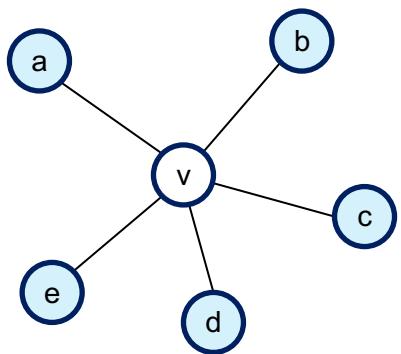
node v 's embedding at layer k

parameters in layer k

Non-linear activation function (e.g., ReLU)

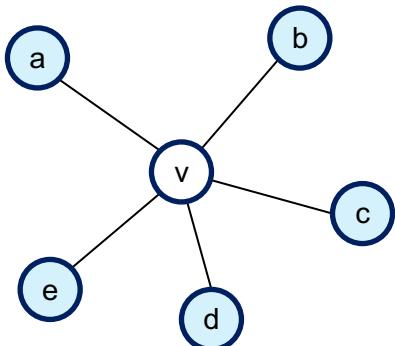
the neighbors of node v

GNN: Graph convolutional networks



$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v) \cup v} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

GNN: Graph convolutional networks

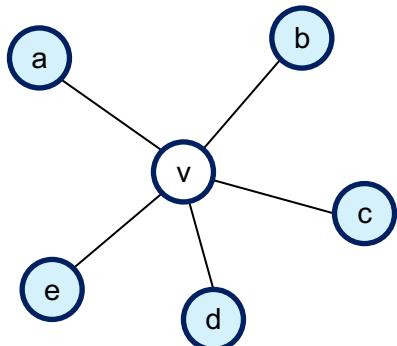


$$h_v^k = \sigma(W^k \sum_{u \in N(v)} \frac{h_u^{k-1}}{\sqrt{|N(u)||N(v)|}} + W^k \sum_v \frac{h_v^{k-1}}{\sqrt{|N(v)||N(v)|}})$$

Aggregate from v 's neighbors

Aggregate from itself

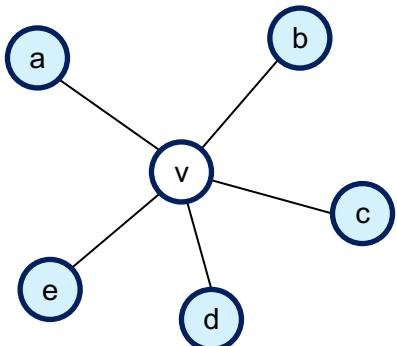
GNN: Graph convolutional networks



The same parameters for both its neighbors & itself

$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}} + \mathbf{W}^k \sum_v \frac{\mathbf{h}_v^{k-1}}{\sqrt{|N(v)||N(v)|}})$$

GNN: Graph convolutional networks

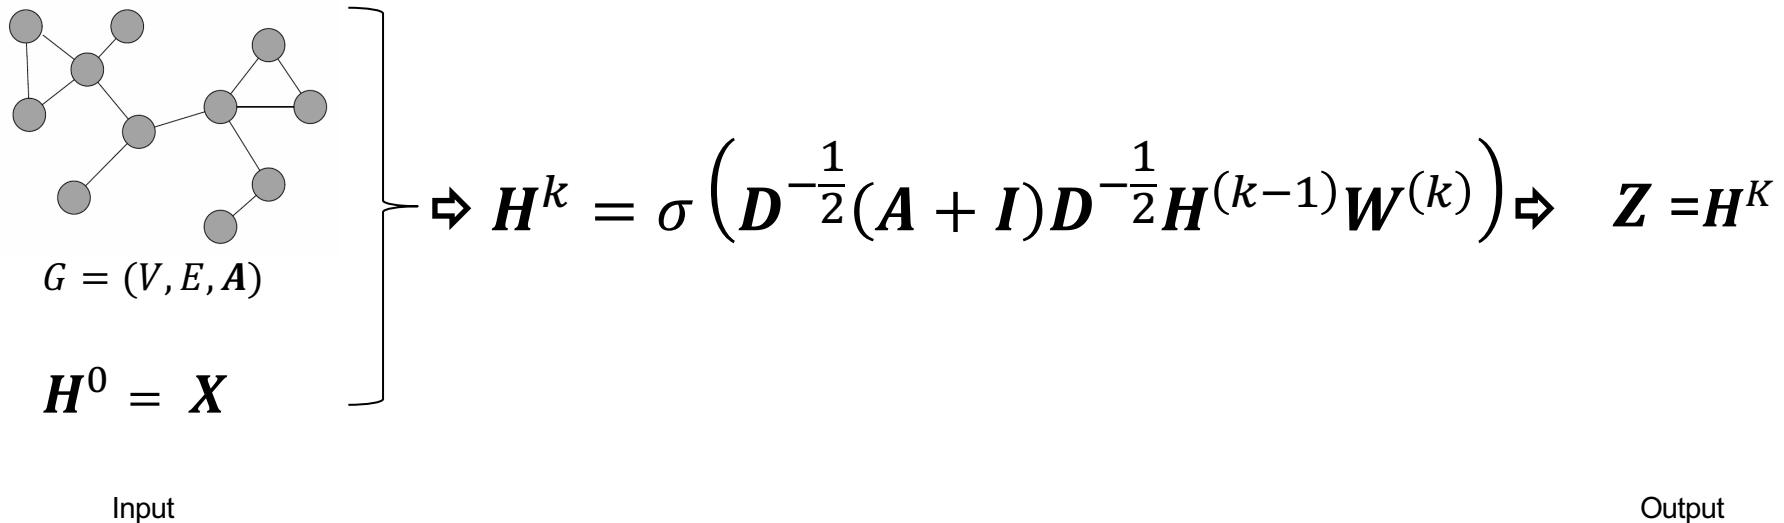


$$\mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}} \mathbf{H}^{(k-1)} \mathbf{W}^{(k)}$$

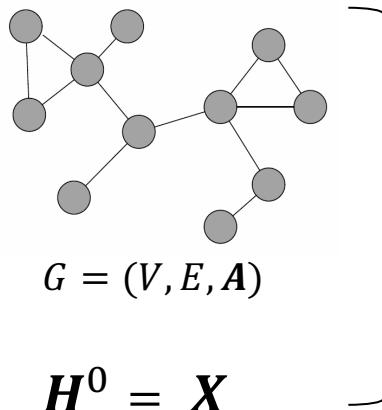
$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}} + \mathbf{W}^k \sum_v \frac{\mathbf{h}_v^{k-1}}{\sqrt{|N(v)||N(v)|}})$$

$$\mathbf{D}^{-\frac{1}{2}} \mathbf{I} \mathbf{D}^{-\frac{1}{2}} \mathbf{H}^{(k-1)} \mathbf{W}^{(k)}$$

GNN: Graph convolutional networks



GNN: Graph convolutional networks



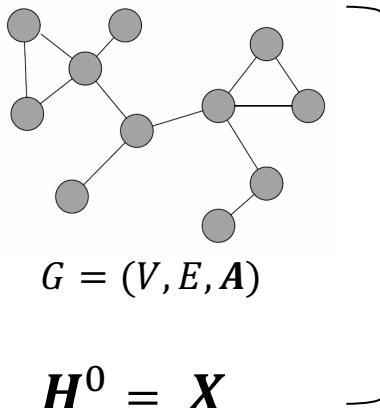
Input

- Model training

- The common setting is to have an end to end training framework with a supervised task
- That is, define a loss function over \mathbf{Z}

Output

GNN: Graph convolutional networks



$$\Rightarrow H^k = \sigma \left(D^{-\frac{1}{2}}(A + I)D^{-\frac{1}{2}}H^{(k-1)}W^{(k)} \right) \Rightarrow Z = H^K$$

Input

- Benefits: Parameter sharing for all nodes
 - #parameters is sublinear in $|V|$
 - Enable inductive learning for new nodes

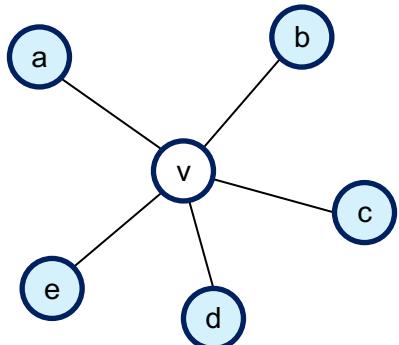
Output

GNN: Graph convolutional networks

$$\mathbf{H}^k = \sigma \left(\mathbf{D}^{-\frac{1}{2}} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-\frac{1}{2}} \mathbf{H}^{(k-1)} \mathbf{W}^{(k)} \right)$$

- GCN is one way of neighbor aggregations
- **GraphSage**
- Graph Attention
-

GraphSage



GCN

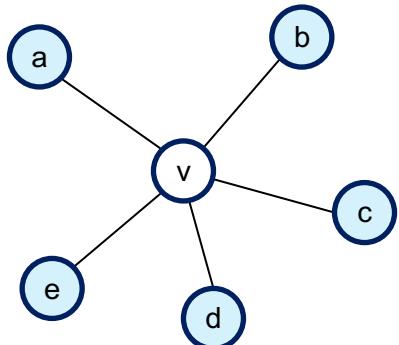
$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v) \cup v} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

GraphSage

Instead of summation, it concatenates
neighbor & self embeddings

$$\mathbf{h}_v^k = \sigma([\mathbf{A}^k \cdot \text{AGG}(\{\mathbf{h}_u^{k-1}, \forall u \in N(v)\}), \mathbf{B}^k \mathbf{h}_v^{k-1}])$$

GraphSage



GCN

$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v) \cup v} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

GraphSage

$$\mathbf{h}_v^k = \sigma([\mathbf{A}^k \cdot \text{AGG}(\{\mathbf{h}_u^{k-1}, \forall u \in N(v)\}), \mathbf{B}^k \mathbf{h}_v^{k-1}])$$

Generalized aggregation: any differentiable function that maps set of vectors to a single vector

GraphSage

$$\mathbf{h}_v^k = \sigma([\mathbf{A}^k \cdot \text{AGG}(\{\mathbf{h}_u^{k-1}, \forall u \in N(v)\}), \mathbf{B}^k \mathbf{h}_v^{k-1}])$$

- Mean:

$$\text{AGG} = \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|}$$

- Pool

- Transform neighbor vectors and apply symmetric vector function.

$$\text{AGG} = \gamma(\{\mathbf{Q}\mathbf{h}_u^{k-1}, \forall u \in N(v)\})$$

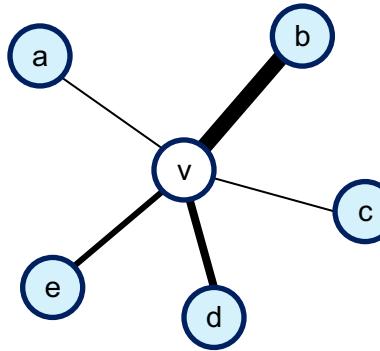
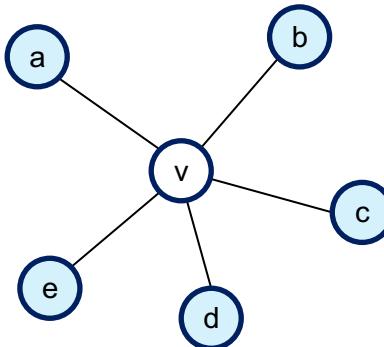
element-wise mean/max

- LSTM:

- Apply LSTM to random permutation of neighbors.

$$\text{AGG} = \text{LSTM}([\mathbf{h}_u^{k-1}, \forall u \in \pi(N(v))])$$

Graph Neural Networks

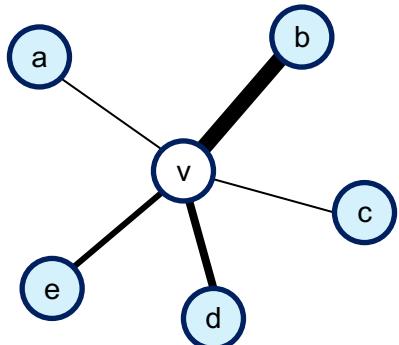


Graph Neural Networks

$$\mathbf{H}^k = \sigma \left(\mathbf{D}^{-\frac{1}{2}} (\mathbf{A} + \mathbf{I}) \mathbf{D}^{-\frac{1}{2}} \mathbf{H}^{(k-1)} \mathbf{W}^{(k)} \right)$$

- GCN is one way of neighbor aggregations
- GraphSage
- **Graph Attention**
-

GNN: Graph Attention



GCN

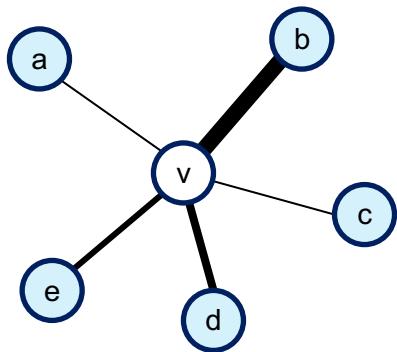
$$\mathbf{h}_v^k = \sigma(\mathbf{W}^k \sum_{u \in N(v) \cup v} \frac{\mathbf{h}_u^{k-1}}{\sqrt{|N(u)||N(v)|}})$$

Graph Attention

$$\mathbf{h}_v^k = \sigma(\sum_{u \in N(v) \cup v} \alpha_{v,u} \mathbf{W}^k \mathbf{h}_u^{k-1})$$

Learned attention weights

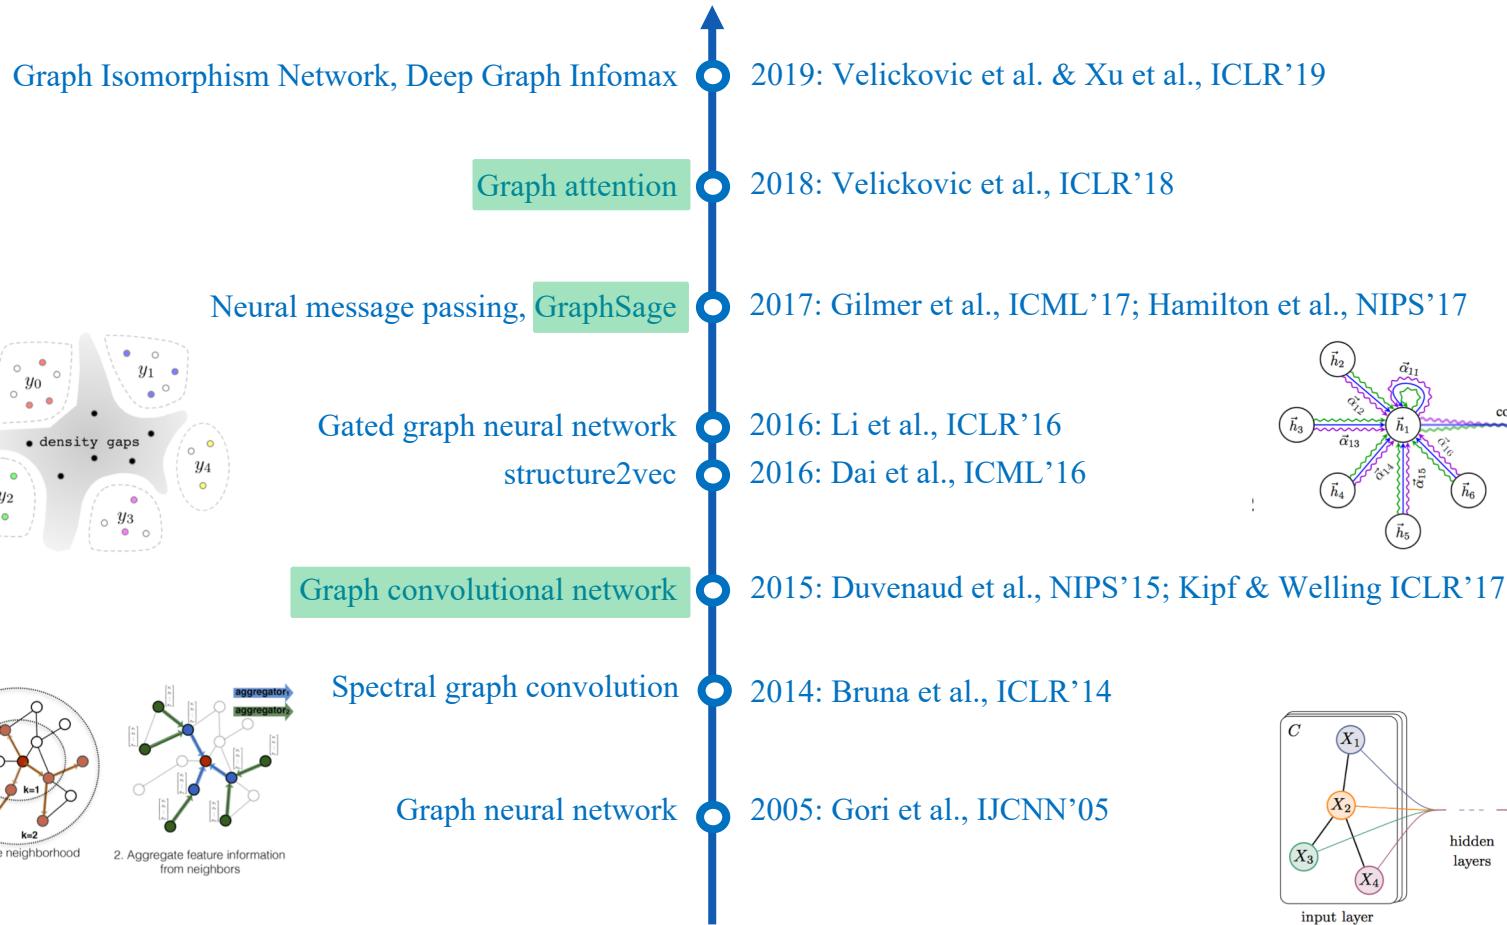
GNN: Graph Attention



$$\alpha_{v,u} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{Q}\mathbf{h}_v, \mathbf{Q}\mathbf{h}_u]))}{\sum_{u' \in N(v) \cup \{v\}} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\mathbf{Q}\mathbf{h}_v, \mathbf{Q}\mathbf{h}_{u'}]))}$$

Various ways to define attention!

Graph neural networks



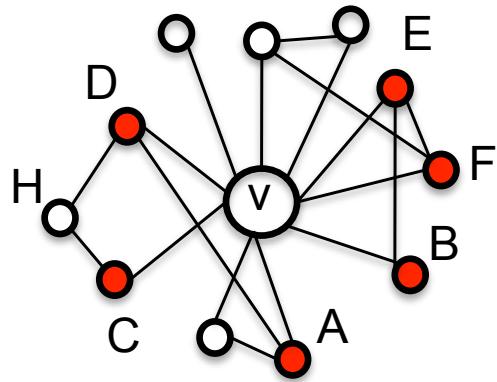
GNN applications & systems

- Deeplnf: Modeling social influence with graph neural networks
- LinKG: Knowledge graph linking with heterogeneous graph attention
- AliGraph: A comprehensive graph neural network platform.
 - **Dr. Hongxia Yang**
 - Applied data science invited talk
 - 10AM--12PM, Thursday, Aug 6th
 - Cook Room, Street Level, Egan Center



1. Qiu et al. Deeplnf: Social Influence Prediction with Deep Learning. In KDD'18.
2. Zhang, et al. OAG: Toward Linking Large-scale Heterogeneous Entity Graphs. In KDD'19.

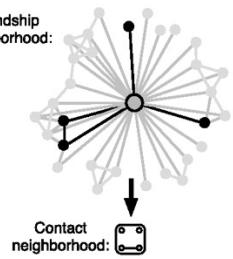
DeepInf: Modeling social influence with graph neural networks



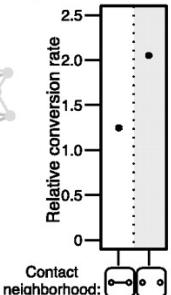
Given the five friends in red did something, whether the central users will do the same thing later, such as retweeting in Twitter, “like” in FB, or product purchase?

Previous Solution

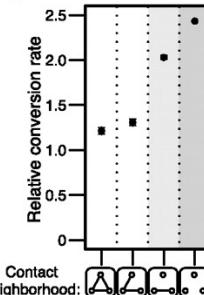
A



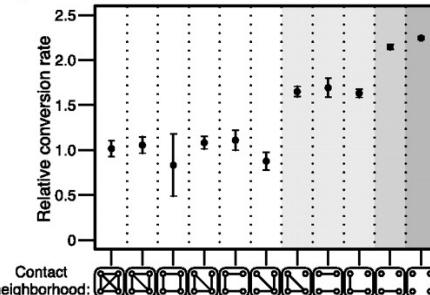
B



C

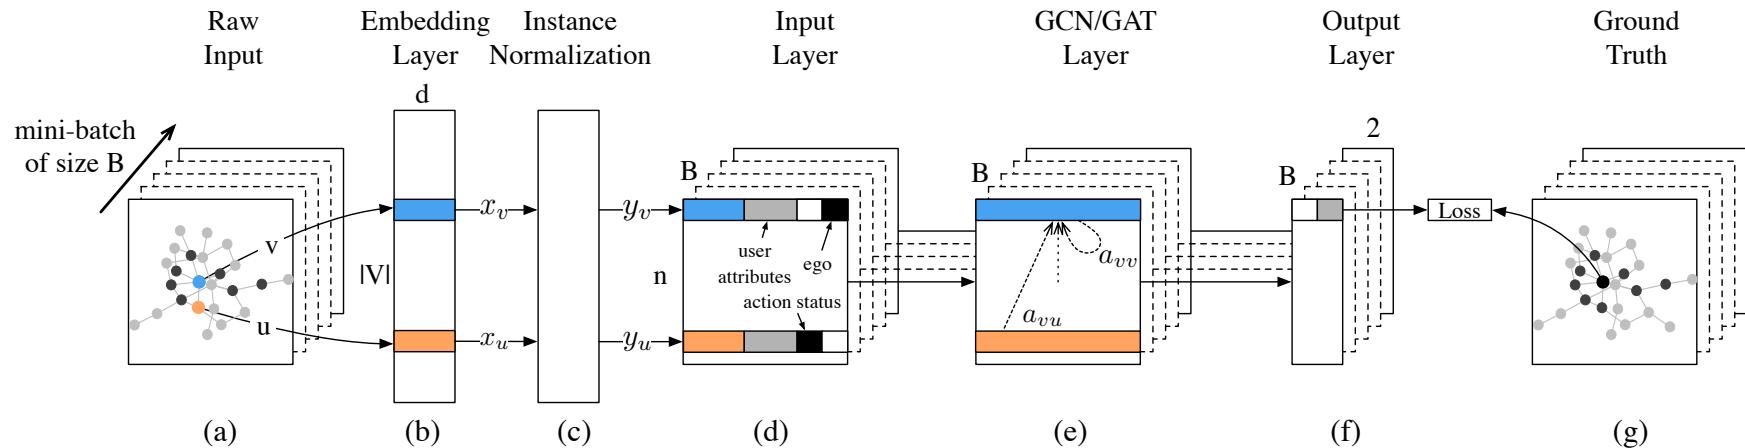


D



Name	Description
Vertex	Coreness [3]. Pagerank [30]. Hub score and authority score [8]. Eigenvector Centrality [5]. Clustering Coefficient [46]. Rarity (reciprocal of ego user's degree) [1]. Network embedding (DeepWalk [31], 64-dim).
Ego	The number/ratio of active neighbors [2]. Density of subnetwork induced by active neighbors [40]. #Connected components formed by active neighbors [40].

Graph Attention Networks

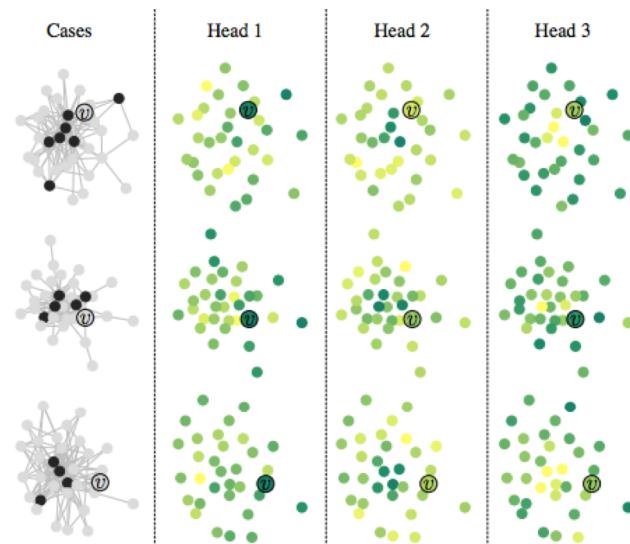


Experiments --- Results

Data	Model	AUC	Prec.	Rec.	F1
OAG	LR	65.55	32.26	69.97	44.16
	SVM	65.48	32.17	69.82	44.04
	PSCN	67.70	36.24	60.46	45.32
	DeepInf-GAT	70.59	38.93	61.29	47.61
Digg	LR	84.72	56.78	73.12	63.92
	SVM	86.01	63.42	67.34	65.32
	PSCN	83.96	62.16	67.34	64.65
	DeepInf-GAT	88.97	68.80	73.79	71.21
Twitter	LR	78.07	45.86	69.81	55.36
	SVM	79.42	49.12	67.31	56.79
	PSCN	79.40	48.43	68.06	56.59
	DeepInf-GAT	80.01	49.39	67.47	57.03
Weibo	LR	77.10	42.34	72.88	53.56
	SVM	77.11	43.27	70.79	53.71
	PSCN	79.54	44.89	73.48	55.73
	DeepInf-GAT	82.75	48.86	74.13	58.90

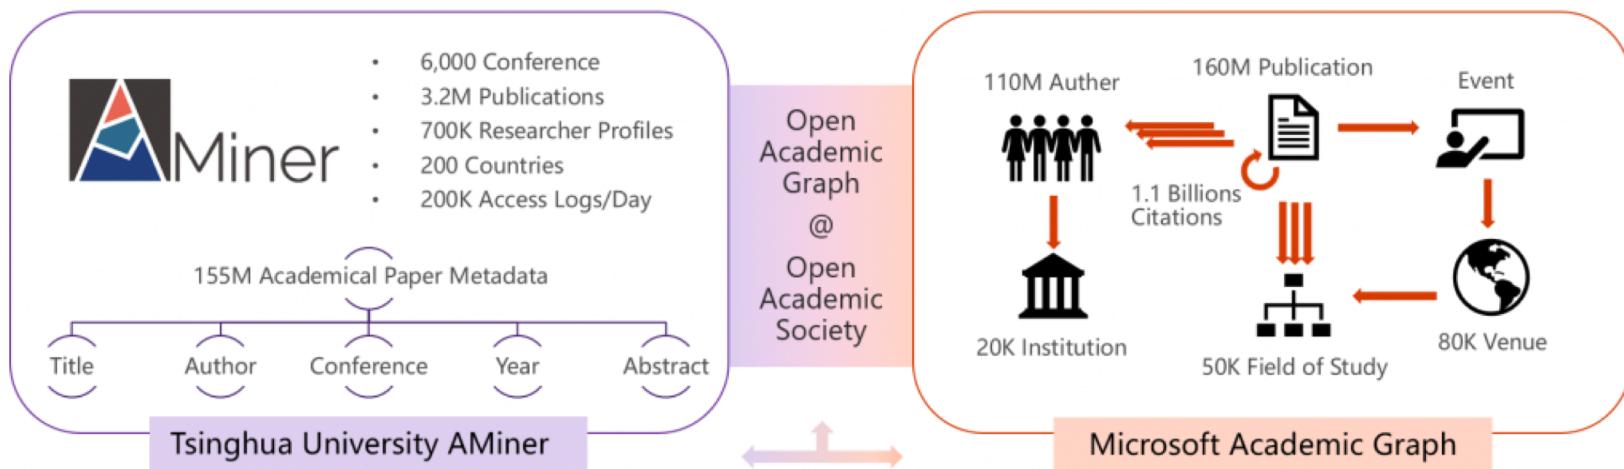
Case Study

- How different graph attention heads highlight different areas of the network.
 - Head 1: Focus on the ego-user
 - Head 2: Highlight active users
 - Head 3: Highlight inactive users

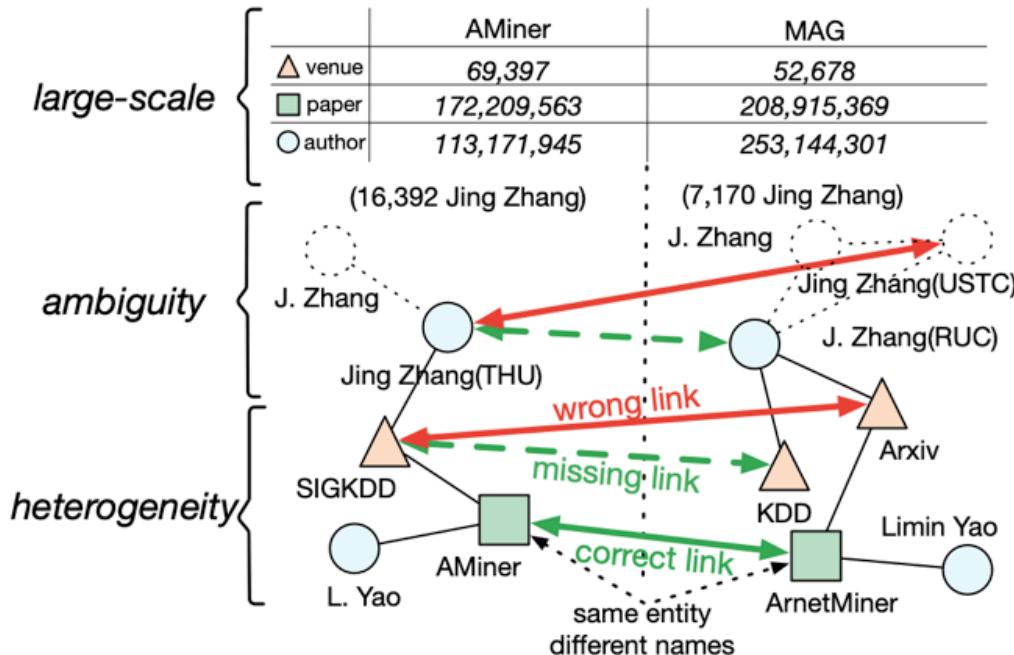


LinKG: Knowledge graph linking with heterogeneous graph attention

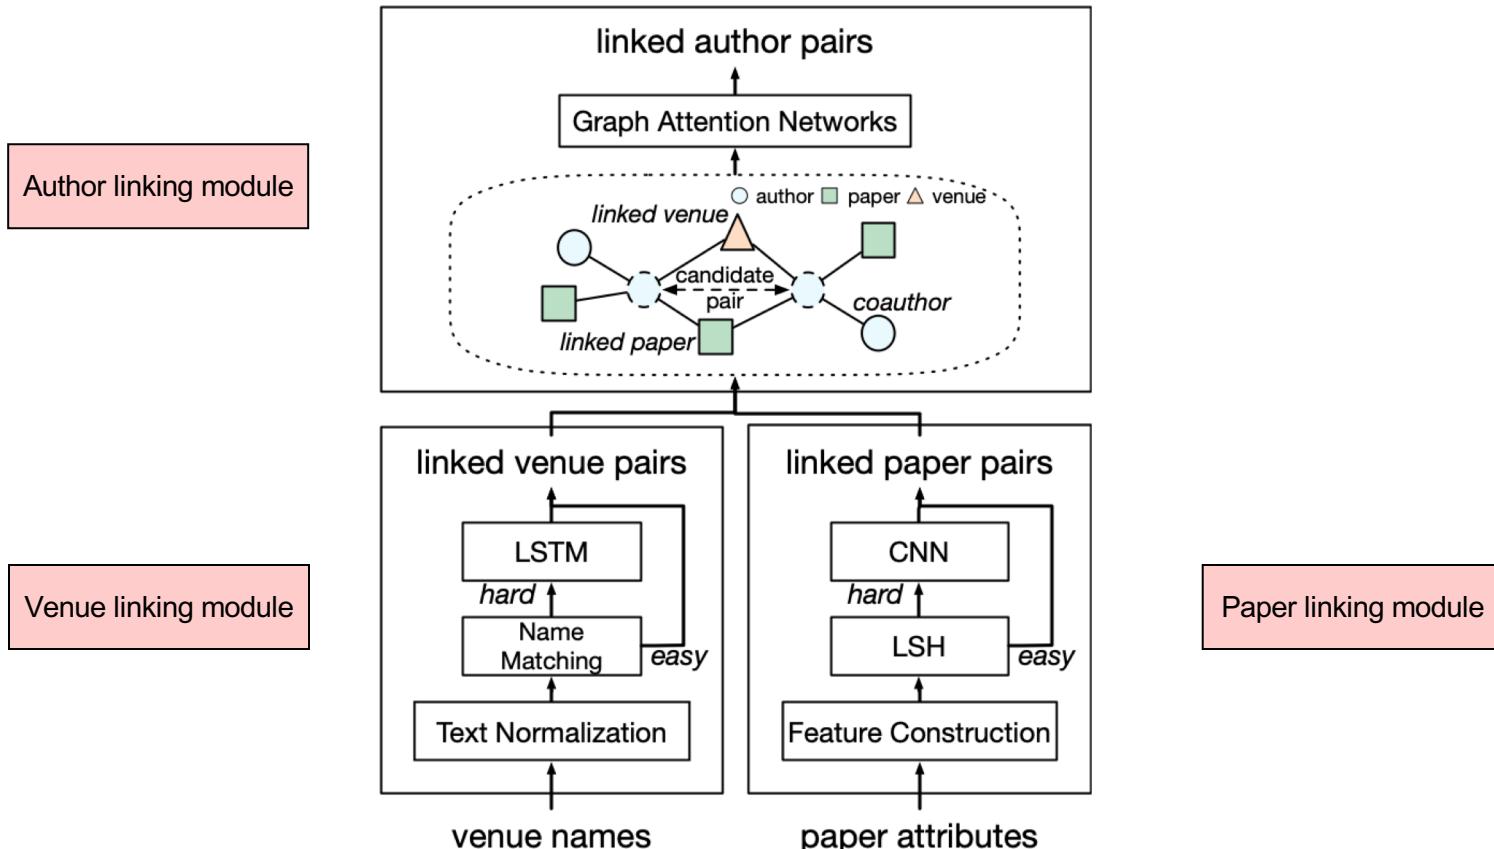
- **Input:** two heterogeneous entity graphs HG_1 and HG_2 .
- **Output:** entity linkings $L = \{(e_1, e_2) | e_1 \in HG_1, e_2 \in HG_2\}$ such that e_1 and e_2 represent exactly the same entity.



Linking large-scale heterogeneous academic graphs



Solution -- LinKG



Author linking model — Heterogenous Graph Attention

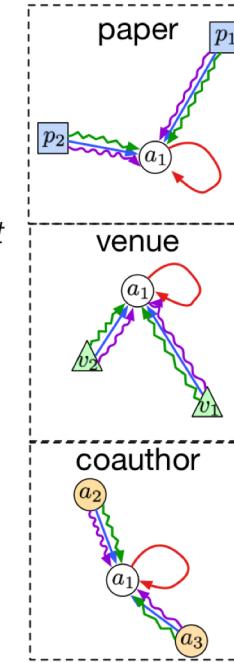
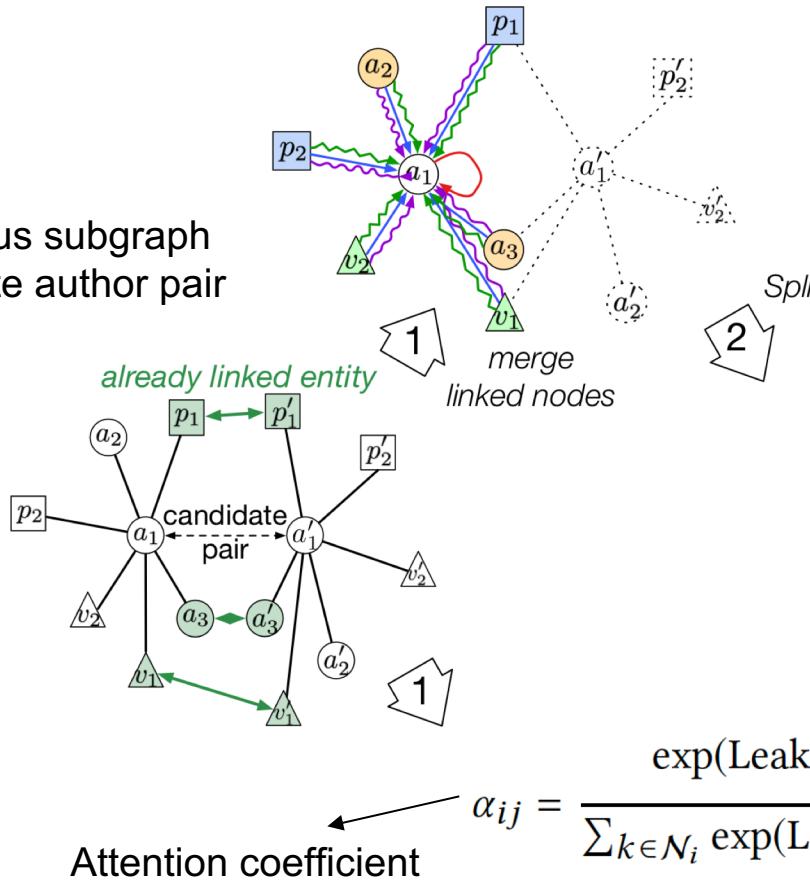
- Encoder layers
 - attention coefficient $\text{attn}(e_i, e_j)$ learnt by self-attention mechanism
$$o_{ij} = \text{attn}(Wh_i, Wh_j)$$
 - Normalized attention coefficient: differentiate different types of entities

$$\alpha_{ij} = \frac{\exp(\text{LeakyReLU}(c_{\tau(e_i)}^\top Wh_i + c_{\tau(e_j)}^\top Wh_j))}{\sum_{k \in \mathcal{N}_i} \exp(\text{LeakyReLU}(c_{\tau(e_i)}^\top Wh_i + c_{\tau(e_k)}^\top Wh_k))}$$

aggregation weight of source entity e_j 's
embedding on target entity e_i

Author linking model — Heterogenous Graph Attention

Heterogeneous subgraph
for a candidate author pair



Different attention
parameters for
different entity types

Experimental Results

Table 1: Results of linking heterogeneous entity graphs. “–” indicates the method does not support the entity linking.

Methods		Keyword	SVM	Dedupe	COSNET	MEgo2Vec	LinKG _C	LinKG _L	LinKG
Venue	Prec.	80.15	81.69	84.25			84.67	91.16	91.16
	Rec.	83.76	83.45	80.92	-	-	85.81	87.58	87.58
	F1	81.91	82.56	82.55			85.23	89.33	89.33
Paper	Prec.	91.01	96.93	99.30			98.68	86.72	98.68
	Rec.	80.53	96.78	87.09	-	-	98.10	86.59	98.10
	F1	85.45	96.86	92.80			98.39	86.66	98.39
Author	Prec.	44.48	84.70	50.65	91.73	91.03	81.30	84.92	95.37
	Rec.	80.63	92.22	85.46	85.33	90.82	84.95	94.75	93.48
	F1	57.33	88.30	63.60	88.42	90.92	83.09	89.57	94.42
Overall	Prec.	74.80	92.36	82.26	91.73	91.03	92.38	86.21	97.36
	Rec.	80.64	94.89	86.38	85.33	90.82	93.29	89.41	96.26
	F1	77.61	93.61	84.27	88.42	90.92	92.83	87.78	96.81

OAG: Open Academic Graph

<https://www.openacademic.ai/oag/>

Data set	#Pairs/Venues	Date
Linking relations	29,841	2018.12
<u>AMiner</u> venues	69,397	2018.07
MAG venues	52,678	2018.11

Table 1: statistics of OAG venue data

Data set	#Pairs/Papers	Date
Linking relations	91,137,597	2018.12
<u>AMiner</u> papers	172,209,563	2019.01
MAG papers	208,915,369	2018.11

Table 2: statistics of OAG paper data

Data set	#Pairs/Authors	Date
Linking relations	1,717,680	2019.01
<u>AMiner</u> authors	113,171,945	2018.07
MAG authors	253,144,301	2018.11

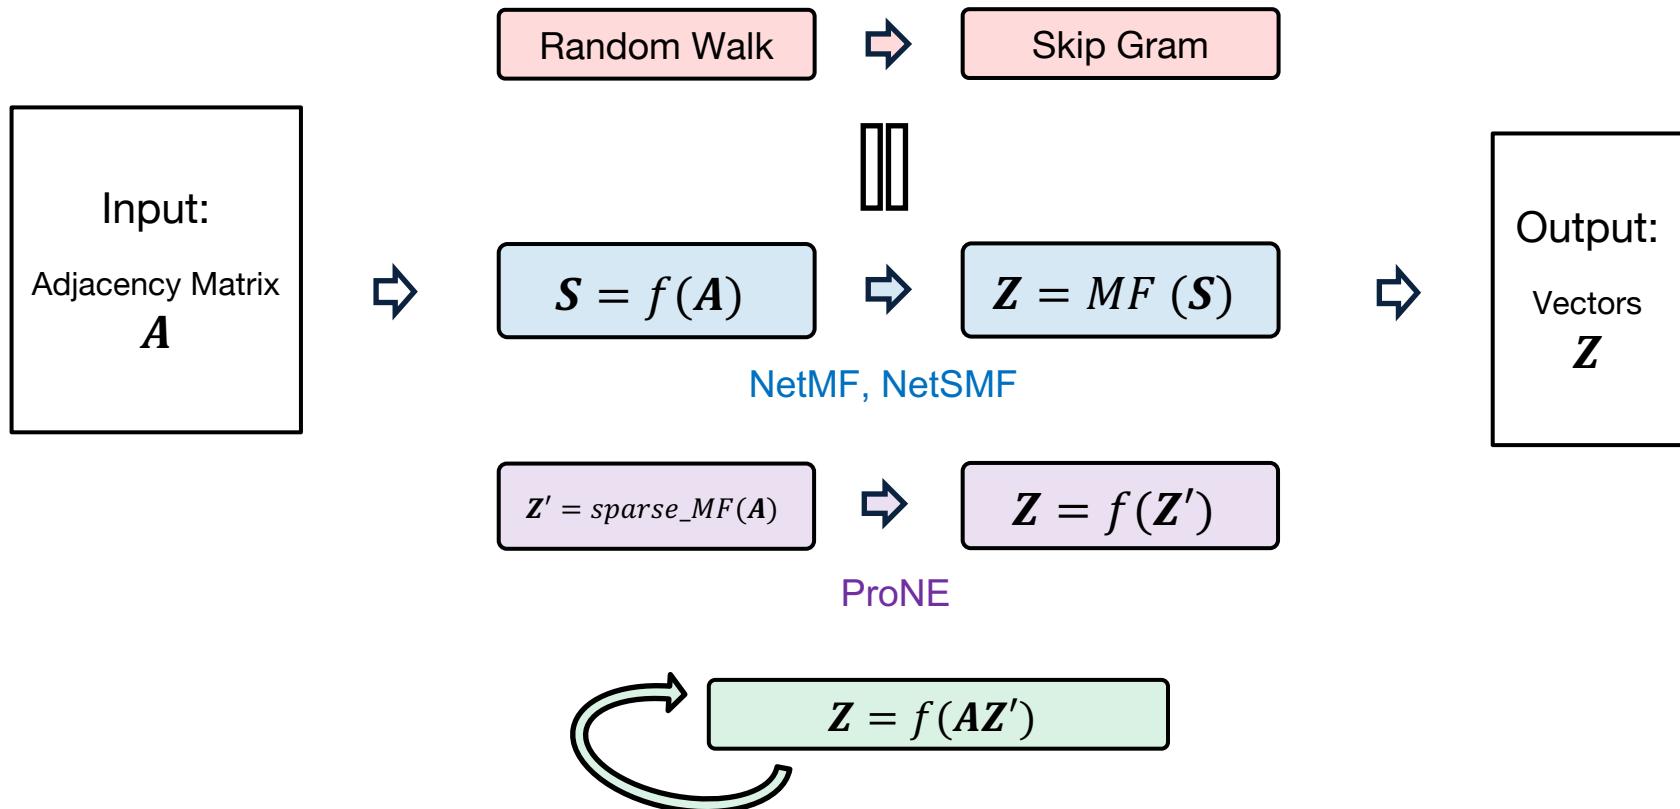
Open Academic Graph

Open Academic Graph (OAG) is a large knowledge graph unifying two billion-scale academic graphs: [Microsoft Academic Graph](#) (MAG) and [AMiner](#). In mid 2017, we published OAG v1, which contains 166,192,182 papers from MAG and 154,771,162 papers from AMiner (see below) and generated 64,639,608 linking (matching) relations between the two graphs. This time, in OAG v2, author, venue and newer publication data and the corresponding matchings are available.

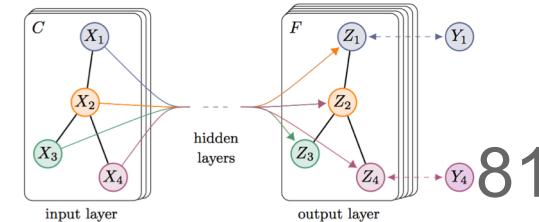
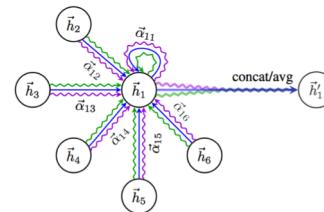
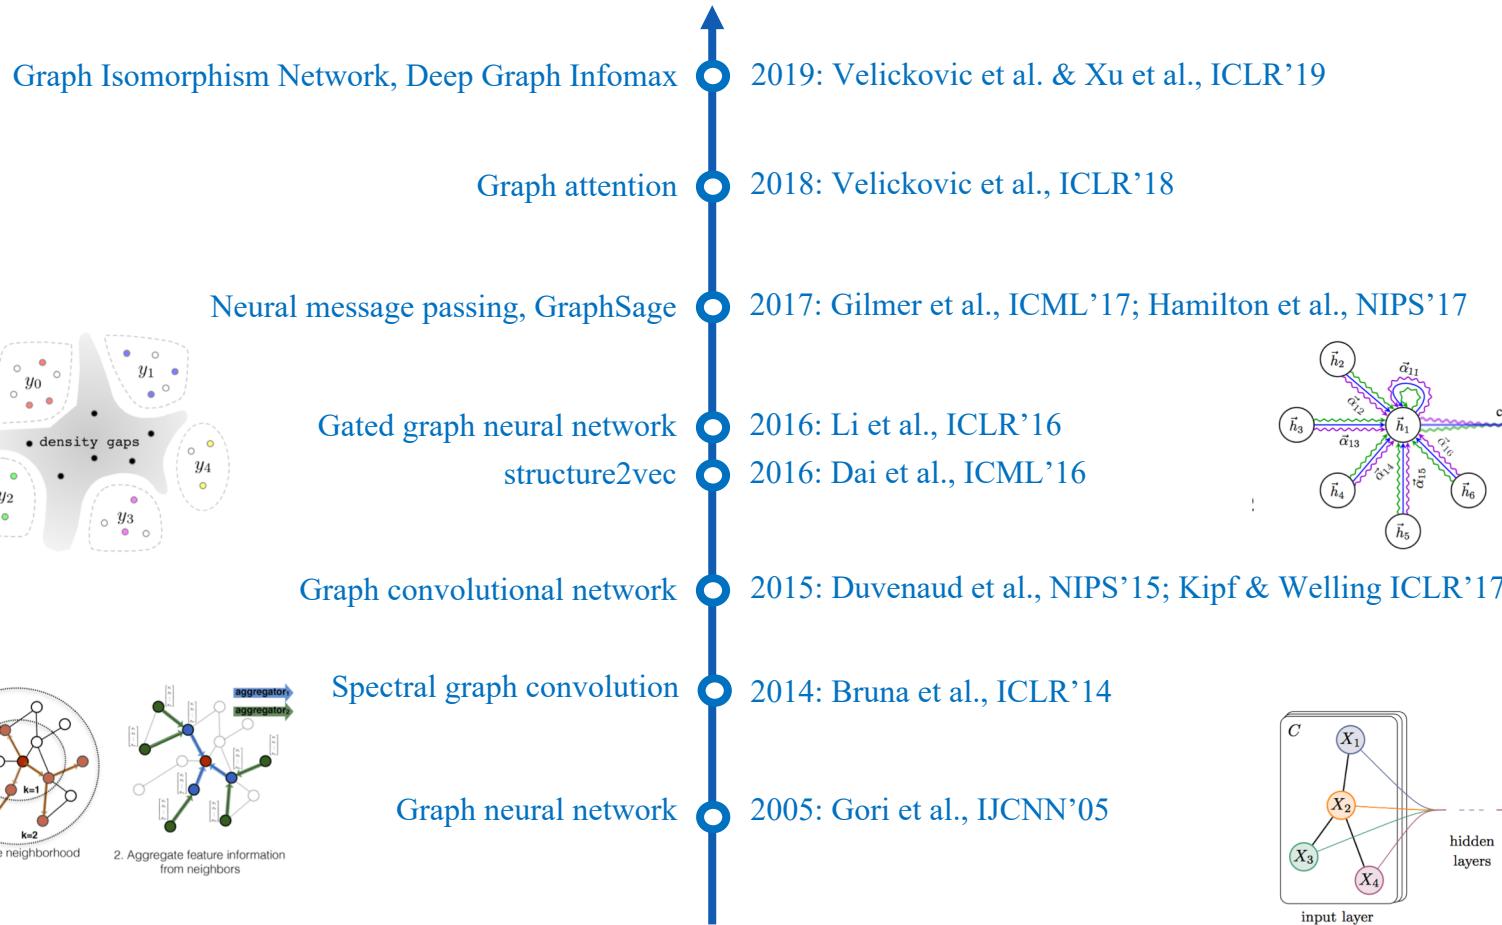
Overview of OAG v2

The statistics of OAG v2 is listed as the three tables below. The two large graphs are both evolving and we take MAG November 2018 snapshot and AMiner July 2018 or January 2019 snapshot for this version.

Connecting NE with graph neural networks



Graph neural networks

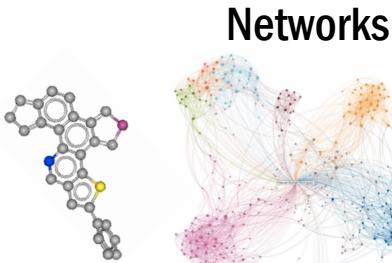


References

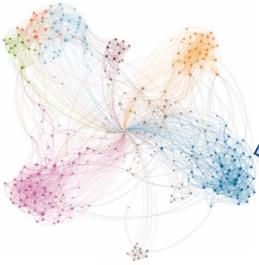
1. Jiezhong Qiu, Yuxiao Dong, Hao Ma, Jian Li, Chi Wang, Kuansan Wang, and Jie Tang. NetSMF: Large-Scale Network Embedding as Sparse Matrix Factorization. *WWW'19*.
2. Jiezhong Qiu, Yuxiao Dong, Hao Ma, Jian Li, Kuansan Wang, and Jie Tang. Network Embedding as Matrix Factorization: Unifying DeepWalk, LINE, PTE, and node2vec. *WSDM'18*.
3. Jie Zhang, Yuxiao Dong, Yan Wang, Jie Tang, and Ming Ding. ProNE: Fast and Scalable Network Representation Learning. *IJCAI'19*.
4. Jiezhong Qiu, Jian Tang, Hao Ma, Yuxiao Dong, Kuansan Wang, and Jie Tang. DeepInfl: Modeling Influence Locality in Large Social Networks. *KDD'18*.
5. Zhang, et al. OAG: Toward Linking Large-scale Heterogeneous Entity Graphs. In *KDD'19*.
6. Hamilton et al. Inductive Representation Learning on Large Graphs. *NIPS 2017*
7. Kipf et al. Semisupervised Classification with Graph Convolutional Networks. *ICLR 2017*
8. Velickovic et al. Graph Attention Networks. *ICLR 2018*
9. Perozzi et al. DeepWalk: Online learning of social representations. In *KDD'14*.
10. Tang et al. LINE: Large scale information network embedding. In *WWW'15*.
11. Grover and Leskovec. node2vec: Scalable feature learning for networks. In *KDD'16*.
12. Dong et al. metapath2vec: scalable representation learning for heterogeneous networks. In *KDD 2017*.

Graph Neural Networks: A Learning Perspective

Learning from graph data



Networks



Timing



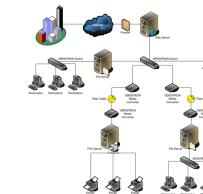
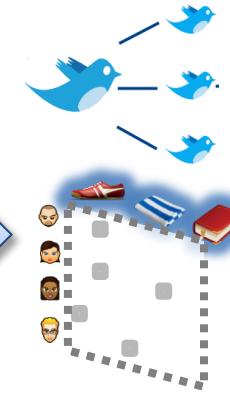
Feature/algorithm design

(city=Atlanta) AND (age=40)
(neighbor=A) XOR (neighbor=B)
(bought=car) OR (time<3 years)

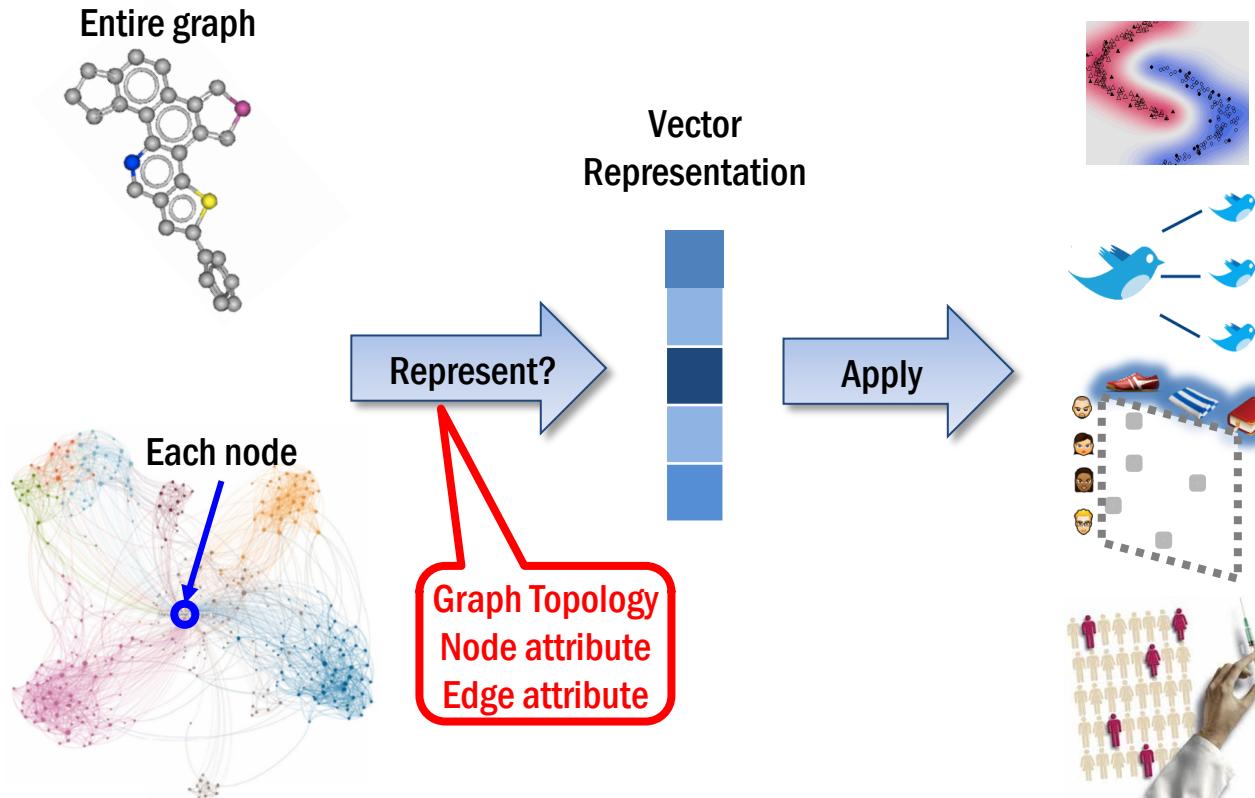
How to represent?
How to update?



Various apps



Fundamental problem and challenge



GCN/GNN/MPN/Structure2Vec

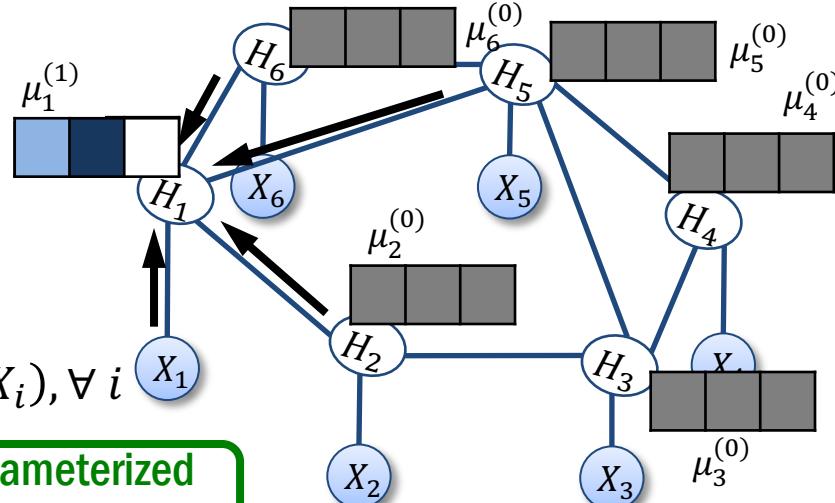
Obtain embedding via

iterative update algorithm:

1. Initialize $\mu_i^{(0)} = \sigma(W_0 X_i), \forall i$

2. Iterate T times

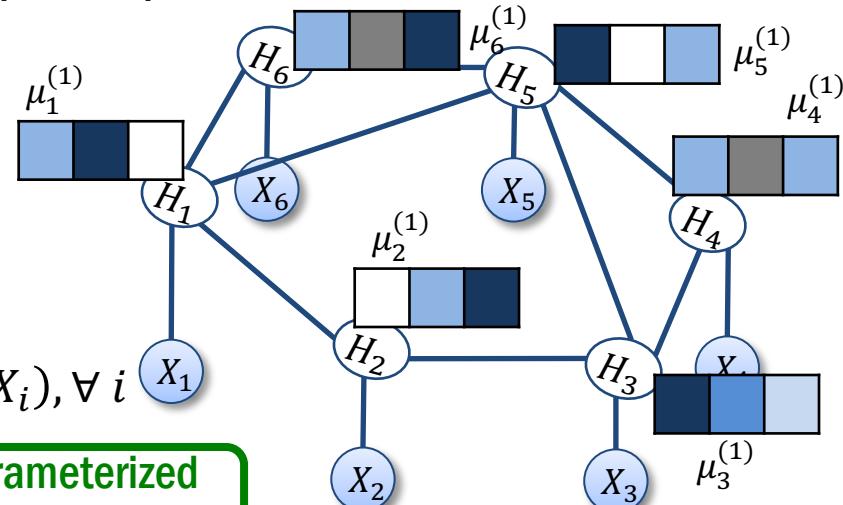
Parameterized
as neural network



$$\mu_i^{(t)} \leftarrow \sigma \left(W_1 \mu_i^{(t-1)} + W_2 \sum_{j \in \mathcal{N}(i)} \mu_j^{(t-1)} \right), \forall i$$

GCN/GNN/MPN/Structure2Vec

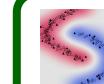
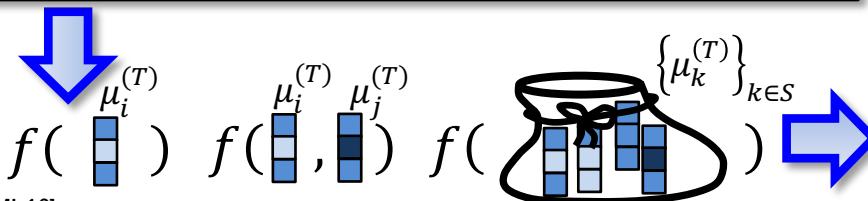
Obtain embedding via



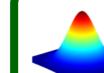
1. Initialize $\mu_i^{(0)} = \sigma(W_0 X_i), \forall i$
2. Iterate T times

Parameterized
as neural network

$$\mu_i^{(t)} \leftarrow \sigma \left(W_1 \mu_i^{(t-1)} + W_2 \sum_{j \in \mathcal{N}(i)} \mu_j^{(t-1)} \right), \forall i$$



Supervised
Learning



Generative
Models



Reinforcement
Learning

Inductive:
Generalize to
New nodes &
Graphs

Variants of graph neural networks

$$\text{Vanilla: } \mu_i^{(t)} \leftarrow \sigma \left(W_1 \mu_i^{(t-1)} + W_2 \sum_{j \in \mathcal{N}(i)} \mu_j^{(t-1)} \right)$$



$$\text{General: } \mu_i^{(t)} \leftarrow Update \left(\mu_i^{(t-1)}, Aggregate \left(\left\{ \mu_j^{(t-1)} \right\}_{j \in S_1}, \left\{ \mu_j^{(t-2)} \right\}_{j \in S_2}, \dots \right) \right)$$



$$\text{Residual: } \mu_i^{(t)} \leftarrow \mu_i^{(t-1)} + \sigma \left(W_1 \mu_i^{(t-1)} + \dots \right)$$

$$\text{Gating: } \mu_i^{(t)} \leftarrow (1 - \beta) \cdot \mu_i^{(t-1)} + \beta \cdot \sigma \left(W_1 \mu_i^{(t-1)} + \dots \right)$$

$$\text{Attention: } \mu_i^{(t)} \leftarrow \sigma \left(\dots + W_2 \sum_{j \in \mathcal{N}(i)} \alpha_j \cdot \mu_j^{(t-1)} \right), \sum_{j \in \mathcal{N}(i)} \alpha_j = 1$$

Different message passing scheme

Obtain embedding via

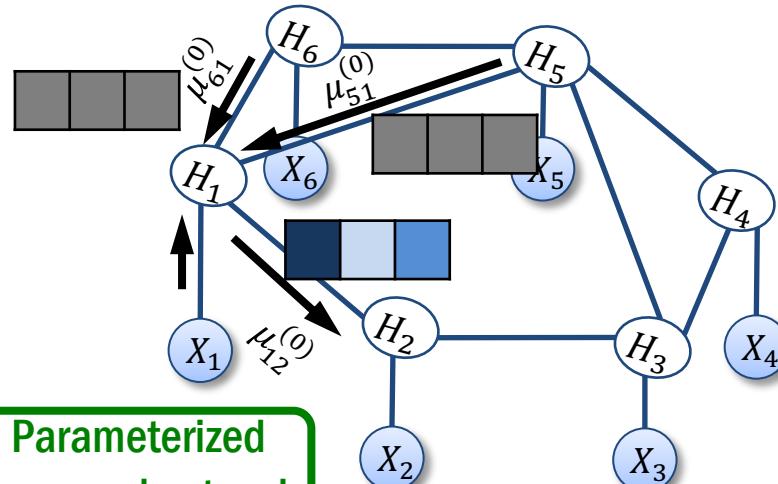
iterative update algorithm:

1. Initialize $\mu_{ij}, \forall (i, j)$

2. Iterate T times

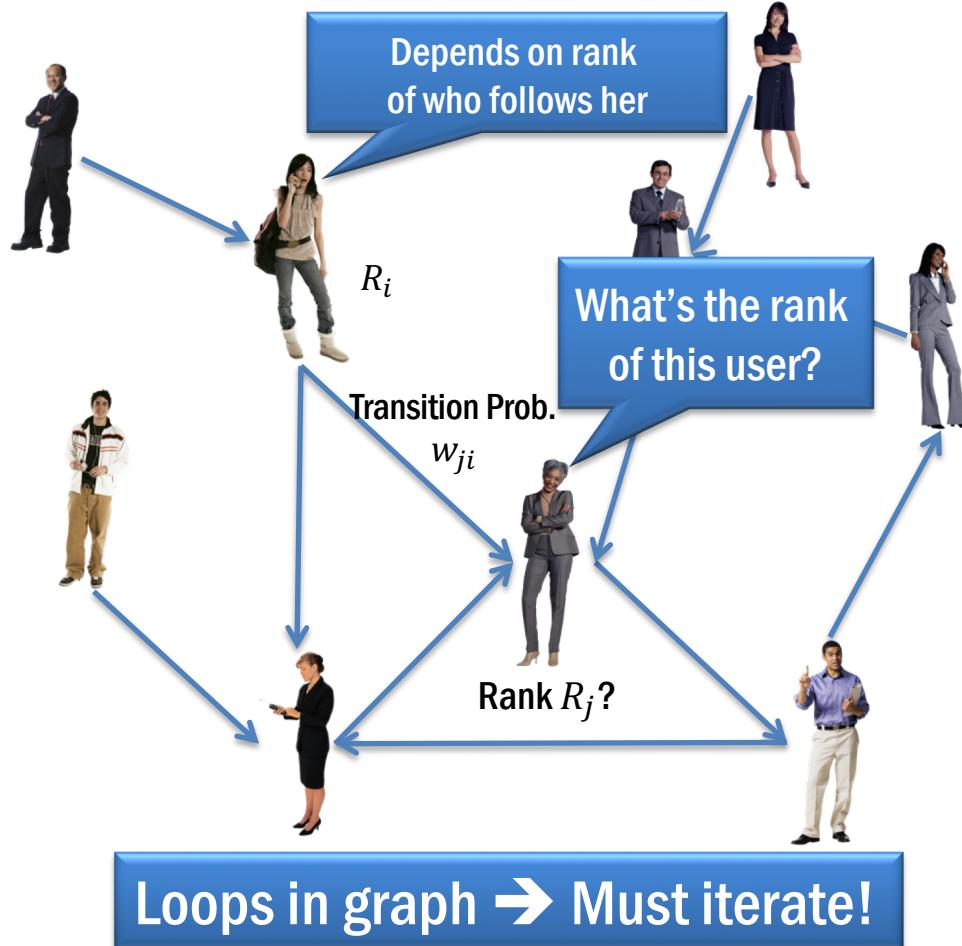
$$\mu_{ij}^{(t)} \leftarrow \sigma \left(W_1 \mu_{ij}^{(t-1)} + W_2 \sum_{\ell \in \mathcal{N}(i) \setminus j} \mu_{\ell i}^{(t-1)} \right), \forall (i, j)$$

3. Aggregate $\mu_i = W_3 \sum_{\ell \in \mathcal{N}(i)} \mu_{\ell i}^{(T)}, \forall i$



Connection to Other Graph Algorithms: More General

PageRank, specific feature



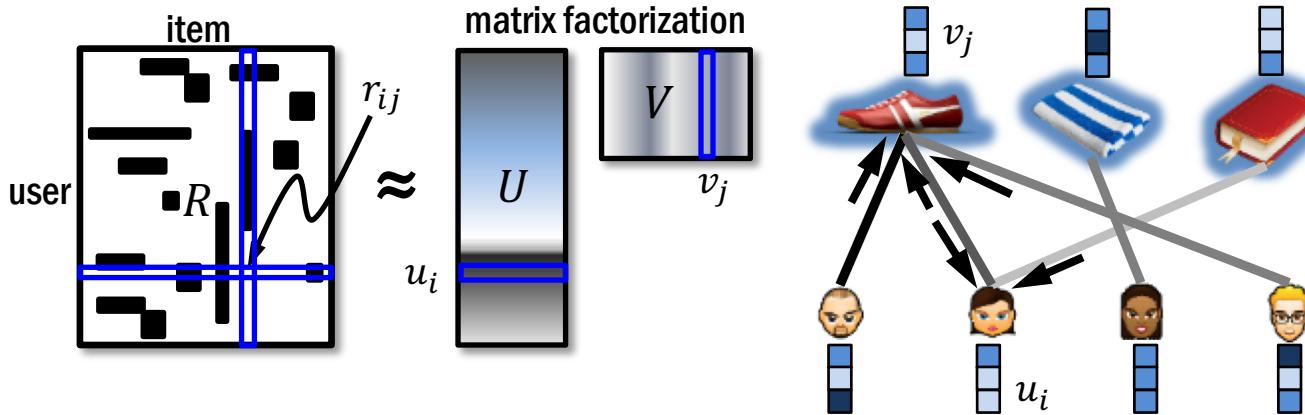
Iterate until convergence

$$R_j \leftarrow \sum_{(j,i) \in E} w_{ji} R_i$$

Fixed update
Extract specifics

Loops in graph → Must iterate!

Matrix factorization, specific feature



Alternating least square $\|R - UV\|_F^2$

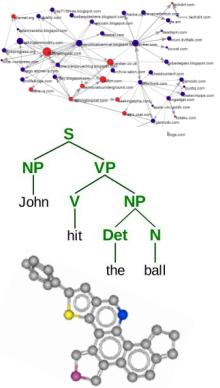
1. Initialize $u_i, v_j, \forall i, j$
2. Iterate T times
 - Update user factors:

$$u_i \leftarrow \operatorname{argmin}_u \sum_{(j,i) \in E} (r_{ij} - u \cdot v_j)^2, \forall i$$

- Update item factors:

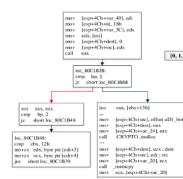
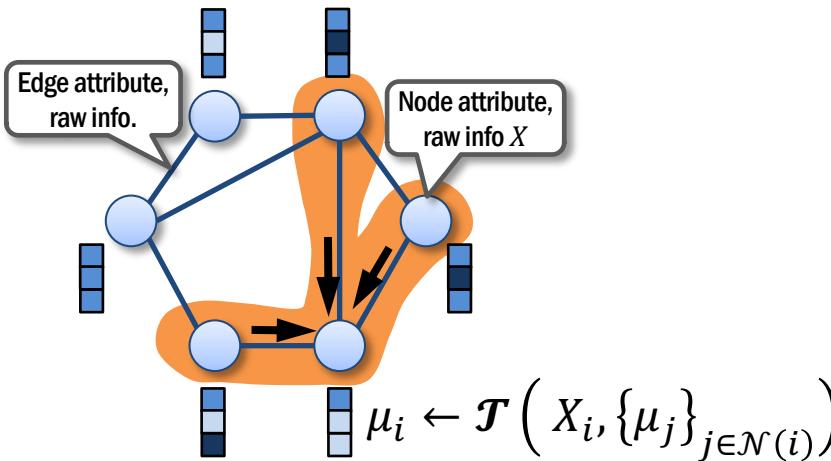
$$v_j \leftarrow \operatorname{argmin}_v \sum_{(i,j) \in E} (r_{ij} - u_i \cdot v)^2, \forall j$$

Fixed update
Extract specifics

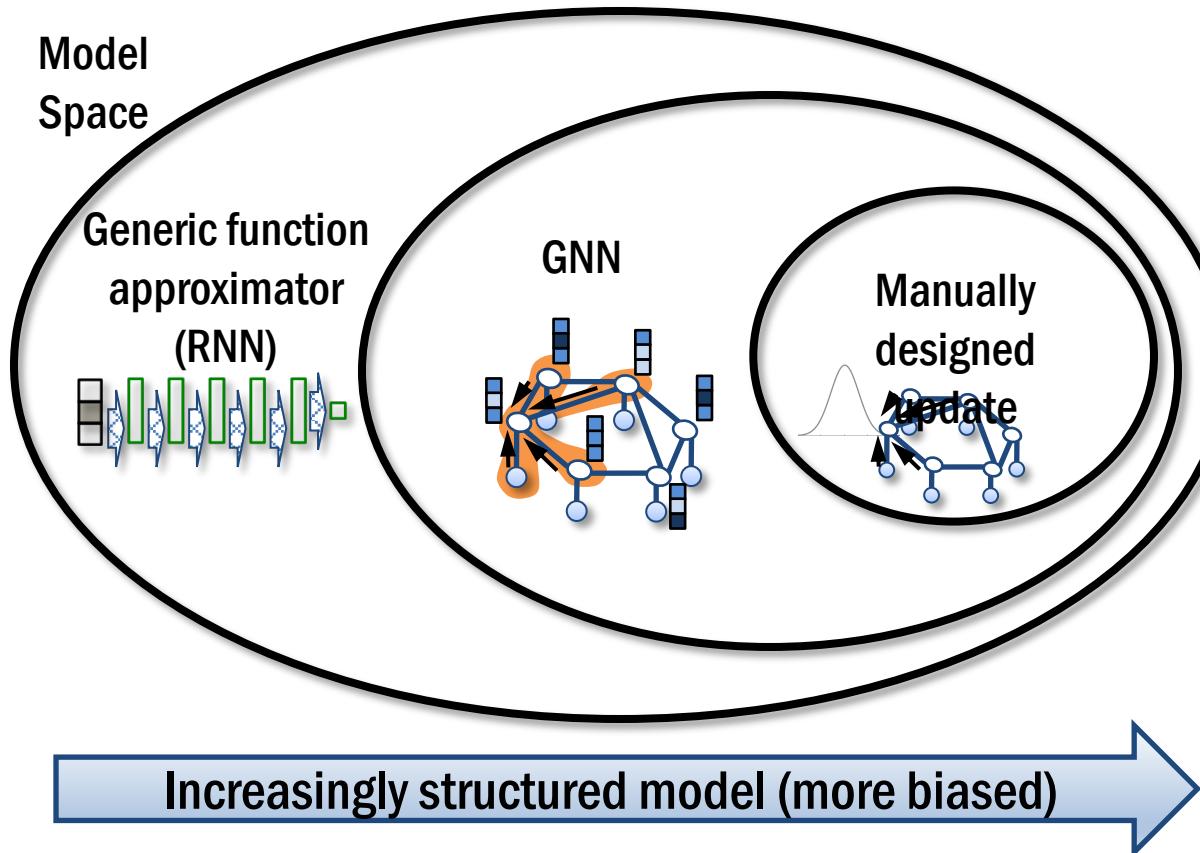


GNN = Parametrized graph algorithm

Graph Algorithm =
Graph Representation + Iterative Update

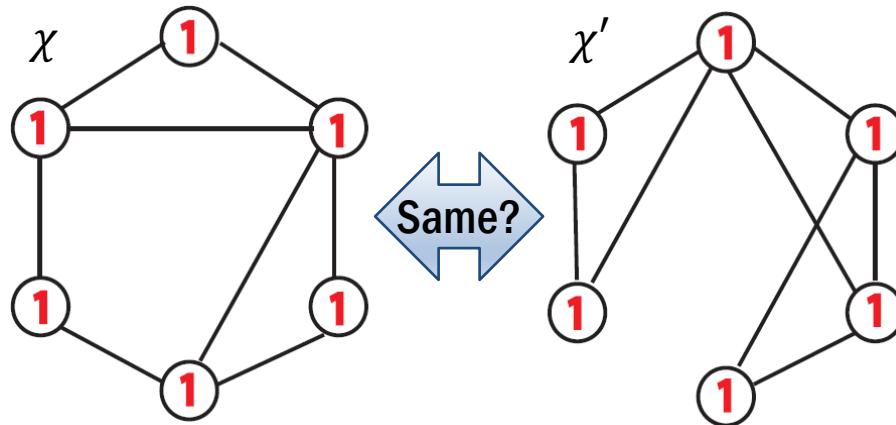


GNN is high structured deep model

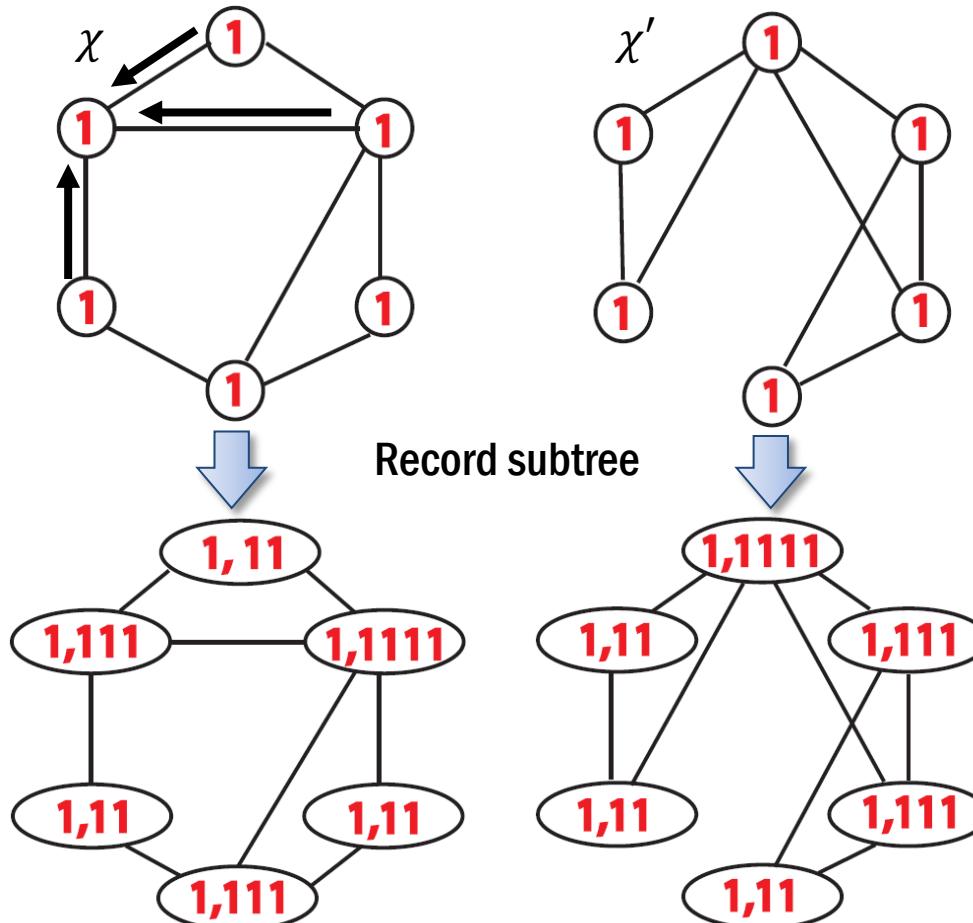


Connection to Graph Isomorphism: Can Represent

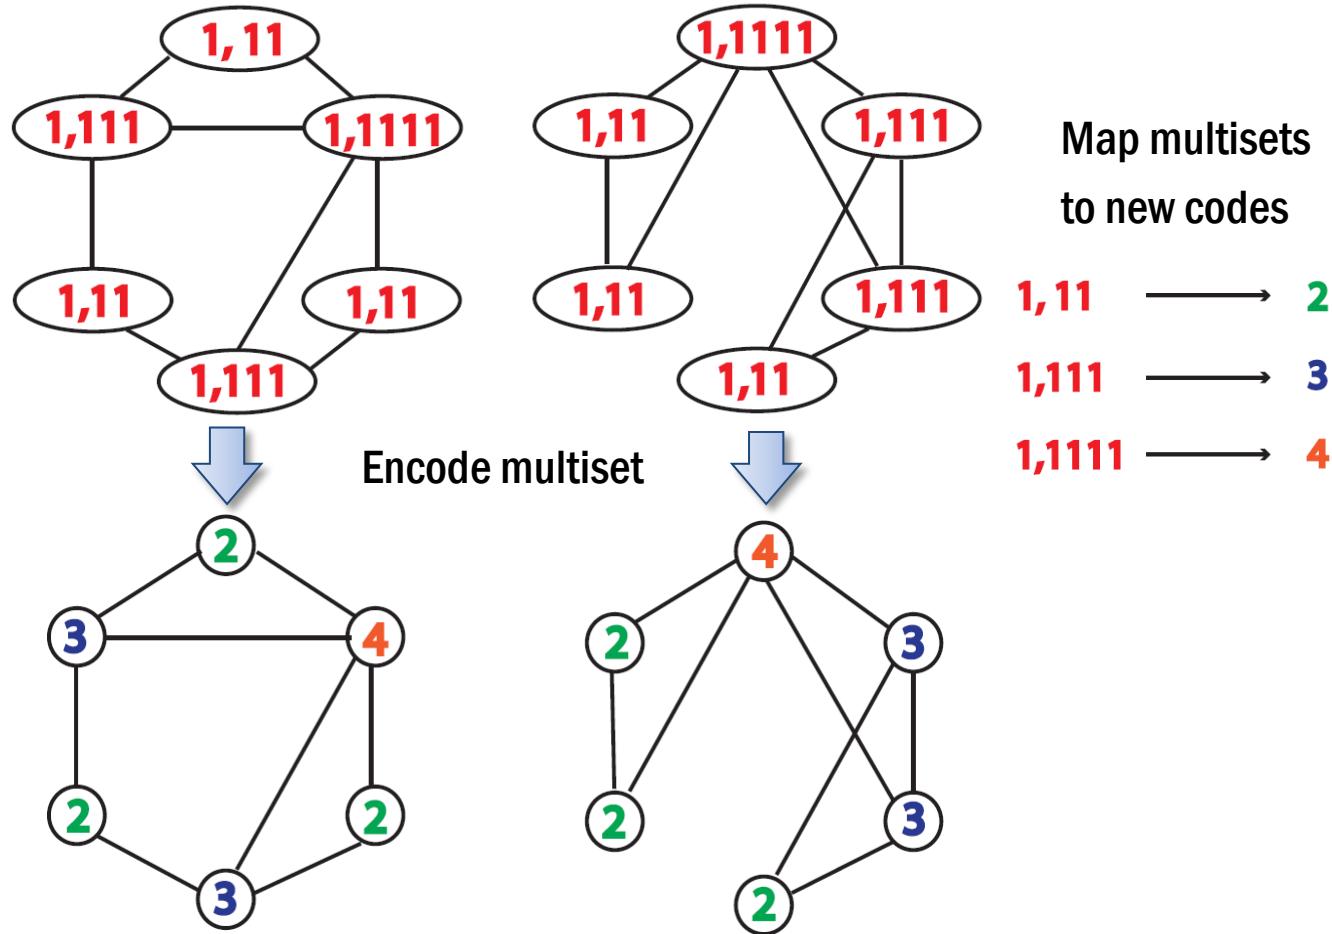
Graph isomorphism test



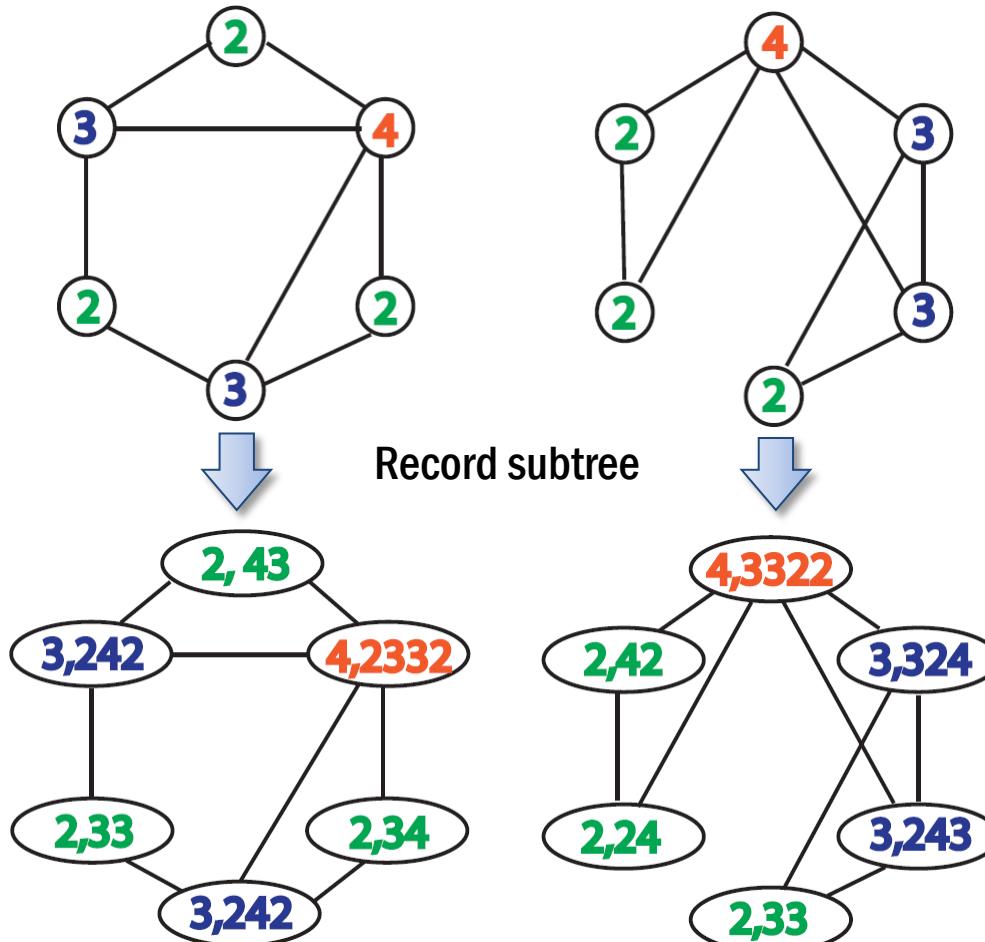
Weisfeiler-Lehmann algorithm: record subtree as multiset



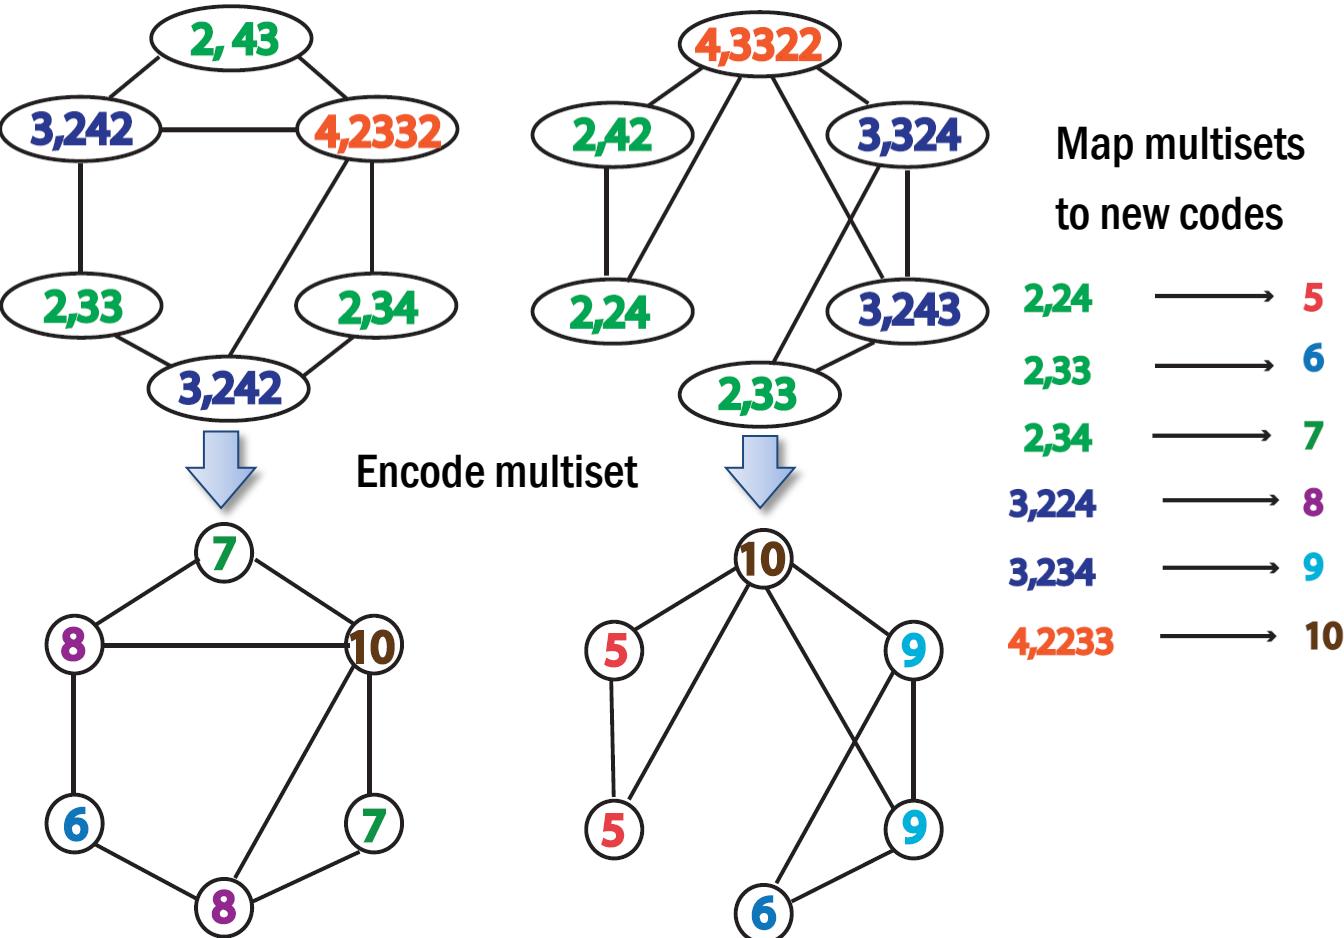
Weisfeiler-Lehmann algorithm: encode (or hash)



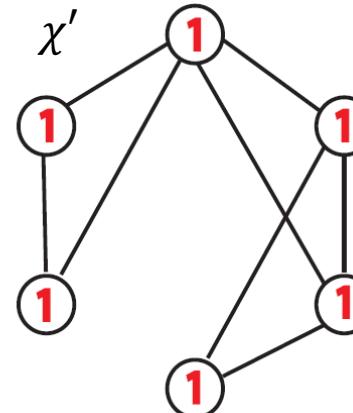
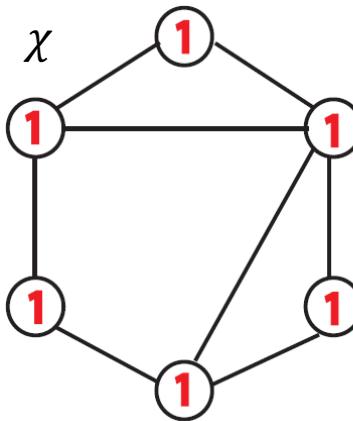
Weisfeiler-Lehmann algorithm: record subtree as multiset again



Weisfeiler-Lehmann algorithm: encode (or hash) again



Weisfeiler-Lehmann algorithm: representation after T iterations



1 2 3 4 5 6 7 8 9 10

$$\phi(\chi) = (6, 3, 2, 1, 0, 1, 2, 2, 0, 1, \dots \dots \dots)$$

$$\phi(\chi') = (6, 3, 2, 1, 2, 1, 0, 0, 2, 1, \dots \dots \dots)$$

Level 0
feature

Level 1
feature

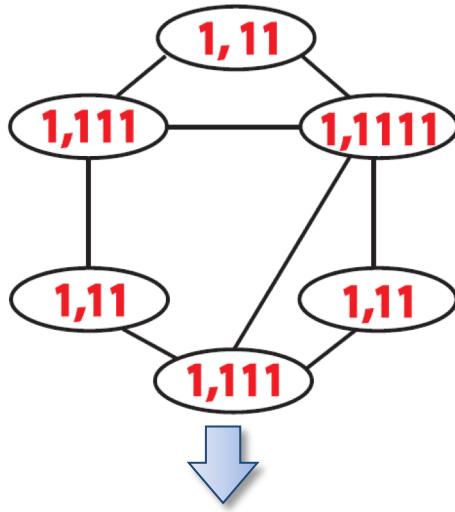
Level 2
feature

Level T
features

Approximate Check!

1. If $\phi(\chi) \neq \phi(\chi')$, graphs not the same
2. Otherwise same graph or can not tell yet

Representation for multiset function



Map multisets
to new codes

$$1,11 \longrightarrow 2$$

$$1,111 \longrightarrow 3$$

$$1,1111 \longrightarrow 4$$

Vocabulary A , multiset $S \subset A$

1. Representation: $\forall g$ multiset function

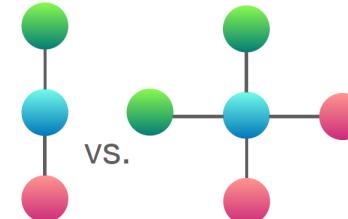
$$g(S) = \gamma(\sum_{x \in S} f(x))$$

with f a vector function

2. One-to-one: $\exists f$ s.t. $\sum_{x \in S} f(x)$ is unique for each finite multiset S

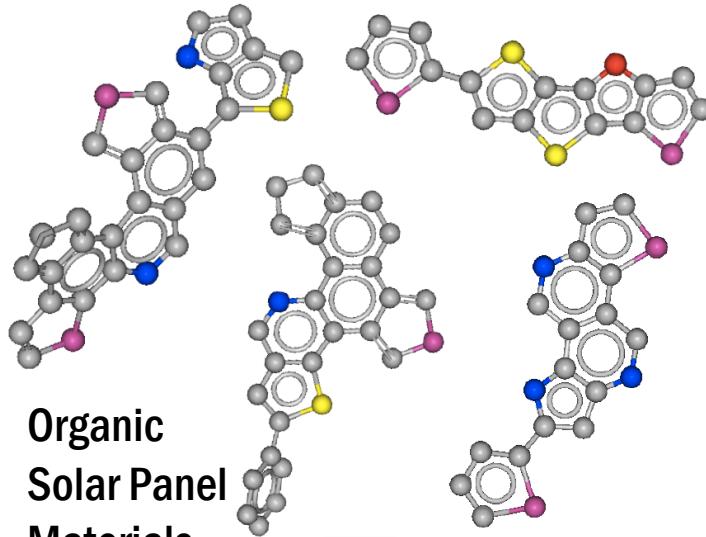
Key update: $\sigma \left(W_1 \mu_i^{(t)} + W_2 \sum_{j \in N(i)} \mu_j^{(t)} \right)$

Average or max-pooling not as expressive



Benefit of GNN for Graph Feature Extraction Algorithm

Materials/Drug design

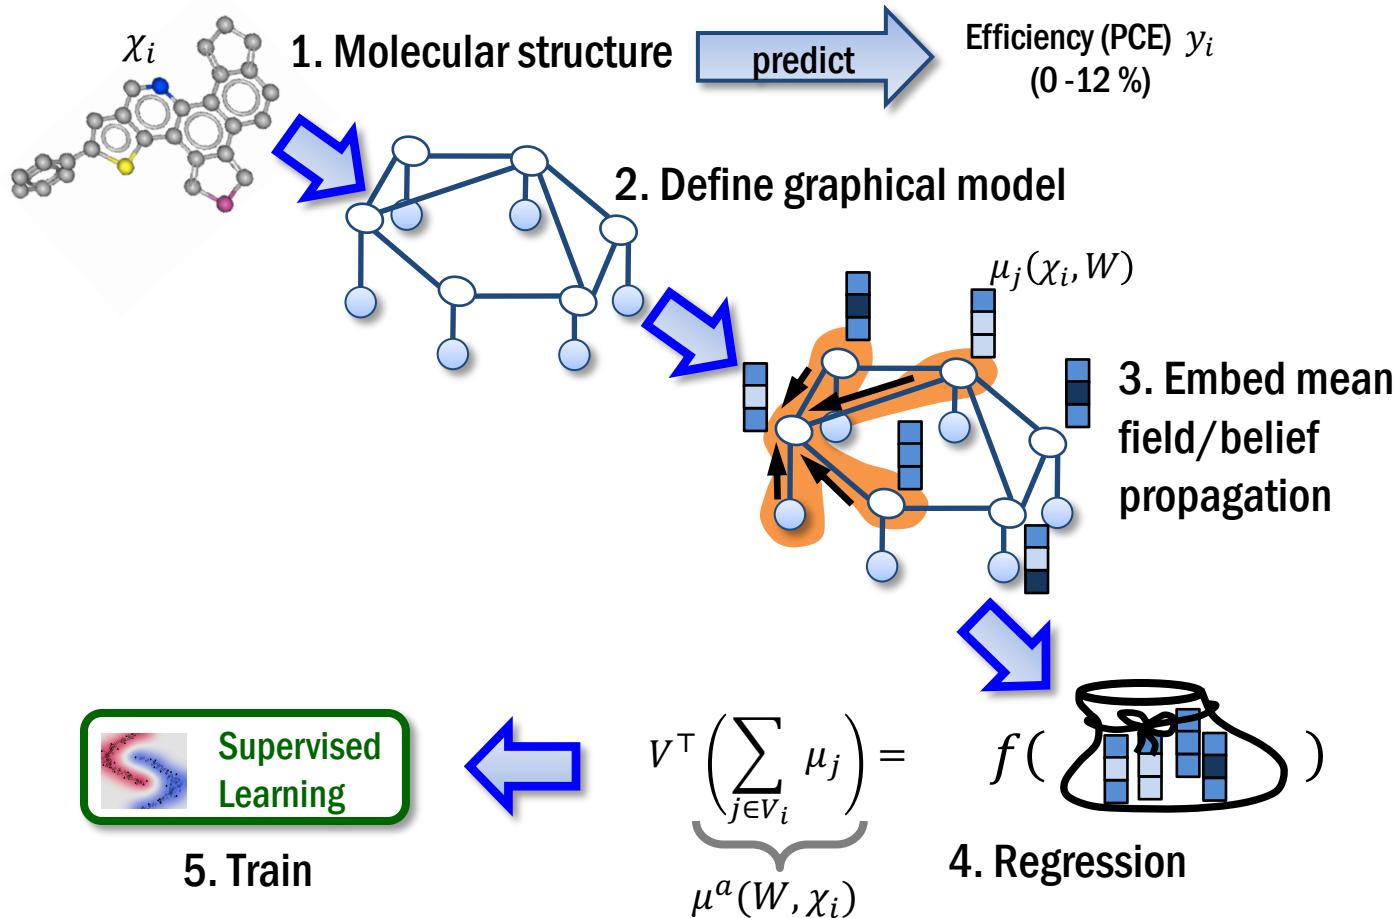


Power Conversion Efficiency (PCE)
(0 - 12 %)

Dataset	Harvard clean energy project
Size	2.3 million
Type	Molecule
Alphabet #	6
Avg node #	28
Avg edge #	33

feature	dimension	MAE
Level 3	1.6 million	0.143
Level 6	1.3 billion	0.096

Prediction for structured data



Parameter learning

Given m data points $\{\chi_1, \chi_2, \dots, \chi_m\}$, estimate parameters W and V which minimize empirical loss

$$\min_{V,W} L(V, W) := \sum_{i=1}^m (y_i - V^\top \mu^a(W, \chi_i))^2$$

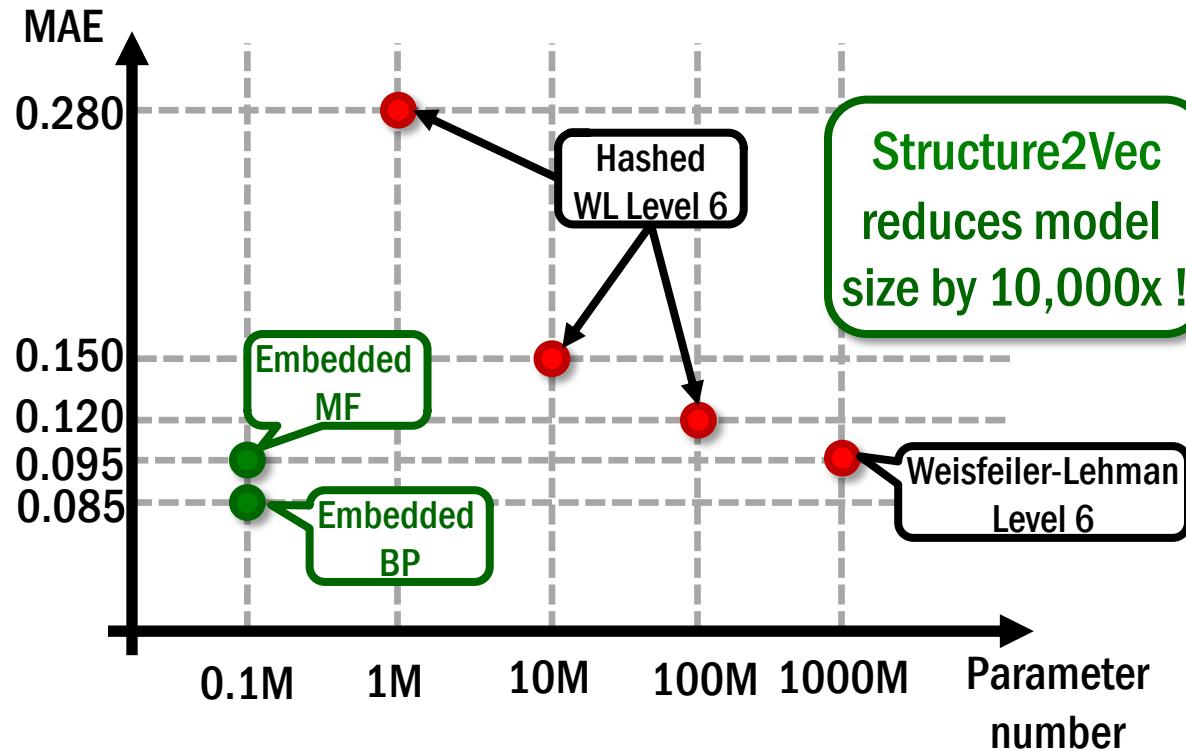
Computation	Operation	Similar to
Objective $L(V, W)$	A forward sequence of nonlinear mappings	Graphical model inference
Gradient $\frac{\partial L}{\partial W}$	Chain rule of derivatives in reverse order	Back propagation in deep learning



More compact model and lower error



Harvard clean energy dataset, 2.3 million organic molecules,
predict power conversion efficiency (0 - 12 %)



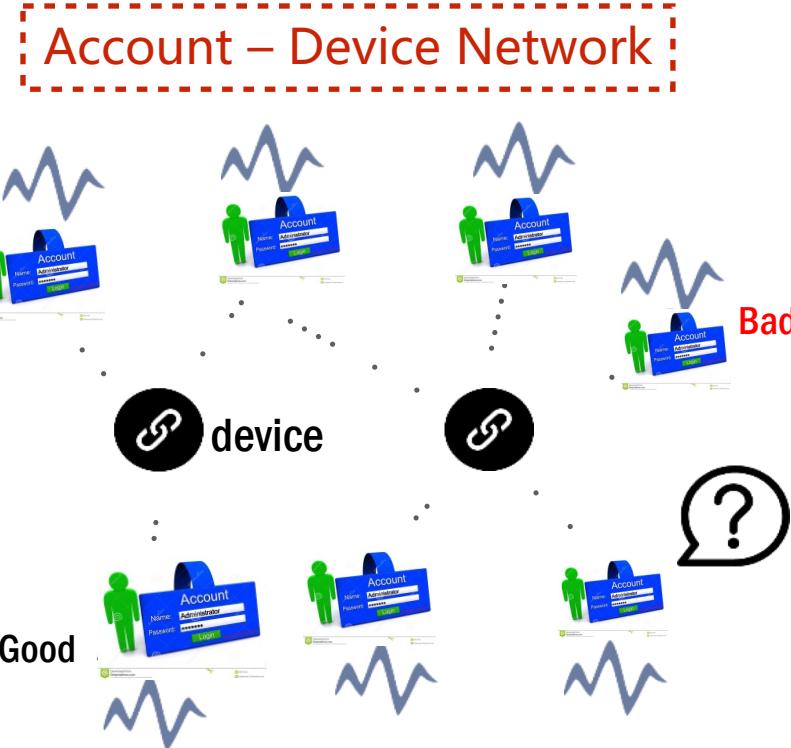
Fraudulent account detection

Alipay: new accounts in a month: millions of nodes and edges.

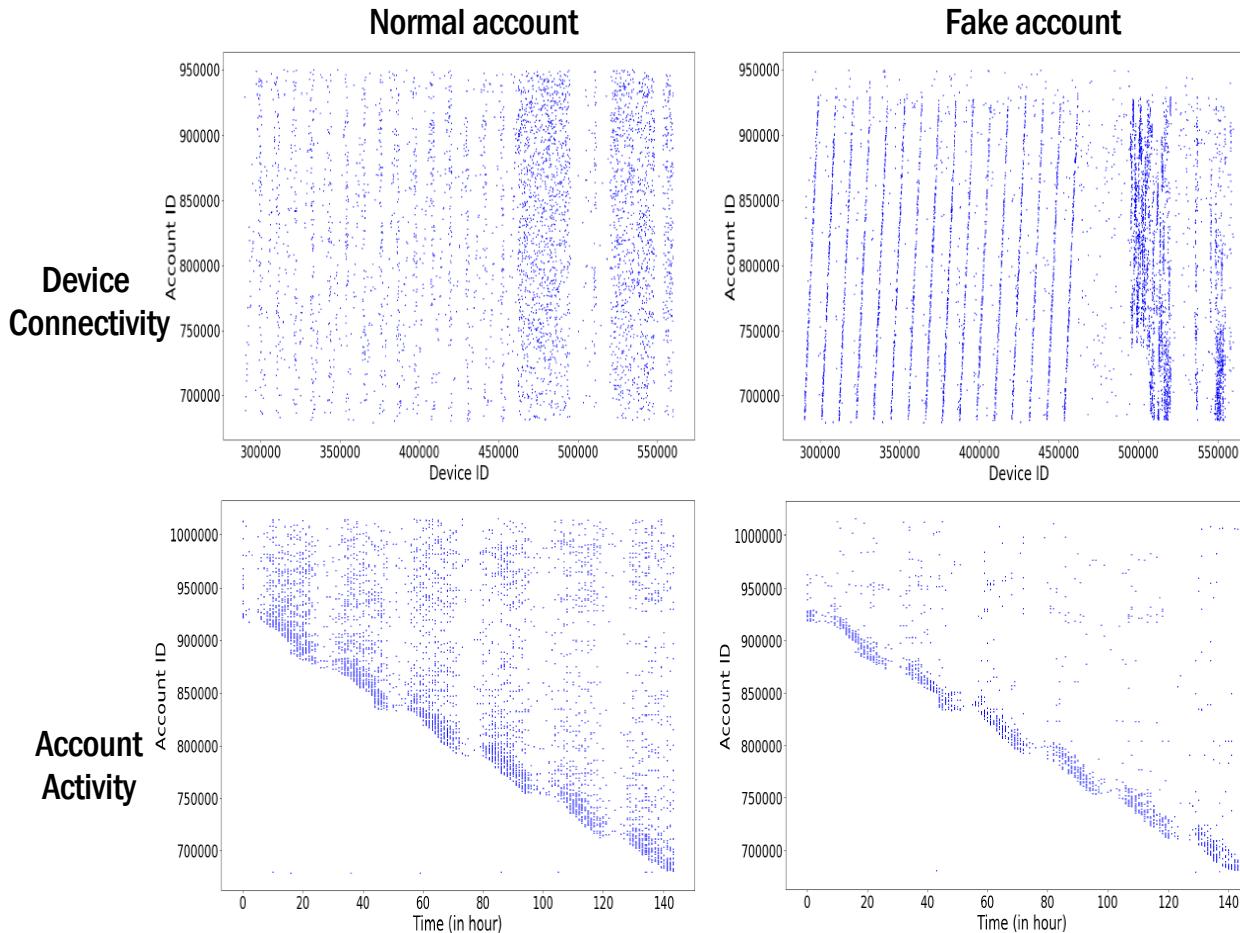
Fake account can increase system level risk



Leverage account activity + connectivity?



Fraudulent account pattern

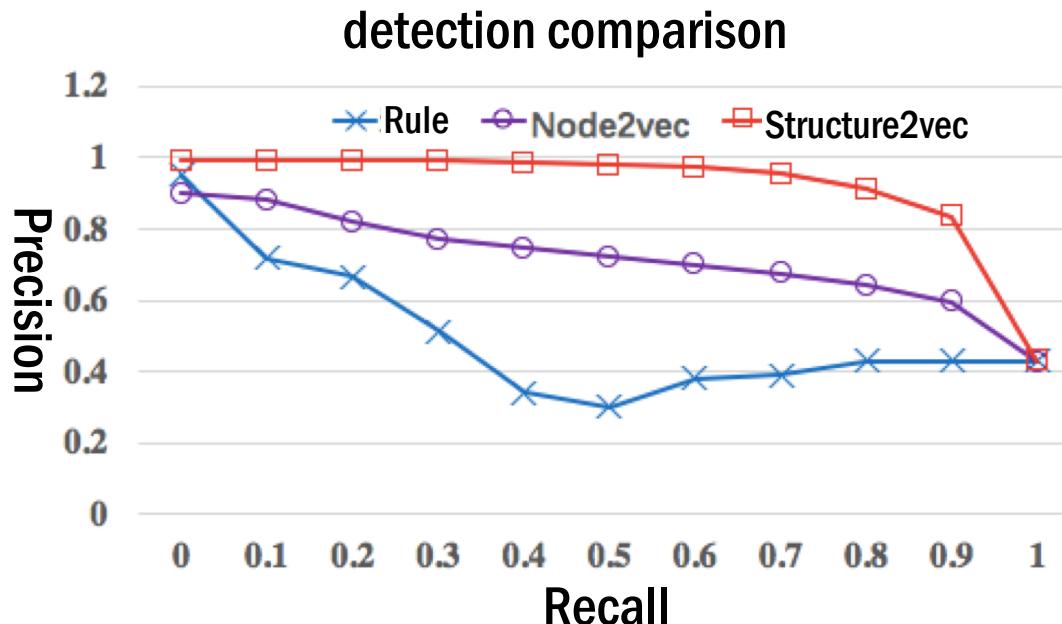


Results

< 1 percent of fraudulent accounts / month

High precision = less disturbance to user experience

High recall = detect more fraudulent account



GNN to Parametrize Combinatorial Optimization Algorithm

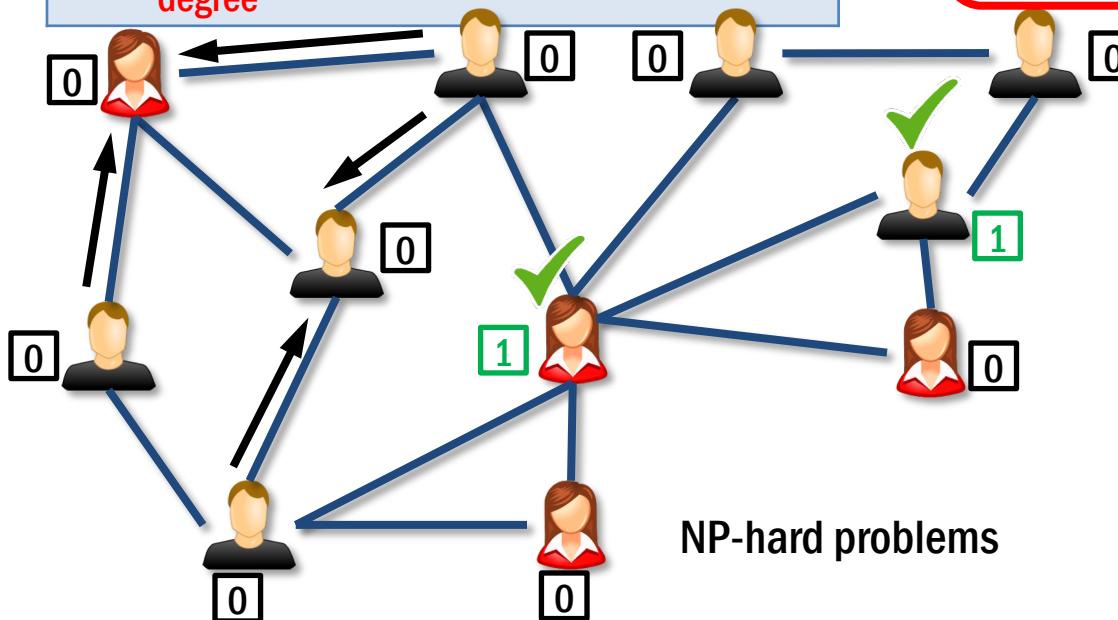
Combinatorial optimization over graph

2 - approximation for minimum vertex cover

Repeat till all edges covered:

- Select uncovered edge with **largest total degree**

Manually
designed rule.
Can we learn
from data?



Greedy algorithm as Markov decision process

Minimum vertex cover: smallest number of nodes to cover all edges

$$\begin{aligned} \min_{x_i \in \{0,1\}} \sum_{i \in \mathcal{V}} x_i \\ s.t. x_i + x_j \geq 1, \forall (i,j) \in \mathcal{E} \end{aligned}$$

Repeat:

1. Compute **total degree** of each uncovered edge

2. Select both ends of uncovered edge with largest total degree

Until all edges are covered

Reward: $r^t = -1$

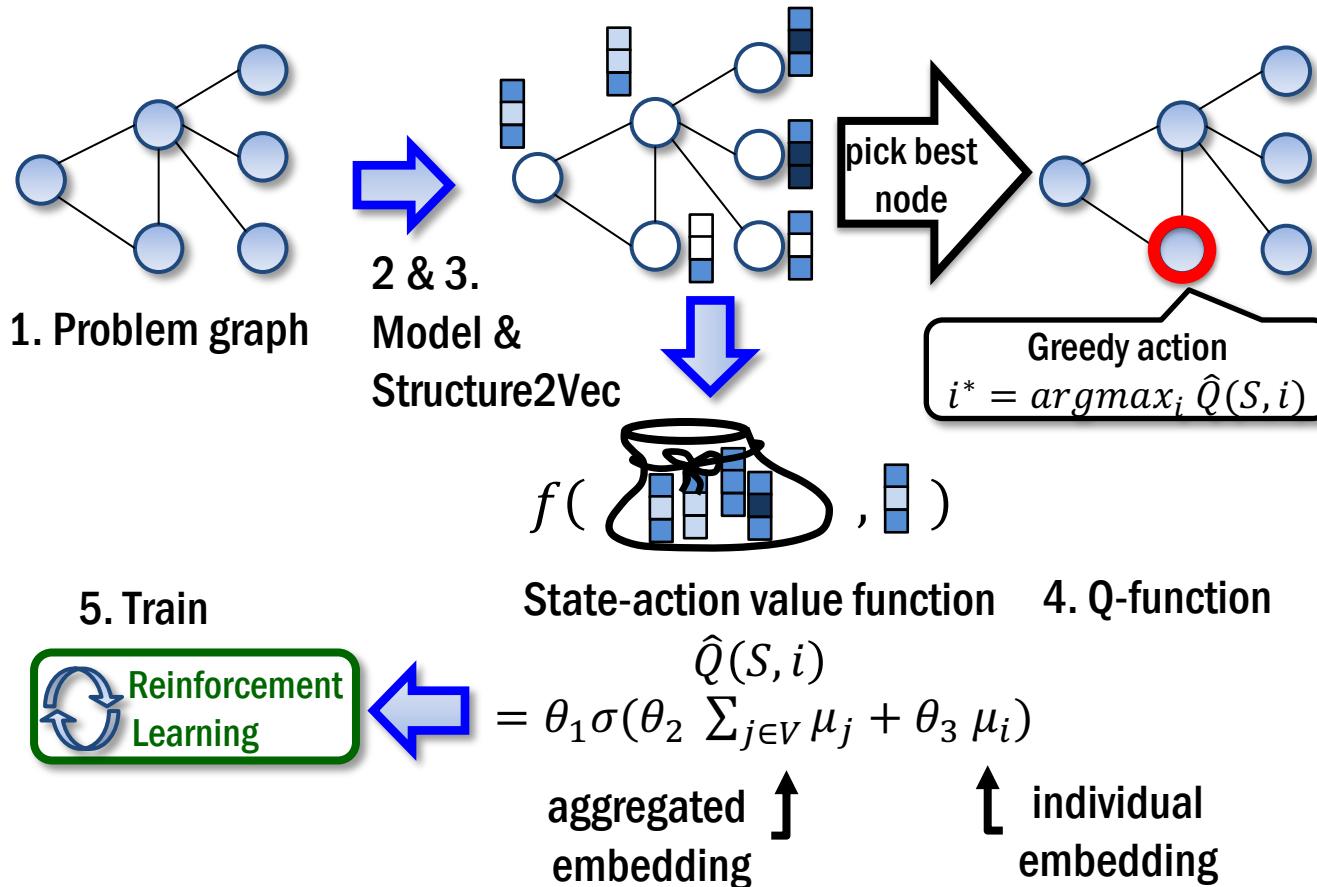
State S : current selected nodes

Action value function: $\hat{Q}(S, i)$

Greedy policy:
 $i^* = \operatorname{argmax}_i \hat{Q}(S, i)$

Update state S

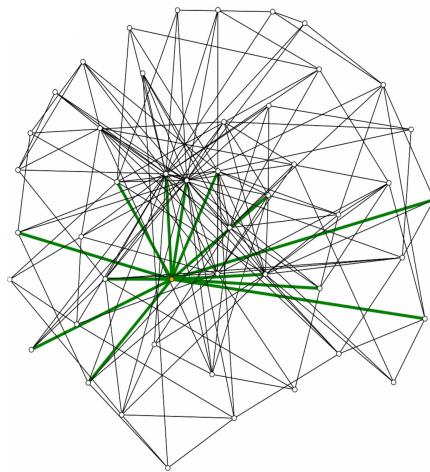
Embedding for state-action value function



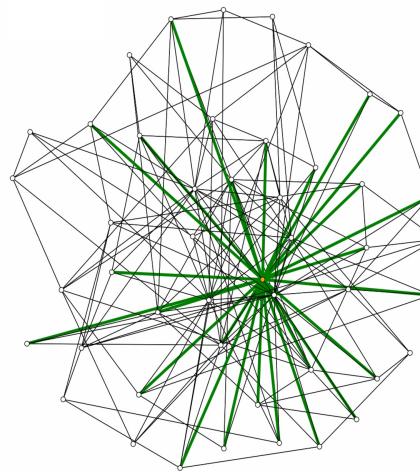
What new algorithm is learned?

Learned algorithm balances between

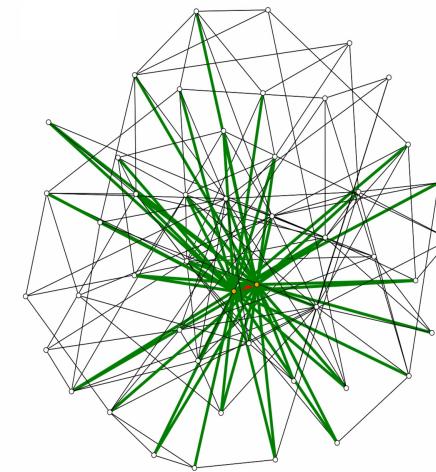
- degree of the picked node and
- fragmentation of the graph



Structure2Vec



Node greedy

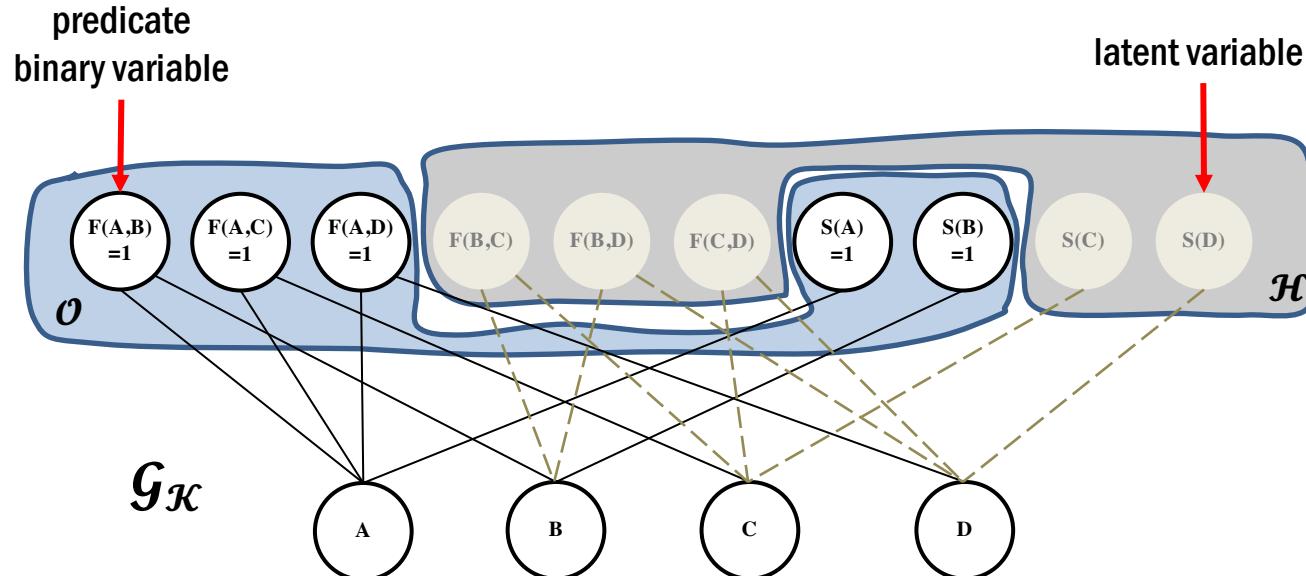


Edge greedy

GNN to Parametrize Variational Inference Algorithm for Probabilistic Logic

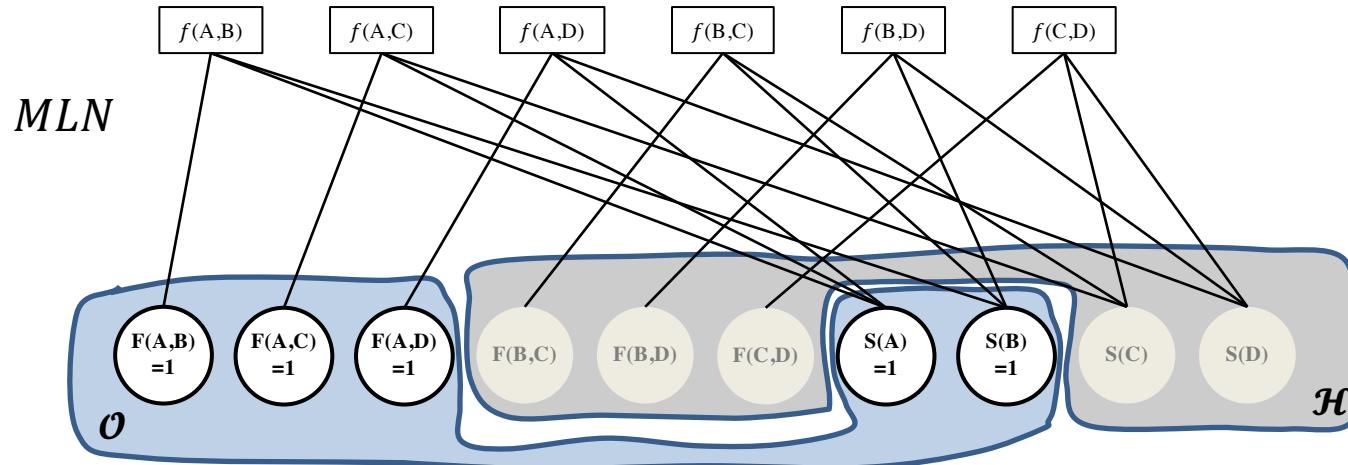
Factor graph representation for knowledge base

- Entity (constant), $\mathcal{C} = \{A, B, C, D \dots\}$
- Predicate (attribute | relation), $r(\cdot): \mathcal{C} \times \mathcal{C} \times \dots \mapsto \{0,1\}$
 - Eg. $\text{Smoke}(x)$, $\text{Friend}(x,x')$, $\text{Like}(x,x')$



Markov logic networks

- Use logic formula $f(\cdot): \mathcal{C} \times \mathcal{C} \times \dots \times \mathcal{C} \mapsto \{0,1\}$ for potential functions
 - Eg. formula $f(A,B): \text{Friend}(A,B) \wedge \text{Smoke}(A) \Rightarrow \text{Smoke}(B)$

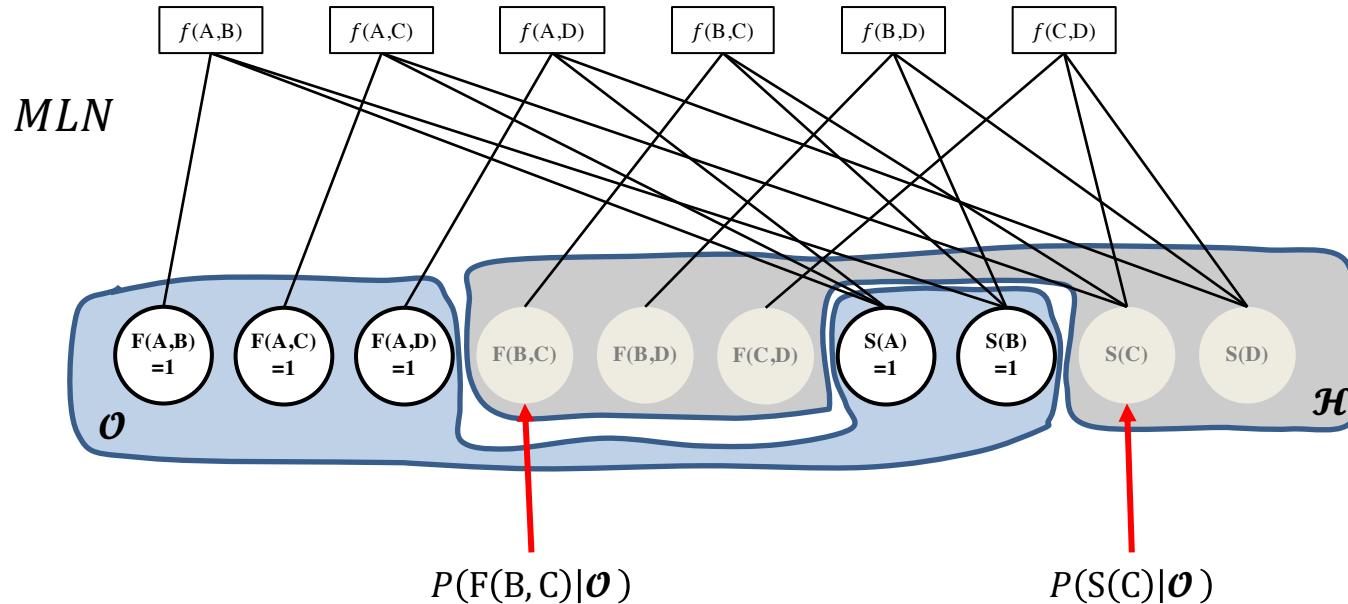


$$P(\mathcal{O}, \mathcal{H}) = \frac{1}{Z} \exp \left(\sum_f w_f \sum_{a_f} \phi_f(a_f) \right)$$

w_f : formula weight, $\phi_f(x, x')$: $\neg F(x, x') \vee \neg S(x) \vee S(x')$, Z : normalization constant

Challenges in inference

- A large grounded network, $O(n^2)$ in the number n of entities!
- Enumerate configuration over $O(n^2)$ binary variables, with $O(2^{n^2})$ possibilities.

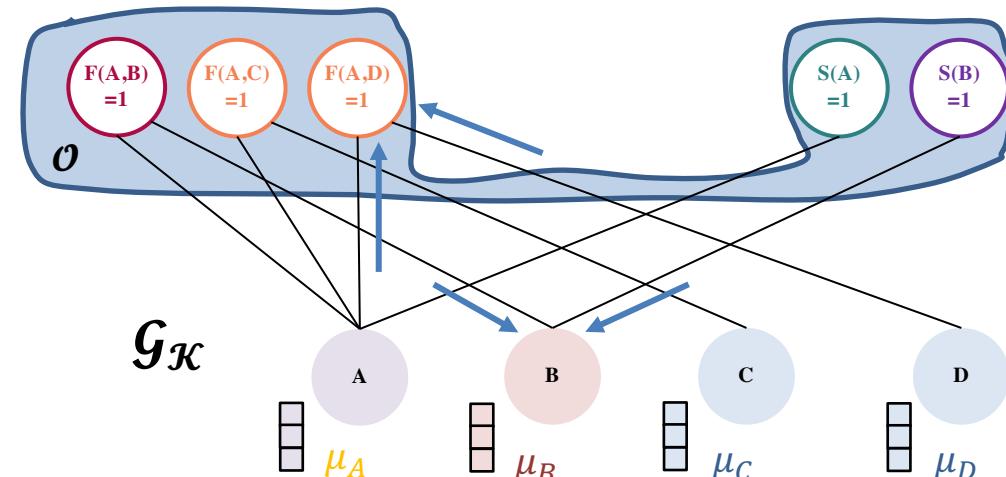


Efficient inference? Most previous works are on grounded networks

Use GNN for variational inference?

- GNN on original KB ($\mathcal{G}_{\mathcal{K}}$) to get embedding $\mu_A, \mu_B, \mu_C, \mu_D$ for entities

$$\mu_A, \mu_B, \mu_C, \mu_D = GNN(\mathcal{G}_{\mathcal{K}}; \theta)$$



Iterative update $t = 0, \dots, T$:

$$\mu_B^{t+1} = MLP \left(\mu_B^t + \sum_{r \in N(B)} \mu_r^t \right)$$

$$\mu_r^{t+1} = MLP \left(\mu_r^t + \sum_{c \in N(r)} \mu_c^t \right)$$

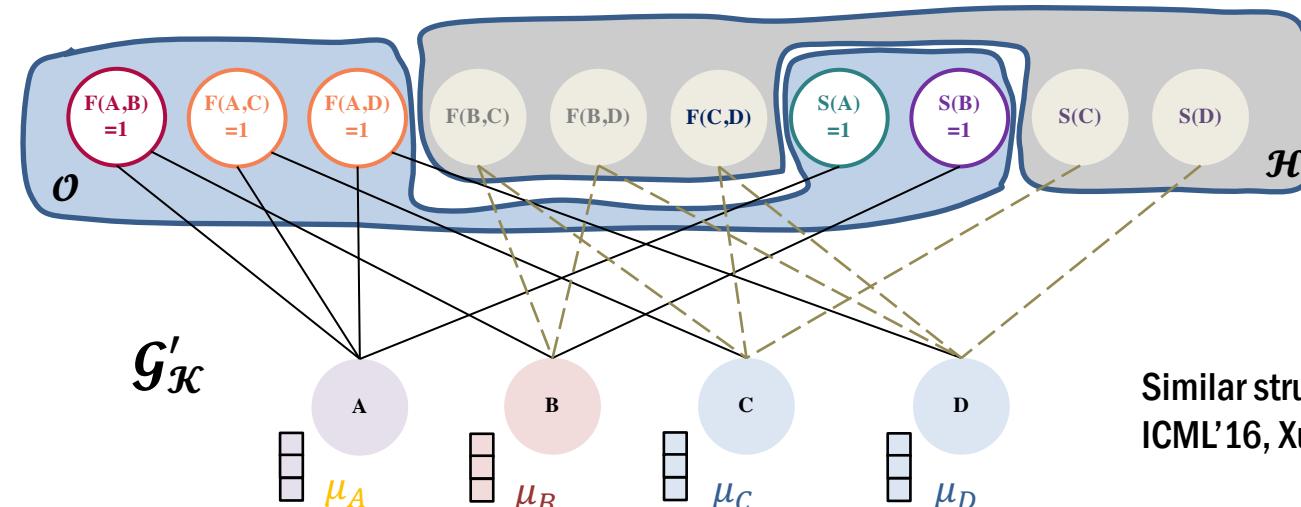
Similar structure = similar GNN embedding (Dai ICML'16, Xu ICLR'19)

Use GNN for variational inference? (eg. Inside VAE)

- GNN on original KB ($\mathcal{G}_{\mathcal{K}}$) to get embedding $\mu_A, \mu_B, \mu_C, \mu_D$ for entities

$$\mu_A, \mu_B, \mu_C, \mu_D = GNN(\mathcal{G}_{\mathcal{K}}; \theta)$$

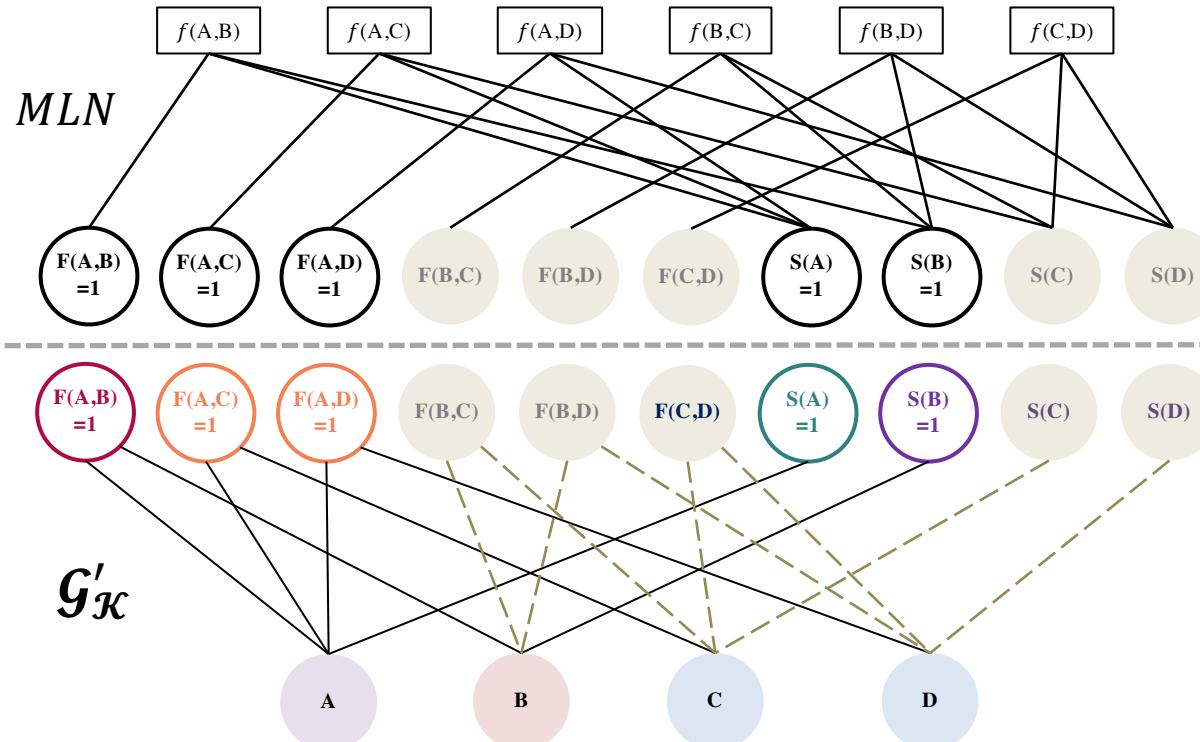
- Define $Q(F(B, C) = 1 | \mathcal{O}) = \frac{1}{1 + \exp(\mu_B^\top \Theta_F \mu_C)}$, $Q(S(C) = 1 | \mathcal{O}) = \frac{1}{1 + \exp(\theta_S^\top \mu_C)}$



Train the parameters
 $\theta, \Theta_F, \theta_S, \dots$
with mean field
variational inference

Similar structure = similar GNN embedding (Dai
ICML'16, Xu ICLR'19)

Is GNN embedding expressive enough?



$$F(x, y) \xrightarrow{MLN} F(x', y')$$

$$F(x, y) \xrightarrow{g'_K} F(x', y')$$

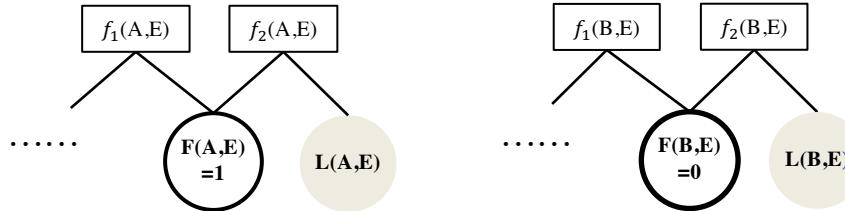
is this true?

if and only if

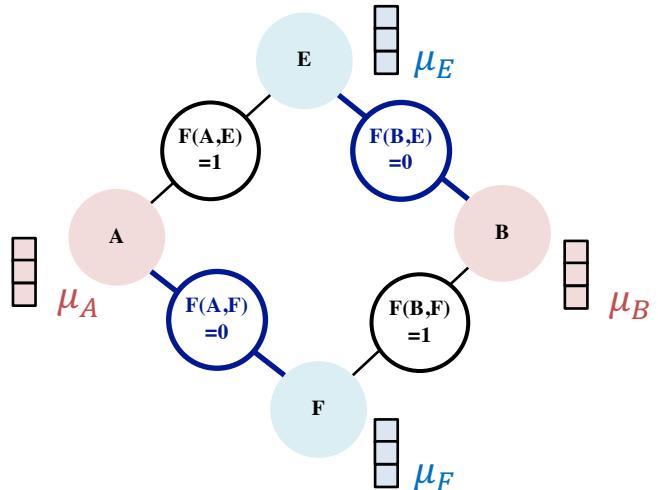
$$x \xrightarrow{g_K} x' \text{ and } y \xrightarrow{g_K} y'$$

ExpressGNN: add a tunable embedding

- formula 1 $f_1(A,B)$: $\text{Friend}(A,B) \wedge \text{Smoke}(A) \Rightarrow \text{Smoke}(B)$
- formula 2 $f_2(A,B)$: $\text{Friend}(A,B) \Rightarrow \text{Like}(A,B)$



$L(A,E)$ and $L(B,E)$ are different.



A and B are the same in KB.

$$\mu_A, \mu_B, \mu_E, \mu_F = GNN(\mathcal{G}_K; \theta)$$

$\omega_A, \omega_B, \omega_E, \omega_F \leftarrow$ tunable low dimensional embedding

$$Q(L(A, E) = 1 | \mathcal{O}) = \frac{1}{1 + \exp(\mu_A^\top \Theta_L \mu_E + \omega_A^\top \omega_E)}$$

$$Q(L(B, E) = 1 | \mathcal{O}) = \frac{1}{1 + \exp(\mu_B^\top \Theta_L \mu_E + \omega_B^\top \omega_E)}$$

- 22 relations
 - Teach, publish ...
- Task goal
 - Predict who is whose advisor
 - Zero observed facts for query predicates

Counting	UW-CSE				
	AI	Graphics	Language	Systems	Theory
# entity	300	195	82	277	174
# relation	22	22	22	22	22
# fact	731	449	182	733	465
# query	4K	4K	1K	5K	2K
# ground atom	95K	70K	15K	95K	51K
# ground formula	73M	64M	9M	121M	54M

- 94 crowd-sourced FOL formulas

$\text{advisedBy}(s, p) \Rightarrow \text{professor}(p)$

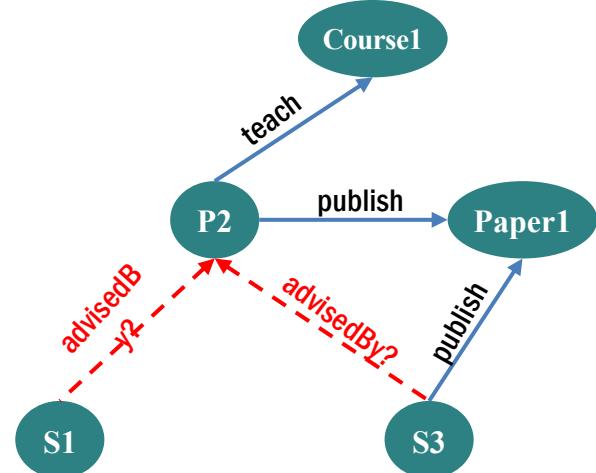
$\text{advisedBy}(s, p) \Rightarrow \neg \text{yearsInProgram}(s, \text{Year_1})$

$\text{professor}(x) \Rightarrow \neg \text{student}(y)$

$\text{publication}(p, x) \vee \text{publication}(p, y) \vee \text{student}(x) \vee \neg \text{student}(y) \Rightarrow \text{professor}(y)$

$\text{student}(x) \vee \neg \text{advisedBy}(x, y) \Rightarrow \text{tempAdvisedBy}(x, y)$

...



Cora dataset details (CS paper citation)

- 10 relation types
 - Author, Title, Venue, HasWordTitle, ...
- Task goal (entity resolution)
 - De-duplicate citations, authors, and venues
 - Zero observed facts for query predicates
- 46 crowd-sourced FOL formulas

$\text{Author}(bc1, a1) \vee \text{Author}(bc2, a2) \vee \text{SameAuthor}(a1, a2) \Rightarrow \text{SameBib}(bc1, bc2)$

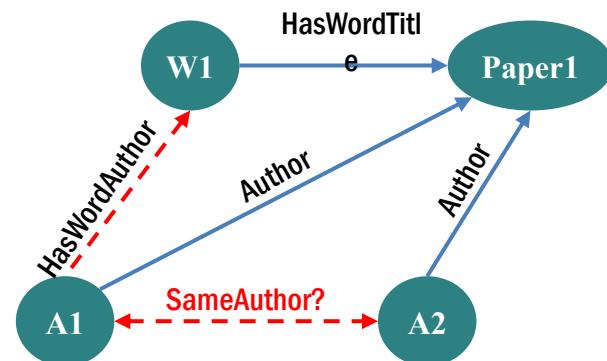
$\text{HasWordAuthor}(a1, w) \vee \text{HasWordAuthor}(a2, w) \Rightarrow \text{SameAuthor}(a1, a2)$

$\text{Title}(bc1, t1) \vee \text{Title}(bc2, t2) \vee \text{SameBib}(bc1, bc2) \Rightarrow \text{SameTitle}(t1, t2)$

$\text{SameVenue}(v1, v2) \vee \text{SameVenue}(v2, v3) \Rightarrow \text{SameVenue}(v1, v3)$

$\text{Title}(bc1, t1) \vee \text{Title}(bc2, t2) \vee \text{HasWordTitle}(t1, +w) \vee \text{HasWordTitle}(t2, +w) \Rightarrow \text{SameBib}(bc1, bc2)$

Counting	Cora (average)
# entity	616
# relation	10
# fact	12K
# query	2K
# ground atom	157K
# ground formula	457M



Inference accuracy and time

- Area under precision-recall curve (AUC-PR)

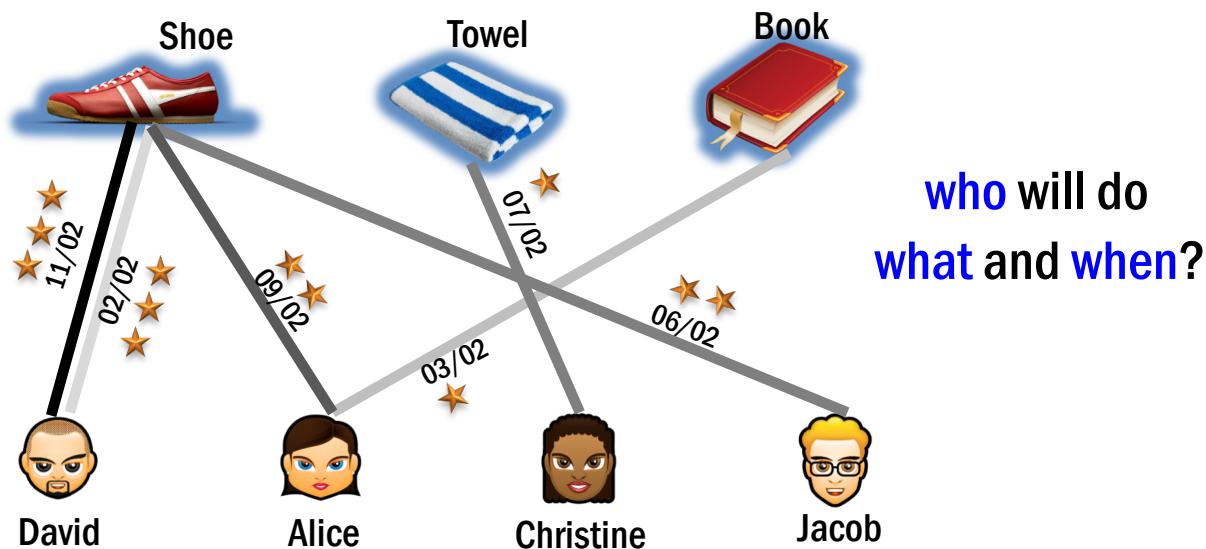
Method	UW-CSE					Cora (avg)
	AI	Graphics	Language	Systems	Theory	
MCMC	-	-	-	-	-	-
BP / Lifted BP	0.01	0.01	0.01	0.01	0.01	-
MC-SAT	0.03	0.05	0.06	0.02	0.02	-
HL-MRF	0.06	0.06	0.02	0.04	0.03	-
ExpressGNN	0.09	0.19	0.14	0.06	0.09	0.64

- Inference wall clock time

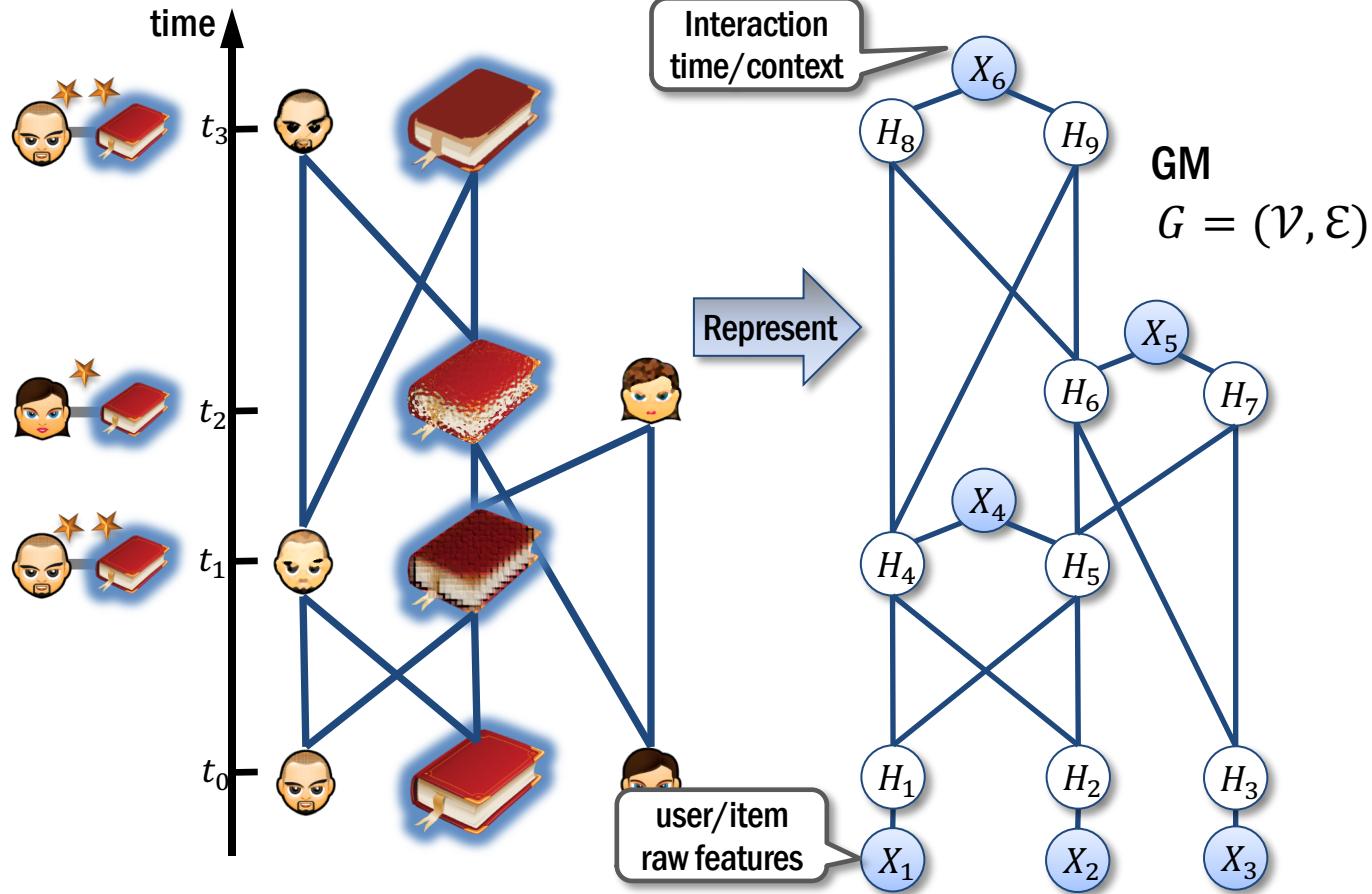
Method	Inference Time (minutes)				
	AI	Graphics	Language	Systems	Theory
MCMC	>24h	>24h	>24h	>24h	>24h
BP	408	352	37	457	190
Lifted BP	321	270	32	525	243
MC-SAT	172	147	14	196	86
HL-MRF	135	132	18	178	72
ExpressGNN	14	20	5	7	13

GNN to Parametrize Algorithm for Dynamic Networks

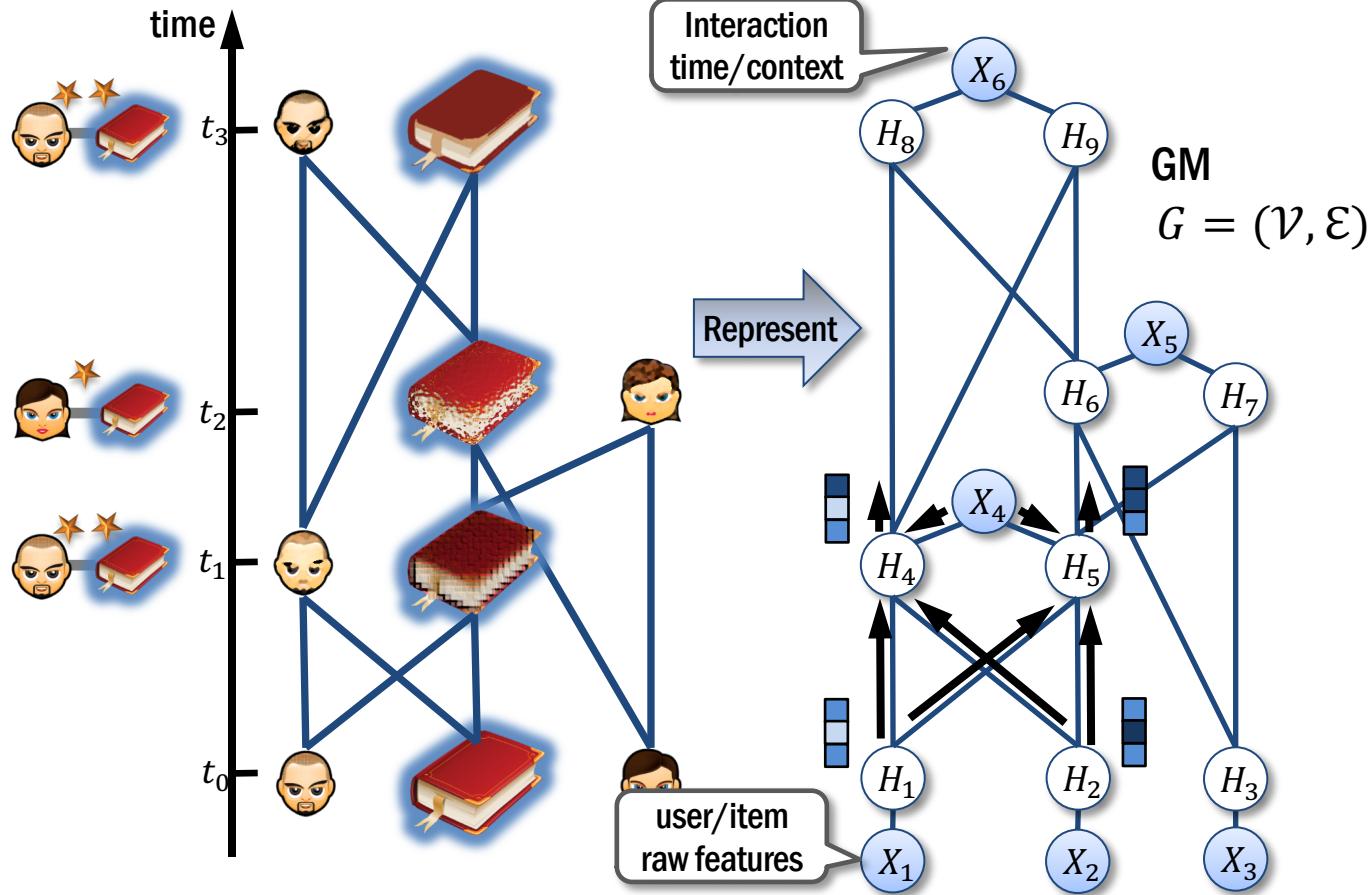
Dynamic processes over networks



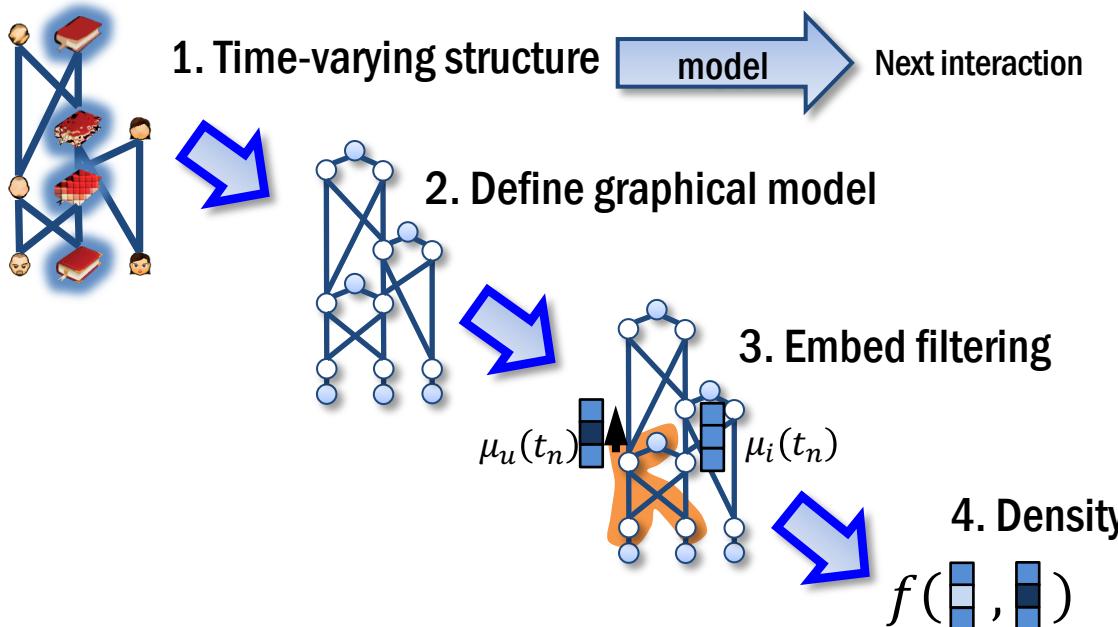
Unroll: time-varying dependency structure



Forward graph neural networks



Embedding filtering algorithm for generative model



Compatibility between user u and item i

$$\alpha_{ui} = \exp(\mu_u^\top(t_n) \mu_i(t_n))$$

Likelihood of next event time $p_{ui}(t)$

$$\alpha_{ui}(t - t_n) \exp\left(-\frac{\alpha_{ui}(t - t_n)^2}{2}\right)$$



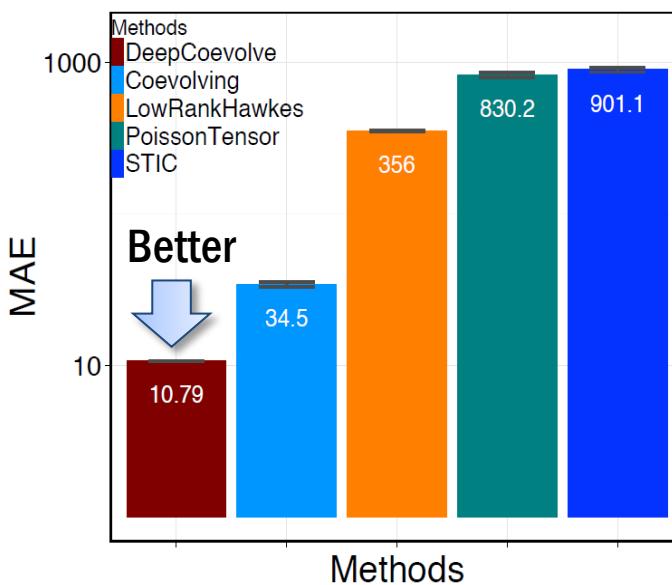
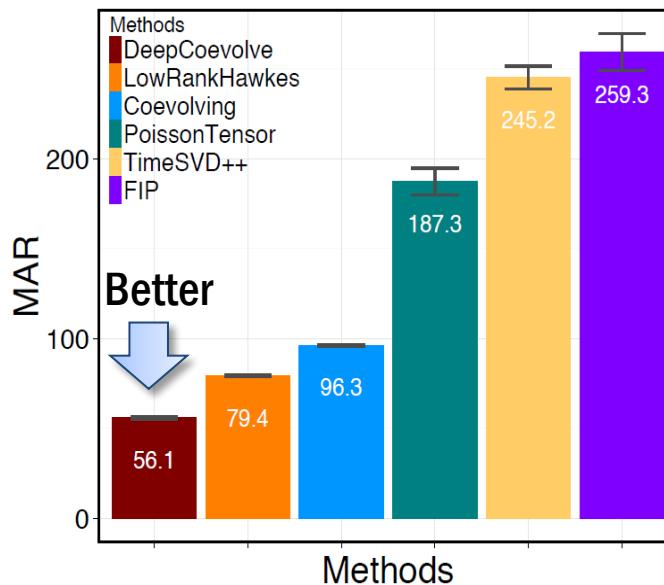
5. Train with
MLE or GAN

IPTV dataset

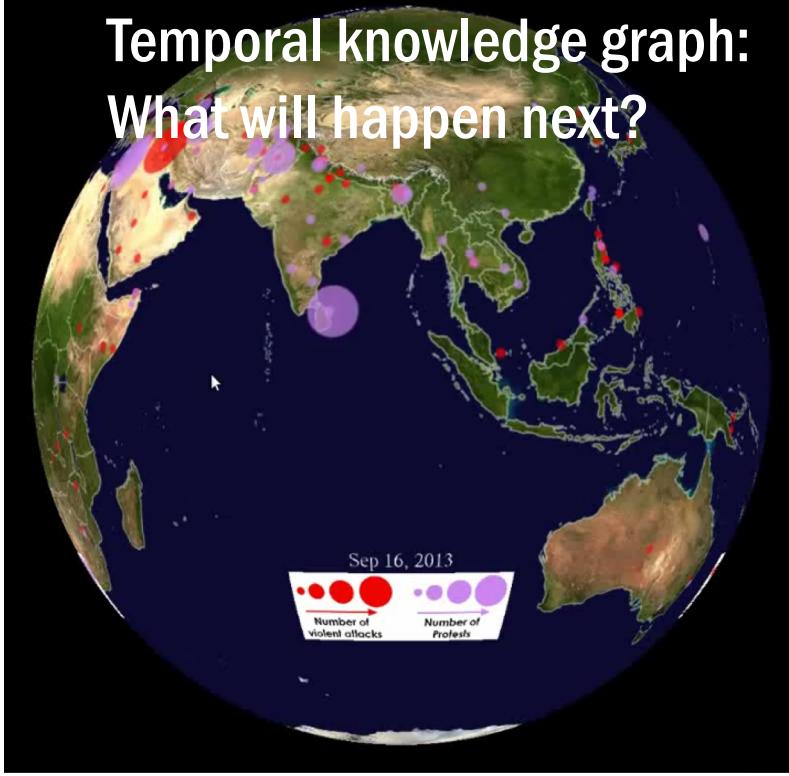
7,100 users, 436 programs, ~2M views

MAR: mean absolute rank difference

MAE: mean absolute error (hours)



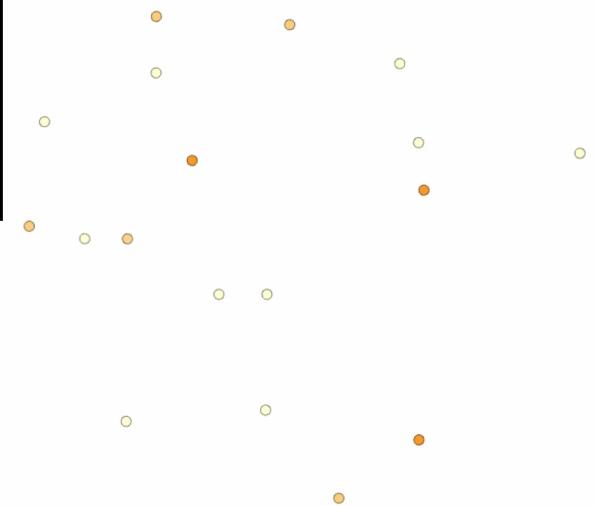
Temporal knowledge graph: What will happen next?



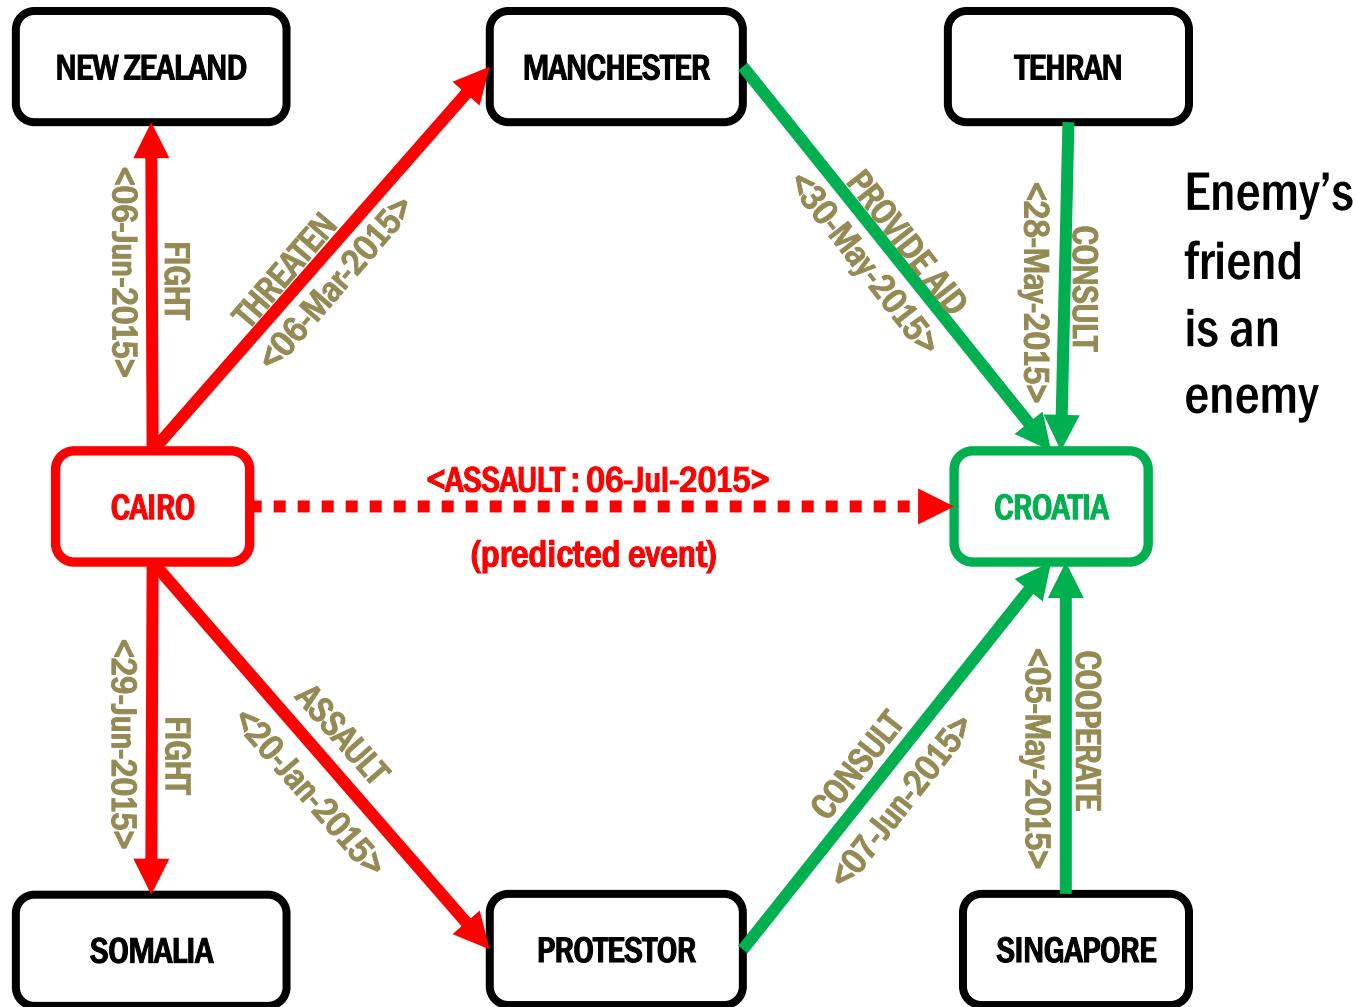
GDELT database

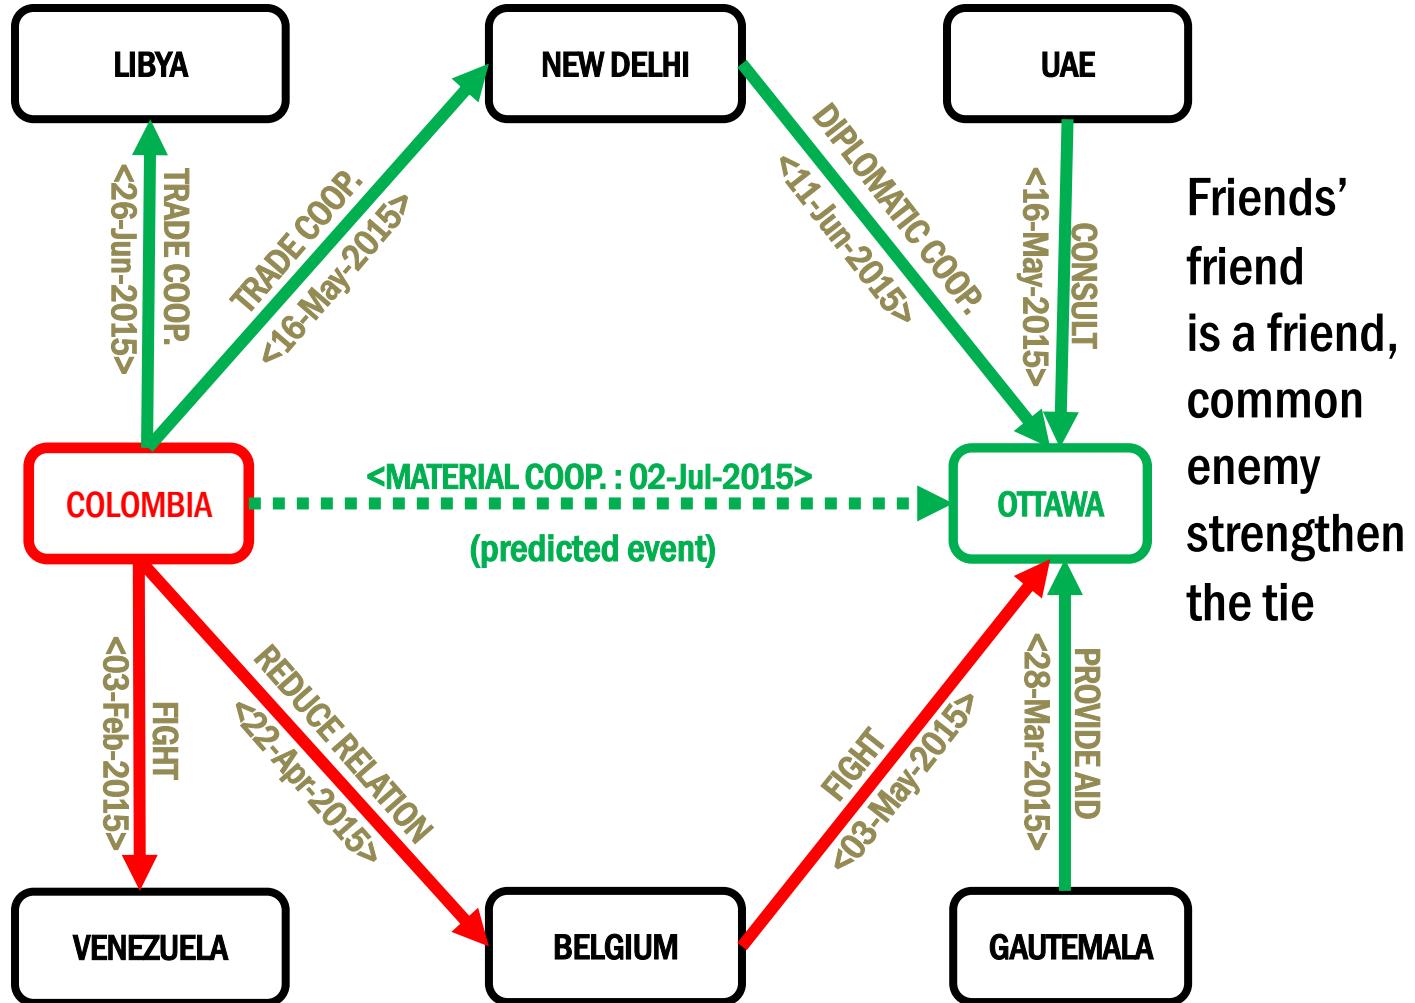
Events in news media
subject – relation – object
and time

Total archives span >215
years, trillion of events



Time-varying dependency structure





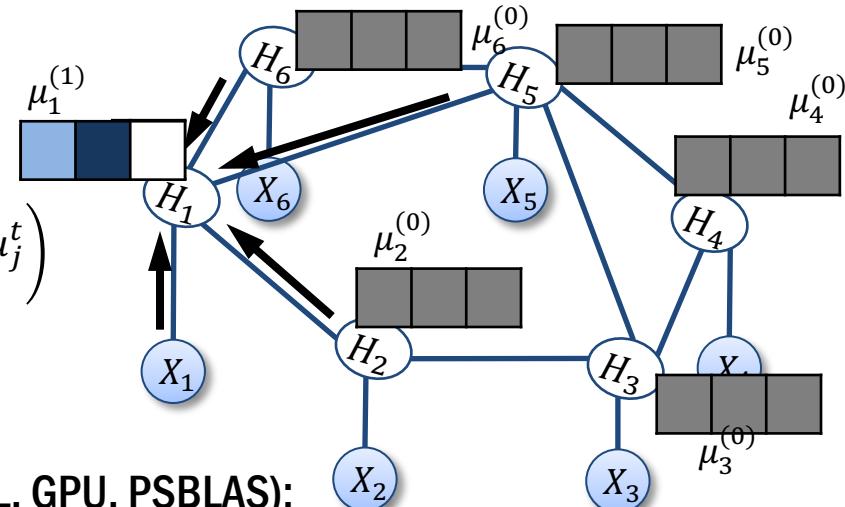
Large Scale Implementation

Sparse matrix formulation

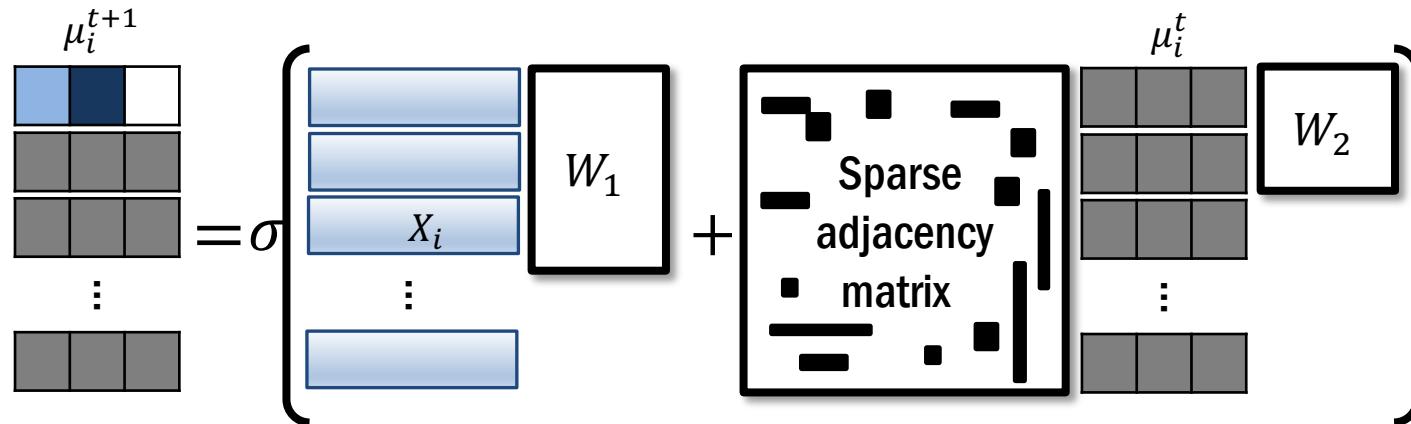
Iterative node update formula:

$$\mu_i^{t+1} \leftarrow \sigma \left(W_1 X_i + W_2 \sum_{j \in \mathcal{N}(i)} \mu_j^t \right)$$

$d \times n$ matrix $d \times d$ matrix



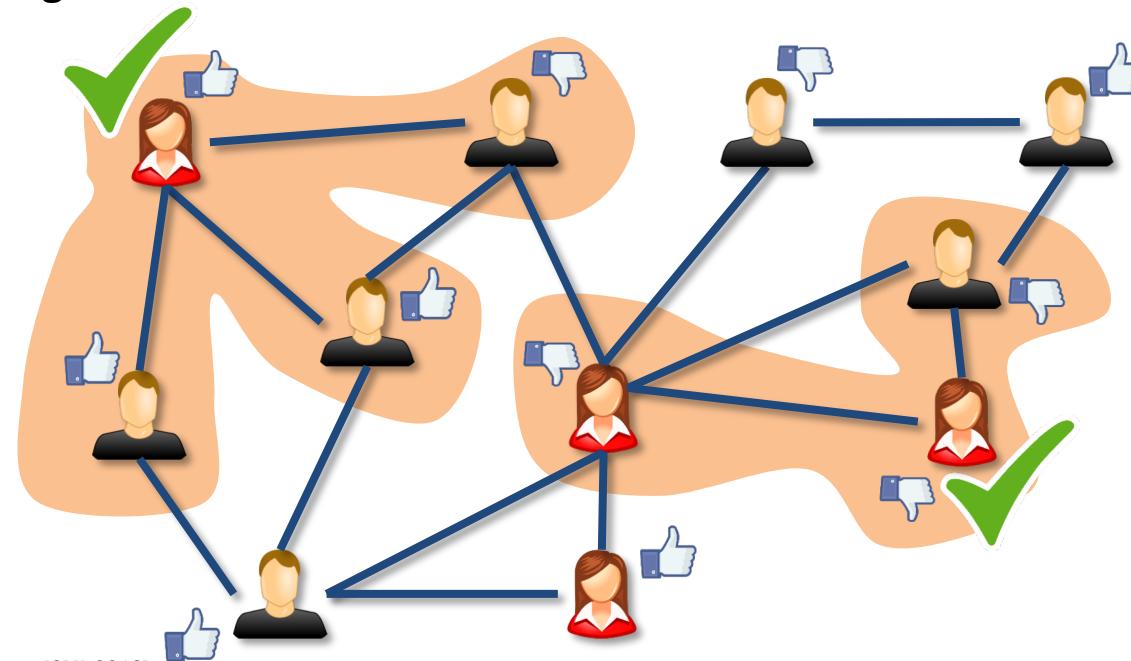
Sparse matrix operation (MKL, GPU, PSBLAS):



Stochastic training

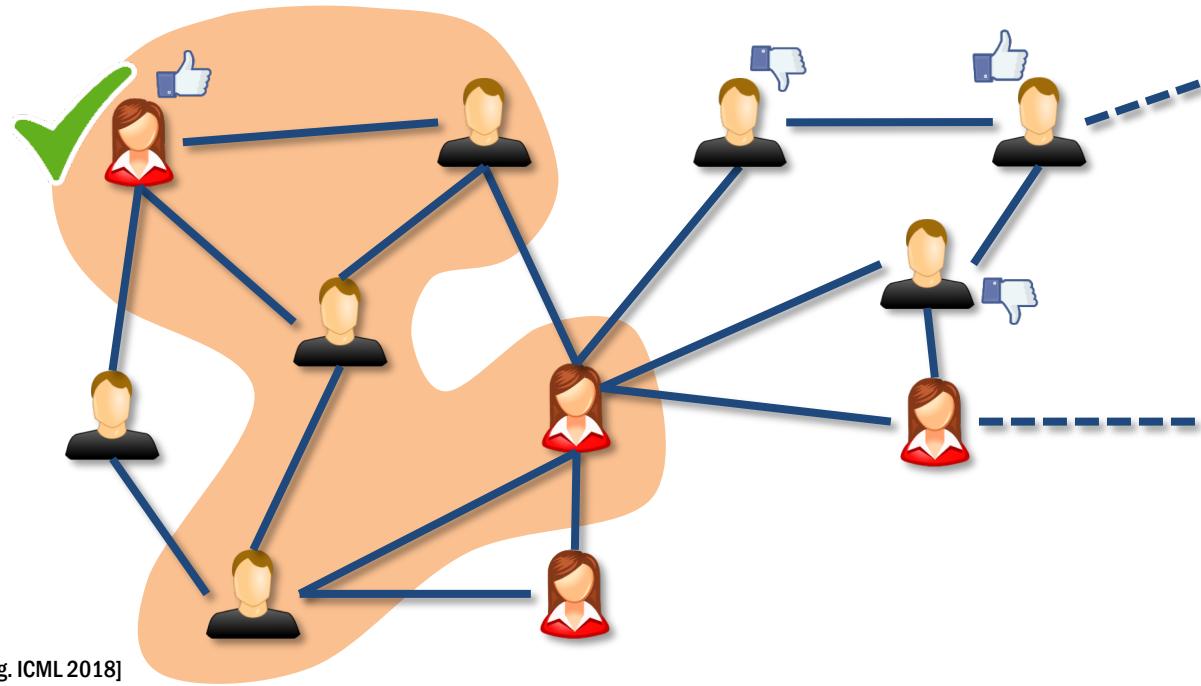
1. Use a mini-batch of nodes for a stochastic loss function
2. Propagation step T determines the subgraph to compute
3. Embedding updates on subgraph

Eg. $T = 1$



Doubly stochastic training

1. Sample a node
2. Propagation T steps to obtain a subgraph (eg. $T = 2$)
3. **Subsample the subgraph**
4. Compute loss using the subsampled subgraph



Variance reduction training algorithm

- Intuition: model parameters W change slowly, so are embeddings
- Idea: approximate embeddings by their historical values
- Maintain history embedding $\bar{\mu}_j^{(t)}$, and $\Delta\mu_j^{(t)} = \mu_j^{(t)} - \bar{\mu}_j^{(t)}$

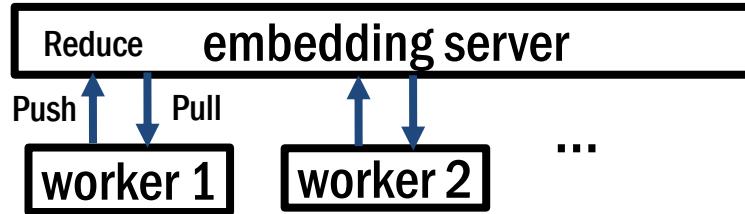
$$\sigma \left(W_1 \mu_i^{(t)} + W_2 \underbrace{\sum_{j \in \mathcal{N}(i)} \mu_j^{(t)}}_{\text{Rewrite}} \right)$$

$$\text{Rewrite} \quad = \sum_{j \in \mathcal{N}(i)} (\Delta\mu_j^{(t)} + \bar{\mu}_j^{(t)})$$

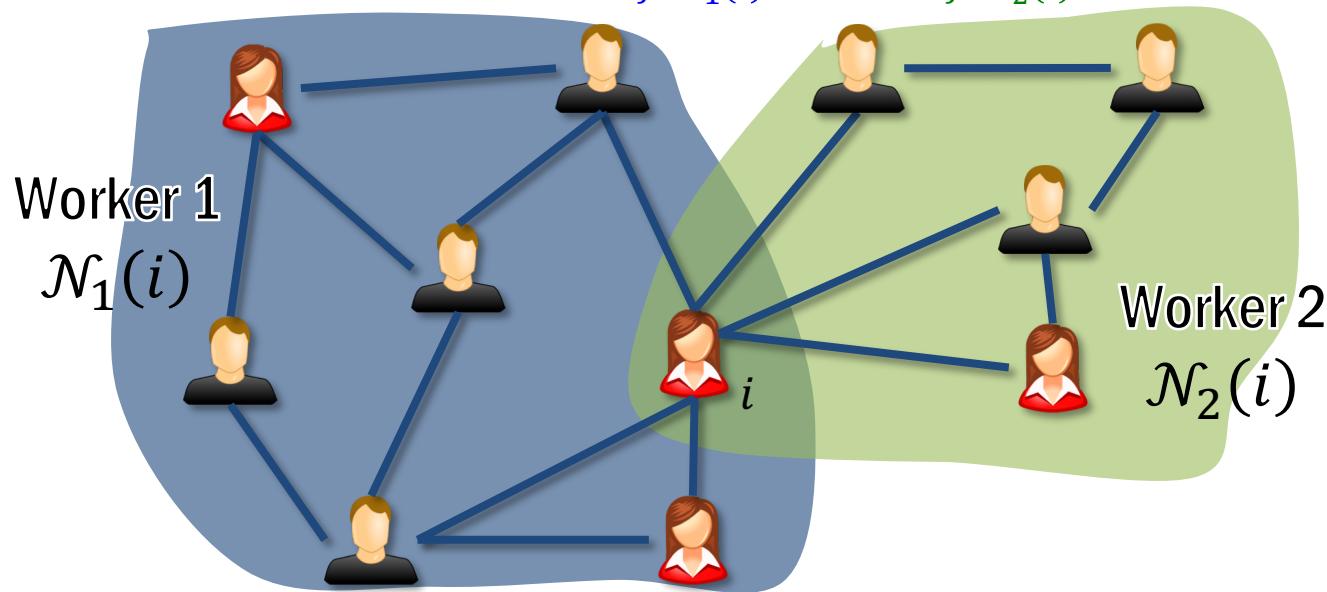
1. $\Delta\mu_j^{(t)} \rightarrow 0$
2. Variance $\rightarrow 0$
3. Gradient unbiased asymptotically

$$\text{Subsample} \quad \approx \frac{|\mathcal{N}(i)|}{|\mathcal{N}'(i)|} \sum_{j \in \mathcal{N}'(i)} \Delta\mu_j^{(t)} + \sum_{j \in \mathcal{N}(i)} \bar{\mu}_j^{(t)}$$

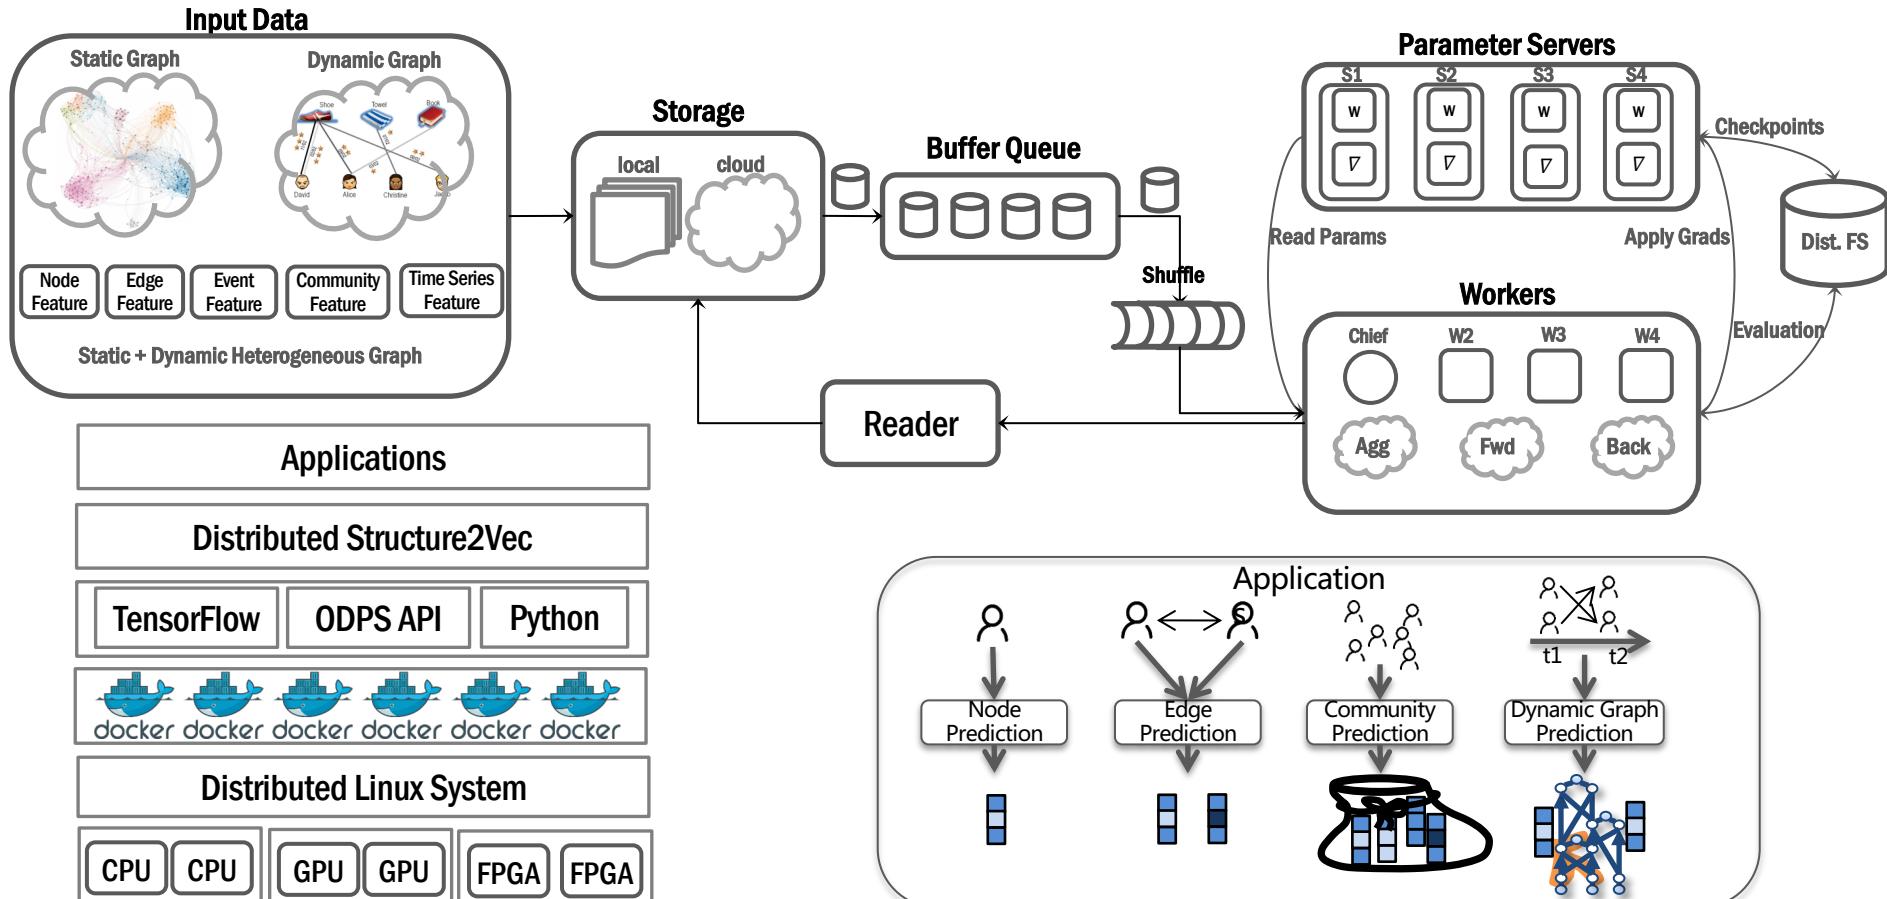
Parameter server architecture



$$\mu_i^{t+1} \leftarrow \sigma \left(W_1 X_i + W_2 \sum_{j \in \mathcal{N}_1(i)} \mu_j^t + W_2 \sum_{j \in \mathcal{N}_2(i)} \mu_j^t \right)$$



Distributed Platform of Structure2Vec

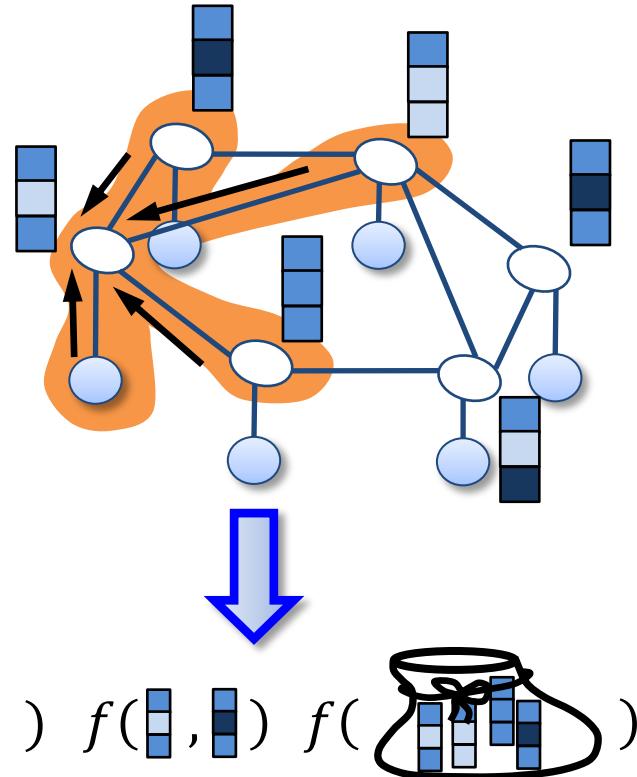


Conclusion

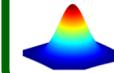
GCN/GNN/Structure2Vec = Parametrized algorithm

Open new possibility to bridge
Deep learning &
Structures (Graph, Logic, Algorithm)

$$\mu_i^{(t)} \leftarrow \sigma \left(W_1 X_i + W_2 \sum_{j \in \mathcal{N}(i)} \mu_j^{(t-1)} \right)$$



Supervised
Learning



Generative
Models



Reinforcement
Learning

$$f(\quad)$$

$$f(\quad , \quad)$$

$$f(\quad , \quad , \quad)$$

References

- H. Dai, B. Dai and L. Song. Structure2Vec: Discriminative Embedding of Latent Variable Models for Structured Data. ICML 2016.
- Y. Zhang, X. Chen, Y. Yang, A. Ramamurthy, B. Li, Y. Qi and L. Song. Can Graph Neural Networks Help Logic Reasoning? Arxiv 2019.
- J. Chen, J. Zhu, and L. Song. Stochastic Training of Graph Convolutional Networks. ICML 2018.
- R Trivedi, H Dai, Y Wang, L Song. Know-Evolve: Deep Temporal Reasoning for Dynamic Knowledge Graph. ICML 2017.
- H Dai, Y Wang, R Trivedi, L Song. Deep Coevolutionary Network: Embedding User and Item Features for Recommendation. Recsys Workshop on Deep Learning. 2017. (BEST PAPER)
- H. Dai, E. Khalil, Y. Zhang, B. Dilkina and L. Song. Learning Combinatorial Optimizations over Graphs. NIPS 2017.
- X. Xu, C. Liu, Q. Feng, H. Yin, L. Song, D. Song. Neural Network-based Graph Embedding for Cross-Platform Binary Code Similarity Detection. CCS. 2017.
- H. Dai, Z. Kozareva, B. Dai, A. Smola, and L. Song. Learning Steady States of Iterative Algorithms over Graphs. ICML 2018.
- H. Dai, H. Li, T. Tian, X. Huang, L. Wang, J. Zhu, and L. Song. Adversarial Attack on Graph Structured Data. ICML 2018.
- Y. Zhang, H. Dai, Z. Kozareva, A. Smola, and L. Song. Variational Reasoning for Question Answering with Knowledge Graph. AAAI 2018.
- Z. Liu, C. Chen, X. Yang, J. Zhou, X. Long and L. Song . Heterogeneous Graph Neural Networks for Malicious Account Detection. CCS2017, CIKM 2018.
- Z. Liu, C. Chen, L. Li, J. Zhou, X. Li and L. Song. Geniepath: Graph Neural Networks with Adaptive Receptive Paths. AAAI 2019.
- H. Dai, Y. Tian, B. Dai, S. Skiena and L. Song. Syntax Directed Variational Autoencoder for Structured Data. ICLR 2018.
- X. Si, H. Dai, M. Raghavanman, M. Naik and L. Song. Learning Loop Invariants for Program Verification. NIPS 2018.